

# The Generic SDLC: Applying STL Principles to AI-Native Engineering

Date: February 24, 2026 Theme: Architectural Restatement of the AI SDLC (v2.8)

---

## The Crisis of the “Agentic” Pipeline

In the first wave of AI-assisted software engineering, we attempted to model the SDLC as a sequence of specialized personas. We built “Requirements Agents,” “Design Agents,” and “Coding Agents,” each with its own hard-coded rules and prompts. This was the “Object-Oriented” phase of AI SDLC—a collection of monolithic classes, each heavily coupled to a specific stage of the development pipeline.

This model failed for the same reason monolithic OO designs often fail: **rigidity**. If you wanted to add a new asset type (like a Security Audit or an API Spec), you had to build a new agent. If the technology stack shifted, every agent’s “expertise” had to be manually retrained.

The breakthrough came when we stopped viewing the SDLC as a sequence of people and started viewing it as **Generic Programming**.

## The STL Analogy: Orthogonality in the SDLC

The C++ Standard Template Library (STL), pioneered by Alexander Stepanov, revolutionized software by decoupling **Algorithms** from **Data Structures**. It proved that an algorithm like `std::sort` didn’t need to know *what* it was sorting, provided the data structure exposed a standard interface (Iterators) and a comparator (Predicates).

The AI SDLC (v2.8) is the application of this “Generic” principle to the entire engineering lifecycle.

### 1. The Algorithm: The Universal `iterate()` Engine

In our model, there is only **one operation**: `iterate()`. This is the universal engine of the methodology. It is “blind” to the domain. It doesn’t know if it is writing a React component or a high-level design document. It only knows how to sense a **Delta ( $\delta$ )** between a current state and a target state, and how to produce a new candidate that reduces that delta.

## 2. The Container: The Asset Graph

The “Data Structures” are our **Asset Types** (Intent, Requirements, Design, Code). These are nodes in a directed cyclic graph. Instead of travelling through a pipeline, a software feature is a **Composite Vector**—a trajectory through this graph. The graph is zoomable and extensible; you can “fold” or “unfold” its complexity without changing the underlying algorithm.

## 3. The Iterator: Admissible Transitions

Iterators in the STL provide a way to traverse containers. In the AI SDLC, **Edges** define the admissible transitions between assets. They bridge the gap between “Design” and “Code,” providing the algorithm with the specific context needed to move from one node to the next.

## 4. The Universal Functor: The Properly Constrained LLM

This is the heart of the system. In the STL, you pass a “Functor” or “Predicate” to an algorithm to define the logic of comparison. In the AI SDLC, the **LLM is the Universal Functor**.

Because the LLM can reason over non-numeric data, it acts as a **Generic Predicate** that can evaluate any edge. However, an unconstrained LLM is prone to hallucination (probability degeneracy). To make it a “Properly Constrained Evaluator,” we surround it with a **Constraint Surface** (the Context []): \* **REQ Keys**: Provide the coordinate system for traceability. \* **ADRs**: Define the formal boundary conditions. \* **Markov Criteria**: Define the objective “stop condition” for stability.

By parameterizing the universal `iterate()` engine with these LLM-driven functors, we achieve a system that can process any “type” of software artifact with formal rigor.

## Functor Escalation: $F_{Dof}F_{Pof}F_H$

One of the most powerful features of generic programming is the ability to specialize algorithms for performance. Our design implements this through **Natural Transformations** of the functor:

- **Deterministic ( $F_D$ )**: When ambiguity is zero, we use cheap, fast “specialized” functors like compilers and linters.
- **Probabilistic ( $F_P$ )**: When ambiguity exists, we escalate to the “generic” LLM functor.
- **Human ( $F_H$ )**: When ambiguity is persistent, we escalate to the human judgment functor.

This “Escalation Chain” ensures the system is as fast as a traditional build tool when things are clear, and as thoughtful as an architect when things are ambiguous.

## Projections: The Methodology as a Generator

Finally, the generic model allows for **Projections**. Just as a C++ template is only instantiated when needed, the AI SDLC is a **Methodology Generator**.

By choosing a **Profile** (e.g., “Minimal” or “Standard”), you are essentially “compiling” a specific version of the methodology. A “Minimal” projection might collapse the Requirements and Design nodes into a single edge, using only the Agent functor. A “Full” projection expands the graph to include 24/7 Sensory Monitoring and Multi-Agent Coordination.

The logic is **Logically Complete** in the blueprint, but only as complex as the project demands at runtime.

## Conclusion: From Tools to Organisms

The shift from specialized agents to **Generic SDLC Programming** marks the transition from “AI Tools” to “Homeostatic Systems.” By decoupling the **Process** (Iteration) from the **Constraint Surface** (Spec + Context), we have built a methodology that doesn’t just assist the developer—it regulates itself.

In this new paradigm, the engineering task is no longer “writing code.” It is **shaping the constraint surface** so that the Universal Evaluator can converge the Asset Graph toward a stable, high-quality product.