

F_D Functor Framework — Design & Implementation Guide

Version: 2.1.0 **Implements:** REQ-ITER-003 (Functor Encoding Tracking), REQ-EVAL-002 (Evaluator Composition) **Package:** imp_claude/code/genisis/

Part I: Current State

Everything in Part I is verifiable against the code. Status badges: [IMPLEMENTED], [PARTIAL], [STUB].

1. Conceptual Overview [IMPLEMENTED]

The functor framework executes the `iterate()` primitive. The primitive has two sub-operations — **construct** (generate/modify assets) and **evaluate** (check convergence) — implemented by three categories of functors.

Three Functor Categories

Category	Symbol	Meaning	Cost	Example
F_D	Deterministic	Subprocess, regex, dict lookup	Free	Run pytest, parse REQ tags
F_P	Probabilistic	LLM / agent call	\$\$\$	“Does this code meet the criterion?”
F_H	Human	Interactive prompt	Blocking	“Do you approve this design?”

Eight Functional Units

The loop decomposes into 8 functional units, each rendered by one of {F_D, F_P, F_H}:

Unit	Default Category	What it does
evaluate	F_D	Run convergence checks (tests, lints, schemas)
construct	F_P	Generate/modify assets (code, specs, tests)
classify	F_D	Categorise signals (REQ tags, source findings)
route	F_H	Select next edge, determine profile
propose	F_P	Suggest changes, generate candidates
sense	F_D	Monitor workspace health (stalls, freshness, coverage)
emit	F_D	Write events to JSONL log — always deterministic
decide	F_H	Final approval/rejection — always human

Two units are **category-fixed**: `emit` is always F_D (the LLM cannot suppress event logging), `decide` is always F_H (humans own final approval).

| Source: `models.py:13–28` — `FunctionalUnit` enum, `CATEGORY_FIXED` dict.

Two Execution Strategies

Two strategies implement `iterate()` today, with a third planned:

Strategy	Status	Construct	Evaluate	LLM Calls
A: E2E Agent	[IMPLEMENTED]	LLM in-session	LLM in-session + subprocess	1 session
B: Deterministic Engine	[PARTIAL]	Not implemented	F_D subprocess + F_P per-check	37 per iteration
C: Hybrid Engine	[PLANNED]	1 LLM call per edge	Batched in same call	4 per iteration

See §13 for full analysis.

2. Profile-Based Routing [IMPLEMENTED]

The framework uses **named profiles** to control which graph edges are active, which evaluators run, and how functional units are encoded. Six profiles cover the projection space.

Profile Selection

`select_profile()` maps feature type to profile name:

```
# fd_route.py:select_profile()
PROFILE_MAP = {
    "feature": "standard",
    "discovery": "poc",
    "spike": "spike",
    "poc": "poc",
    "hotfix": "hotfix",
}
# default → "standard"
```

| Source: fd_route.py — select_profile(feature_type, profiles_dir) —
| > str

What the Engine Reads from a Profile

Each profile YAML (`config/profiles/{name}.yaml`) provides three things the engine uses:

1. **graph.include** — required edge traversal order
2. **graph.optional** — optional edges (selected only after required edges converge)
3. **encoding.functional_units** — category (F_D/F_P/F_H) per functional unit

Encoding Matrix (from Profile YAMLs)

Unit	full	standard	poc	spike	hotfix	minimal
evaluate	F_D	F_D	F_P	F_D	F_D	F_D
construct	F_P	F_P	F_P	F_P	F_P	F_P
classify	F_D	F_D	F_D	F_D	F_D	F_D
route	F_H	F_D	F_H	F_P	F_D	F_P
propose	F_P	F_P	F_P	F_P	F_P	F_P
sense	F_D	F_D	F_D	F_D	F_D	F_D
emit	F_D	F_D	F_D	F_D	F_D	F_D

Unit	full	standard	poc	spike	hotfix	minimal
decide	F_H	F_H	F_H	F_H	F_H	F_H

Source: config/profiles/*.yml — encoding.functional_units section of each file.

Graph Subsets by Profile

Profile	Required Edges	Optional
full	all 10 edges	none
standard	intent→req, req→design, design→code, code↔unit_tests	design→test_cases, design→uat_tests, code→cicd
poc	intent→req, req→design, design→code	none
spike	intent→req, req→design, design→code	none
hotfix	intent→req, design→code, code↔unit_tests	none
minimal	intent→req, design→code	none

Which Strategy Reads What

Profile Field	Engine (Strategy B)	E2E Agent (Strategy A)
graph.include / optional	Yes — select_next_edge()	Yes — agent reads profile
encoding.functional_units	Yes — lookup_encoding() → dispatch	No — agent is the LLM
evaluators	No	Yes — determines check density
convergence	No	Yes — drives convergence rules
context	No	Yes — determines context loading


```

        FunctionalUnit.DECIDE: Category.F_H,    # decide is ALWAYS human
    }

```

This is enforced at all profile levels and validated in tests.

5. Configuration Resolution Pipeline [IMPLEMENTED]

The configuration system composes constraints from four layers:

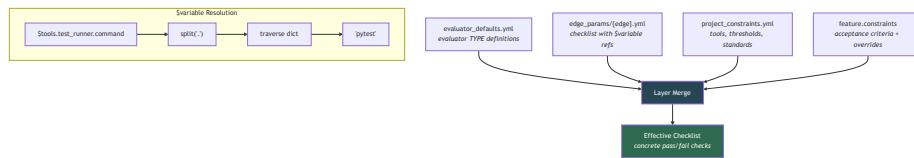


Diagram 2

\$Variable Resolution Regex

Pattern: `\$(\w+(?::.\w+)*)`

Matches:

```

$tools.test_runner.command      → constraints["tools"]
["test_runner"]["command"]
$thresholds.test_coverage_minimum → constraints["thresholds"]
["test_coverage_minimum"]
$standards.style_guide          → constraints["standards"]
["style_guide"]

```

Resolution Rules

1. Edge checklist defines default checks
2. \$variables resolve from project_constraints.yml
3. Feature threshold_overrides apply on top
4. Feature acceptance_criteria append to checklist
5. required=true at any layer stays true (most restrictive wins)
6. Unresolved \$variables → check **SKIPPED** with warning (tracked in unresolved[])

Sequence: resolve_checklist()

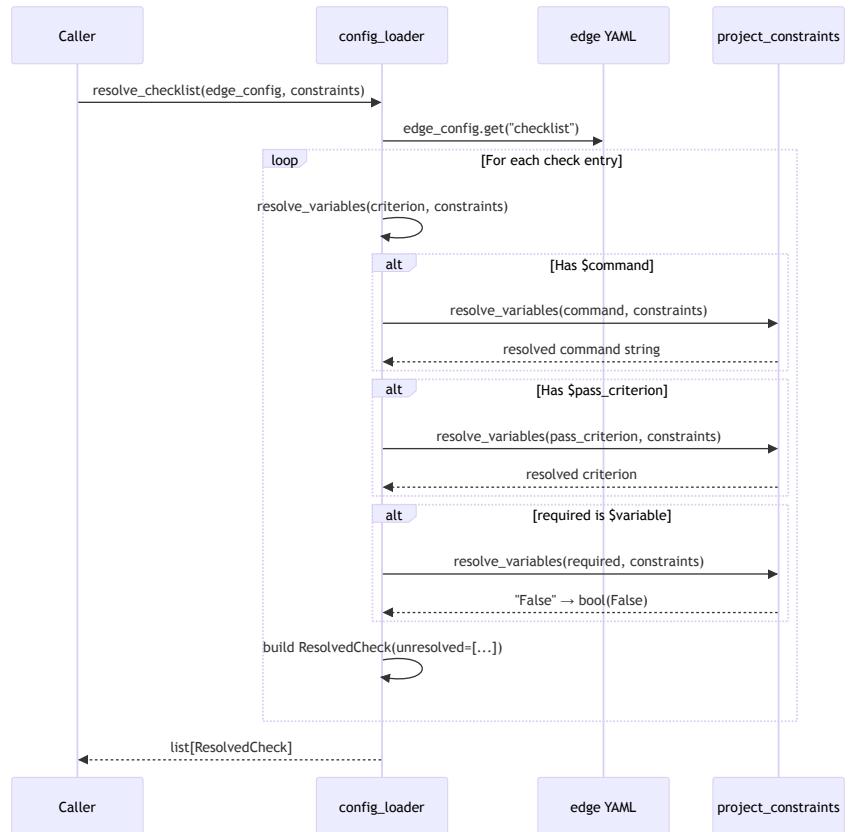


Diagram 3

6. F_D Evaluate – The Deterministic Evaluator [IMPLEMENTED]

State Machine: Check Execution

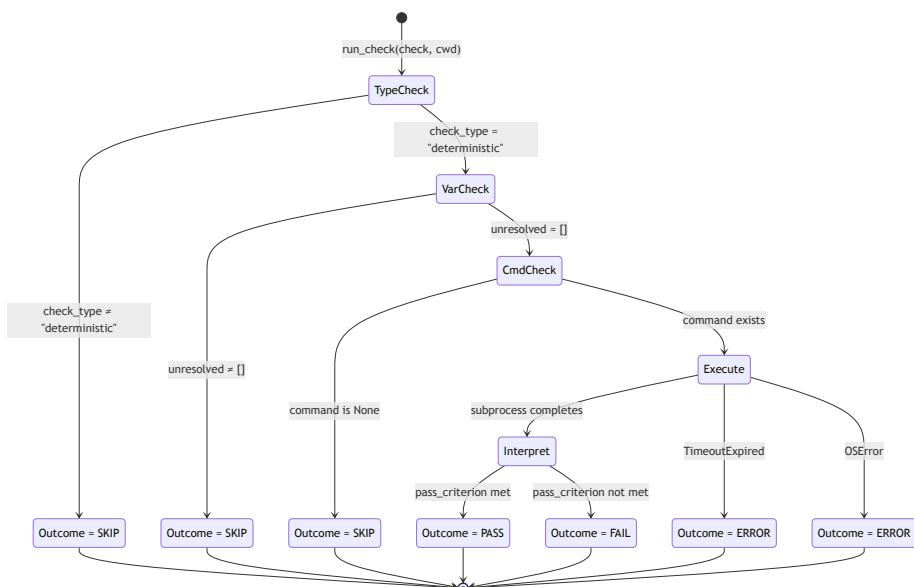


Diagram 4

Pass Criterion Interpretation

The `_interpret_result()` function interprets subprocess output against the `pass_criterion` string:

Criterion Pattern	Interpretation
"exit code 0" (or empty)	<code>returncode == 0 → PASS</code>
"coverage percentage >= N"	Parse <code>(\d+)%</code> from stdout, compare to threshold
"zero violations" / "zero errors"	<code>returncode == 0 → PASS</code>
Default fallback	<code>returncode == 0 → PASS</code>

Checklist Aggregation

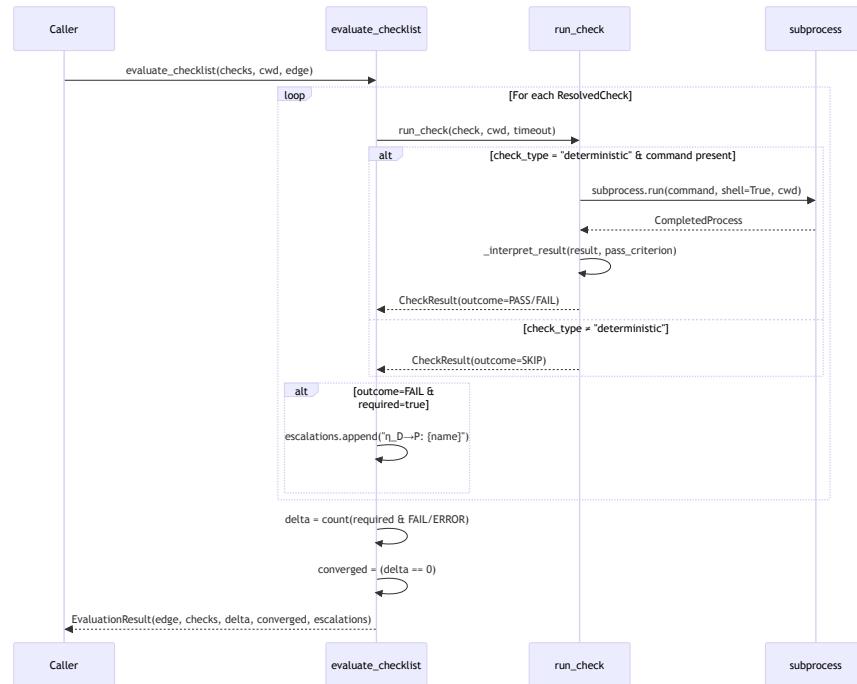


Diagram 5

Delta Formula

```

delta = Σ{ 1 | check ∈ checks, check.required ∧ check.outcome ∈ {FAIL, ERROR} }
converged = (delta == 0)

```

- SKIP outcomes (agent, human, unresolved) do **not** count toward delta
- Non-required (`required=false`) failures do **not** count toward delta
- Delta is a **non-negative integer** — the distance from convergence

7. The η (Natural Transformation) — Escalation Boundary [IMPLEMENTED] signals, [STUB] dispatch

When a check fails, the framework surfaces an **escalation signal** that hands off to the next-higher category. This is the natural transformation $\eta: F_D \rightarrow F_P \rightarrow F_H$.

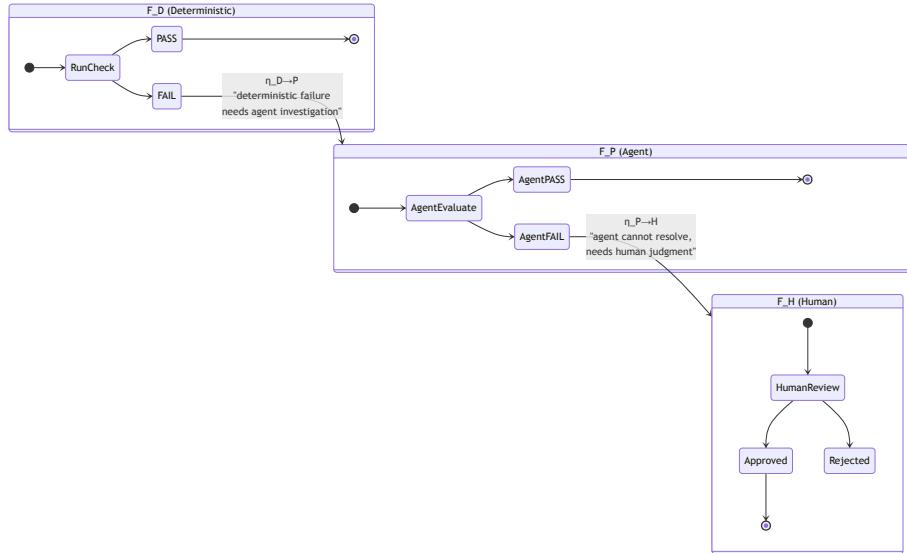


Diagram 6

Where η fires in the code

`fd_evaluate.py` — in `evaluate_checklist()`:

```
if cr.check_type == "deterministic" and cr.outcome in (FAIL, ERROR)
    and cr.required:
    escalations.append(f"\u03b7_D→P: {cr.name} failed - may need agent
investigation")
```

`engine.py` — in `iterate_edge()`:

```
if cr.check_type == "deterministic":
    escalations.append(f"\u03b7_D→P: {cr.name} - deterministic failure")
elif cr.check_type == "agent":
    escalations.append(f"\u03b7_P→H: {cr.name} - agent evaluation failed")
```

The escalation signals are **informational** — the engine records them, but the current implementation does not automatically dispatch to F_P or F_H. Automatic dispatch is future work.

8. F_D Emit — Event Emission [IMPLEMENTED]

Emit is **category-fixed F_D** — it always fires, regardless of profile. The LLM cannot skip it.

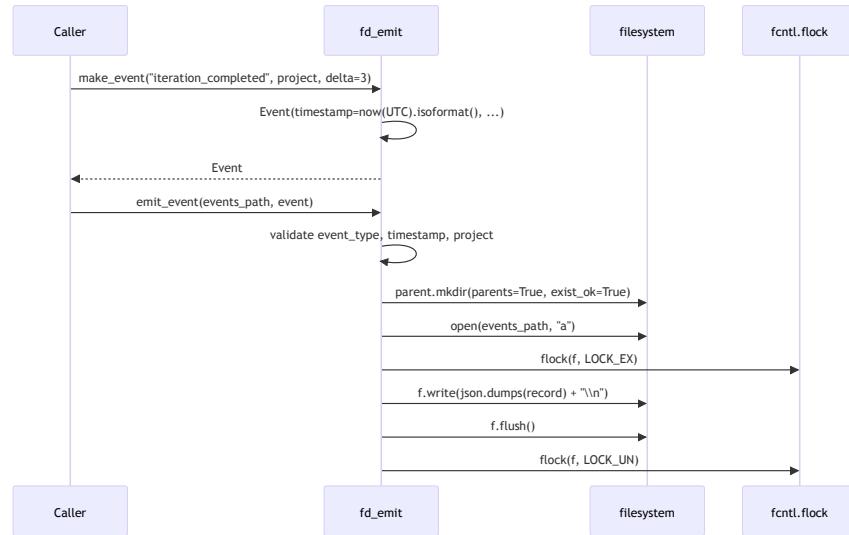


Diagram 7

Event Format (JSONL)

```
{"event_type": "iteration_completed", "timestamp": "2026-02-24T10:30:00+00:00", "project": "my_proj", "feature": "REQ-F-AUTH-001", "edge": "code↔unit_tests", "delta": 3, "status": "iterating", "depth": 0}
```

Event Types

Event Type	When Emitted
project_initialized	/gen-init
iteration_completed	Every iteration boundary
edge_started	Edge traversal begins
edge_converged	All required checks pass
spawn_created	Child vector spawned
spawn_folded_back	Child results returned
checkpoint_created	Session snapshot
review_completed	Human review done
gaps_validated	Traceability check

Event Type	When Emitted
release_created	Release package

9. F_D Classify – Deterministic Classification [IMPLEMENTED]

Three classifiers, all regex/keyword-based (no LLM):

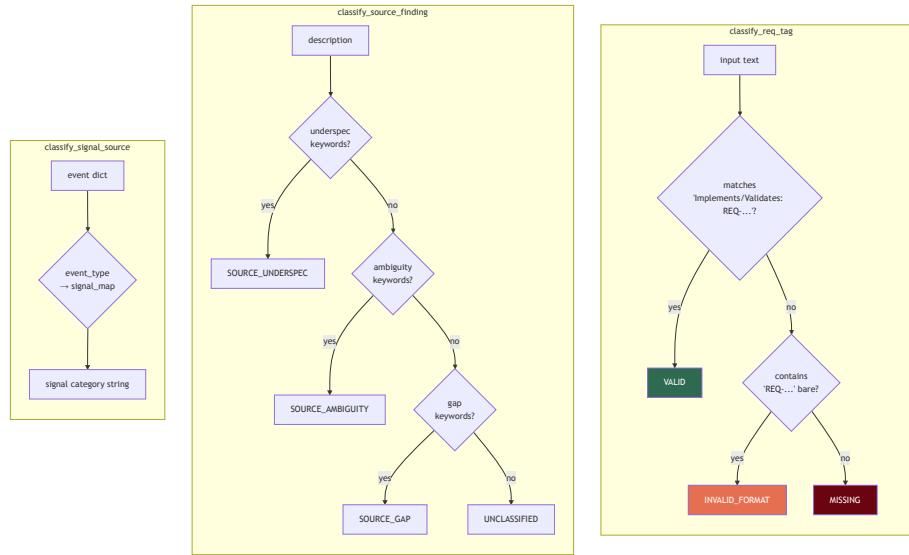


Diagram 8

Keyword Sets

Classification	Keywords
SOURCE_UNDERSPEC	underspecified, insufficient detail, needs clarification, placeholder
SOURCE_AMBIGUITY	unclear, ambiguous, vague, undefined, unspecified, unknown, tbd
SOURCE_GAP	missing, absent, gap, omitted, incomplete, not defined, lacks

Priority order: underspec → ambiguity → gap (first match wins).

10. F_D Sense — Interoceptive Monitors [IMPLEMENTED]

Five monitors map to the spec's sensory system (INTRO-001 through INTRO-007):

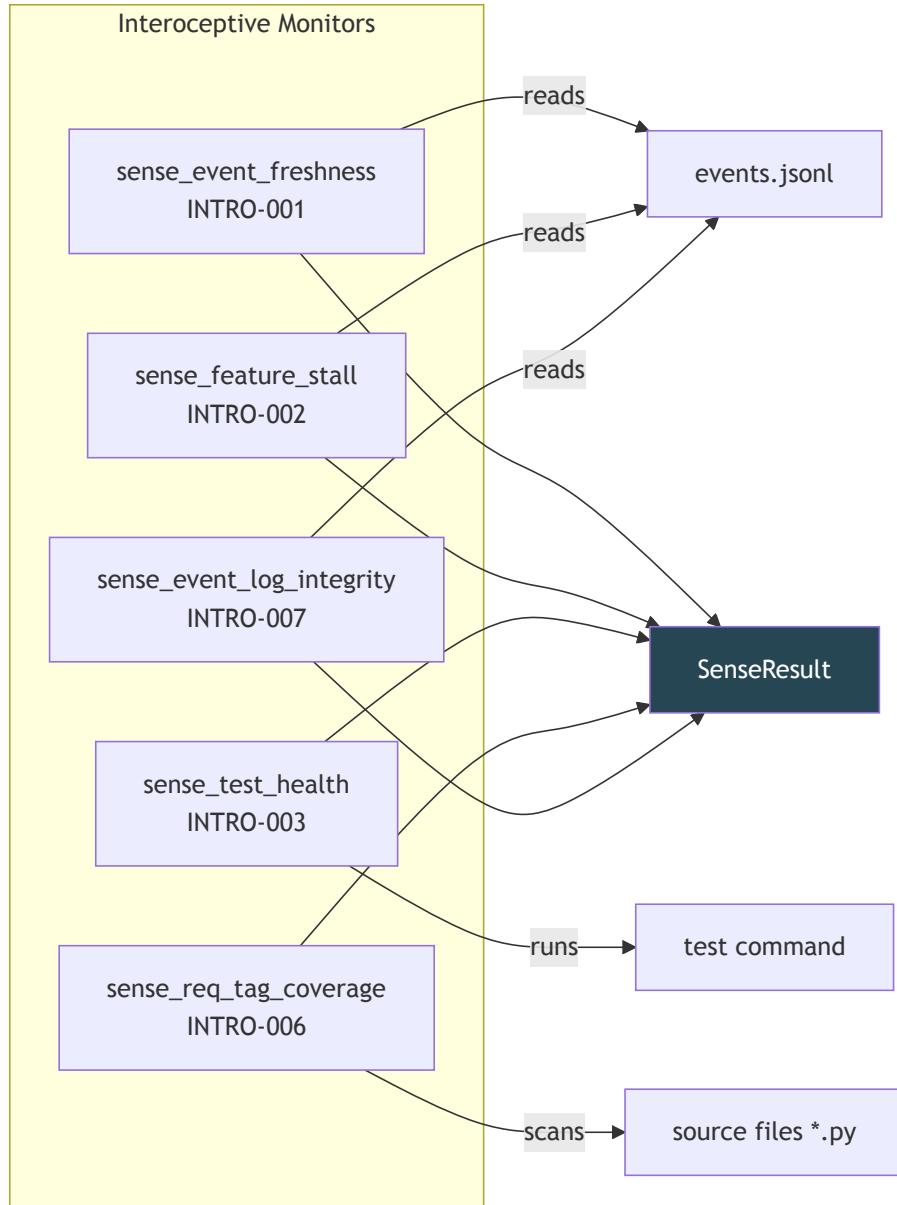


Diagram 9

Stall Detection State Machine

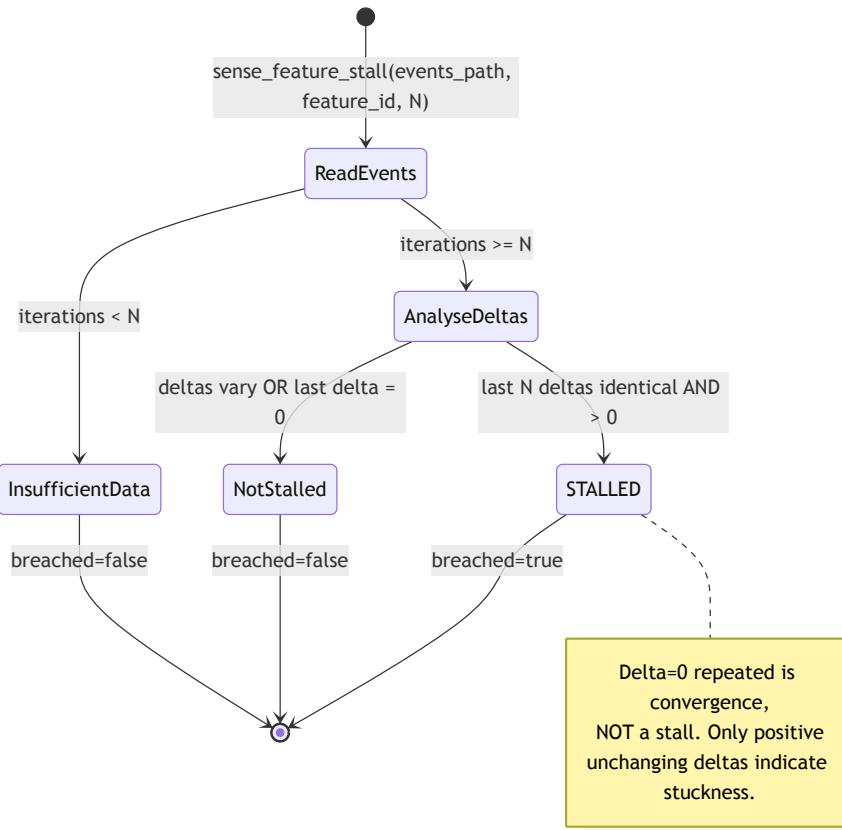


Diagram 10

11. F_D Route – Profile & Edge Selection [IMPLEMENTED]

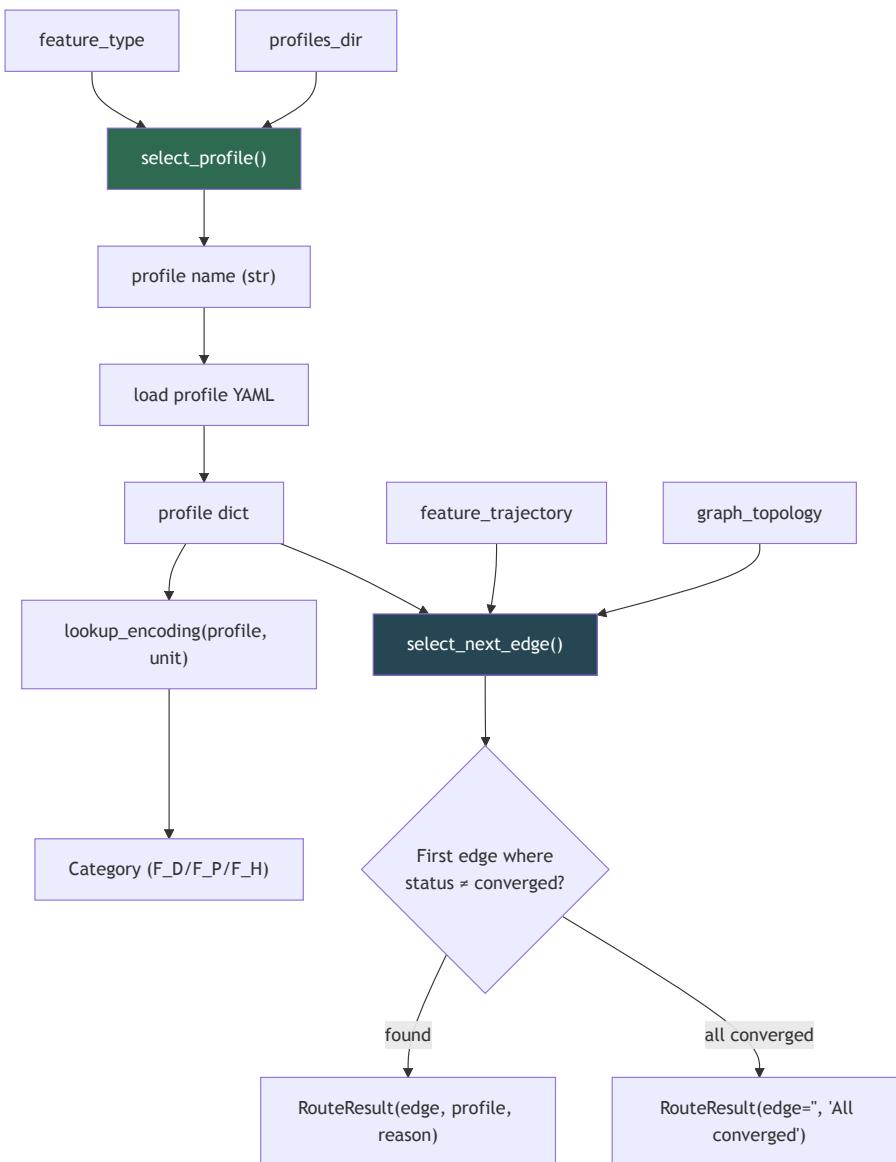


Diagram 11

Source: `fd_route.py` — `select_profile()`, `select_next_edge()`, `lookup_encoding()`.

Edge Naming Convention

Graph edges use Unicode arrows (\rightarrow , \leftarrow). Trajectory keys normalise these:

"code \leftrightarrow unit_tests" \rightarrow trajectory key "code_unit_tests"
"intent \rightarrow requirements" \rightarrow trajectory key "intent_requirements"

12. Dispatch Table [IMPLEMENTED]

The dispatch table maps (FunctionalUnit, Category) to a callable:

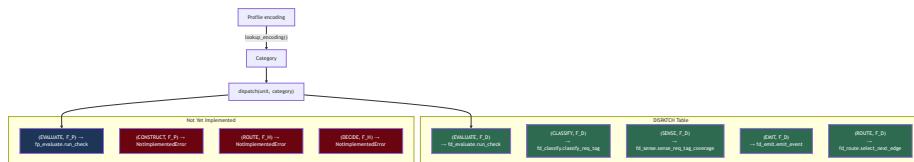


Diagram 12

lookup_and_dispatch(unit, profile) — End-to-End

```
def lookup_and_dispatch(unit: FunctionalUnit, profile: dict) ->
    Callable:
    category = lookup_encoding(profile, unit.value)           # Step 1:
    profile -> category
    return dispatch(unit, category)                            # Step 2:
    table lookup
```

Source: dispatch.py — `lookup_and_dispatch(unit, profile)`, `dispatch(unit, category)`.

13. Execution Strategies & Engine [PARTIAL]

Abstract iterate() Loop

Both strategies implement the same abstract loop:

```
while not converged and budget > 0:
    candidate = construct(asset, context, edge_config)    # F_P:
    generate/modify
    evaluation = evaluate(candidate, evaluators)          # F_D + F_P:
    check
    delta = compute_delta(evaluation)                    # F_D: count
    failures
    emit(iteration_event)                             # F_D:
always fires
    if delta == 0: return candidate                   # converged
    asset = candidate                                # feed back
```

The critical distinction is **who does construct**.

Strategy A: E2E Agent [IMPLEMENTED]

One `claude -p` session drives the **entire** methodology loop. The LLM constructs artifacts, evaluates them (agent checks happen in the same session), and routes to the next edge. The REPL loop operates within a single LLM context window.

```

Single claude -p session

for each edge in profile.graph:
    construct artifact (LLM generates)
    evaluate all checks (in-session)
    emit events
    if converged: route to next edge

```

Property	Value
LLM calls	1 session (all edges)
Construct	Yes — LLM generates artifacts in-session
Context coherence	Full — later edges see earlier artifacts
Cost	~\$2-5 per feature (780s typical)
Testability	Requires live Claude (~13 min per run)
Observability	Opaque — decision path in LLM reasoning

| **Source:** E2E tests at `imp_claude/tests/e2e/test_e2e_convergence.py`.
 Historical runs: 400-890s, \$2-5, 34 tests.

Strategy B: Deterministic Engine [**PARTIAL – evaluate only**]

Python code drives the loop deterministically. The engine resolves configs, dispatches checks by type, computes delta, emits events — but **it cannot construct artifacts**.

```

engine.run(feature_id, feature_type)

profile = select_profile(feature_type)
while edge = select_next_edge():
    for i in range(budget):
        ✗ NO CONSTRUCT STEP
        checks = resolve_checklist(edge)
        for check in checks:
            if deterministic: subprocess
            if agent: claude -p (1 per chk)
            if human: SKIP
        delta = count(required & FAIL)
        emit(iteration_event)
        if converged: break

```

Property	Value
LLM calls	37 per iteration (33 agent evals + 4 construct = 0)
Construct	No — re-evaluates unchanged assets
Context coherence	None — each <code>claude -p</code> is a cold start
Cost	~\$2-8 per run ($37 \times$ cold start overhead)
Testability	F_D: 0.76s, no LLM. Full: 18.9s (agent checks ERROR in test env)
Observability	Full — every step is a Python function call

The engine loop re-evaluates unchanged assets. Without construct, the loop cannot converge on its own — delta stays the same across iterations because nothing changes between them. This is the fundamental gap.

| Source: `engine.py` — `run()`, `run_edge()`, `iterate_edge()`.

Strategy C: Hybrid Engine [PLANNED]

One `claude -p` call per edge that both constructs the artifact AND evaluates all agent checks. Deterministic checks run as subprocess before the LLM call. The engine validates structured output deterministically.

Property	Target Value
LLM calls	4 per iteration (1 per edge)
Construct	Yes — LLM generates per-edge
Context coherence	Per-edge (constructed artifacts from edge N become context for N+1)
Cost	~\$0.50 per feature
Testability	F_D tests unchanged. LLM tests: 4 calls vs 37

See Appendix A for implementation path.

Full Traversal Sequence (Engine – Current)

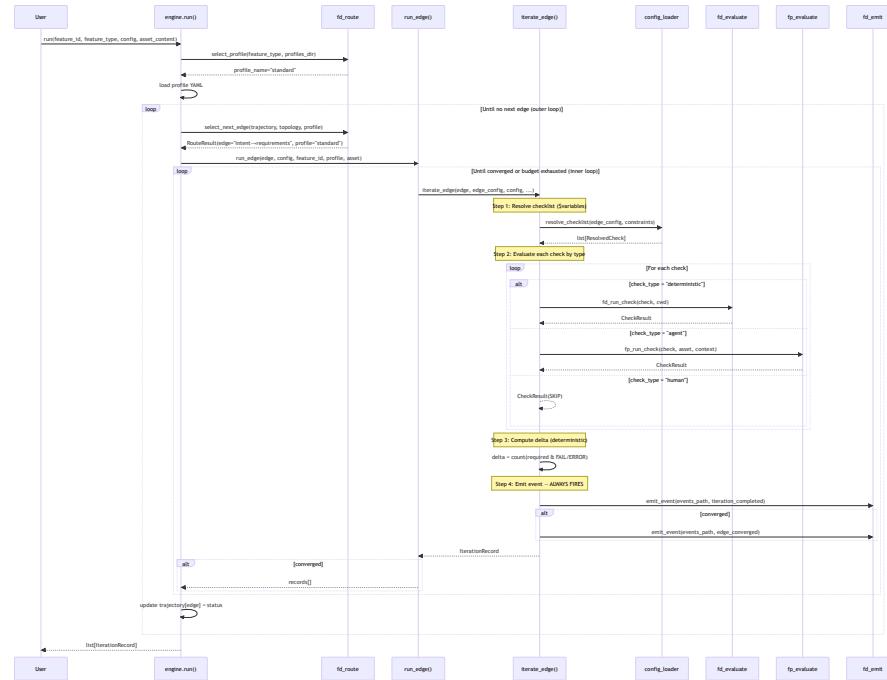


Diagram 13

Engine State Machine

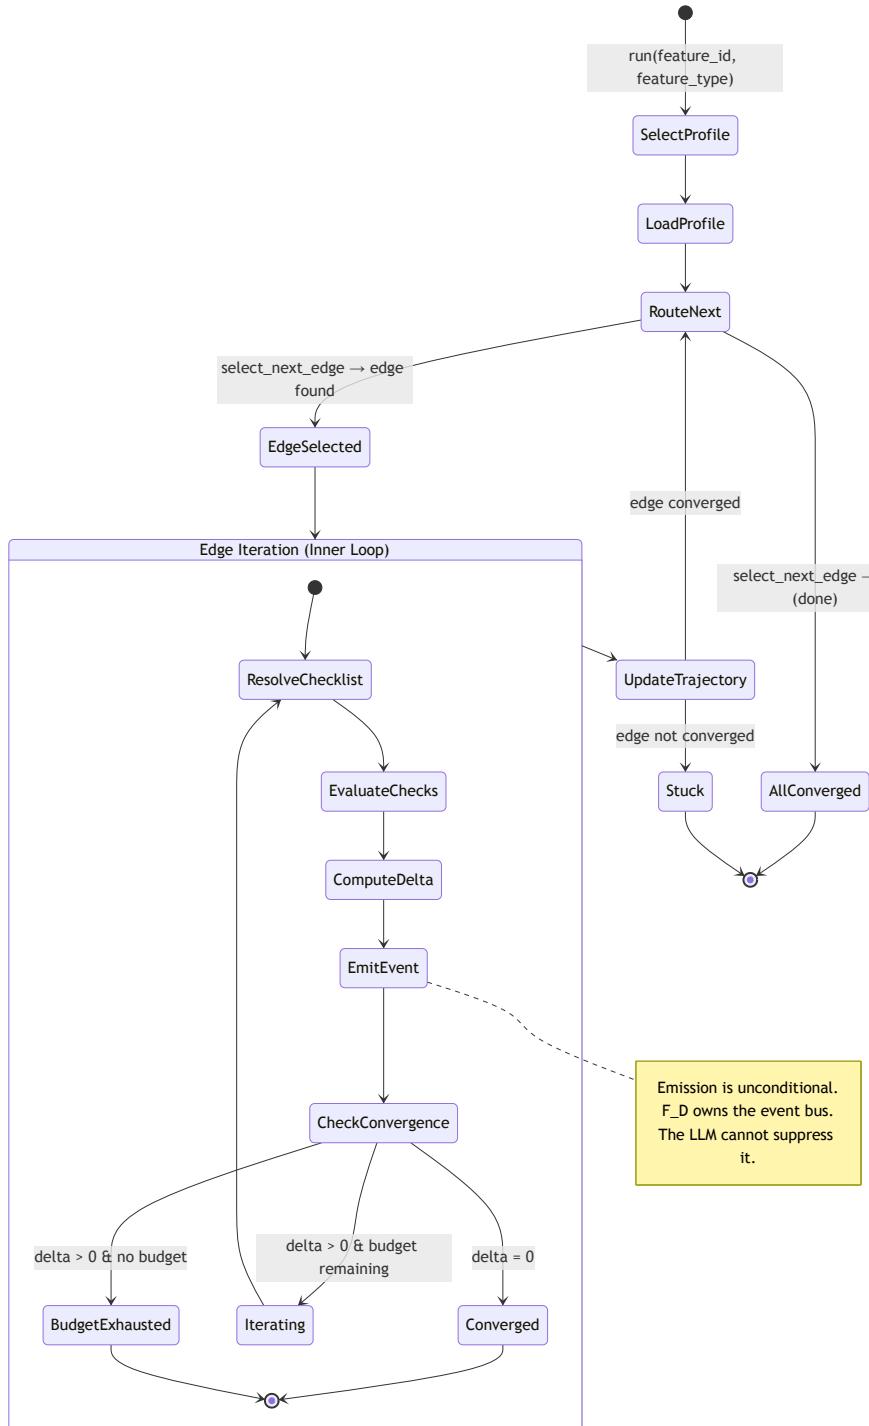


Diagram 14

Three Processing Layers (Cost Model)

The spec (§4.6) defines three processing phases:

Reflex (F_D) – templates, transforms, rules → FREE
(deterministic)

Affect (F_P) – judgment, generation, evaluation → COSTLY (LLM calls)

Conscious (F_H) – decisions, approvals
wait)

→ BLOCKING (human

The engine already has the **full reflex layer** (route, emit, delta, subprocess, classify, sense). The cost question is how to structure the **affect layer's LLM calls** — per-check (Strategy B, 37 calls) vs per-edge (Strategy C, 4 calls) vs in-session (Strategy A, 1 session).

14. Spawn — Child Vectors [IMPLEMENTED] manual, [STUB] automatic

Spawning creates a child vector to investigate a sub-problem. Today this is **manual** via /gen-spawn. The engine does not call run() recursively.

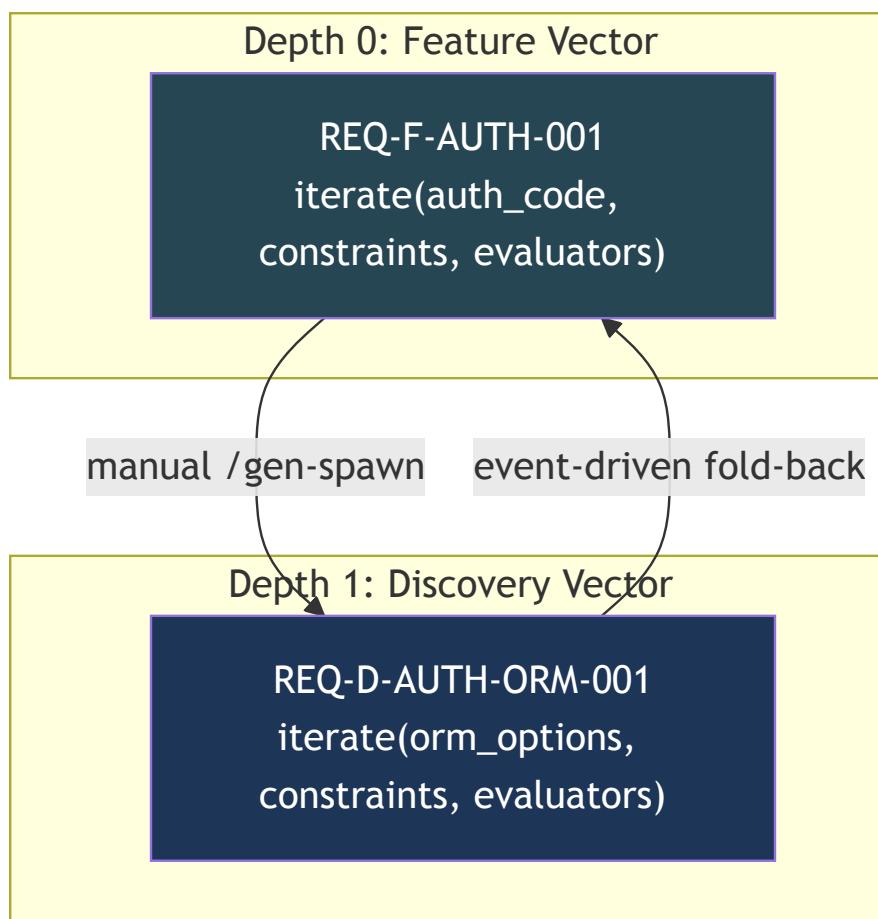


Diagram 15

What exists today: - /gen-spawn creates a child feature vector with its own trajectory - Fold-back via events: child emits spawn_folded_back, parent reads from events.jsonl - Spawn is event-driven and asynchronous — no engine code path

What does not exist: - Engine calling run() recursively with child config - Automatic depth tracking or zoom decay - Programmatic fold-back of child artifacts into parent context

See Appendix D for the target architecture.

15. F_P Evaluate – LLM Integration [IMPLEMENTED]

`fp_evaluate.py` wraps the Claude Code CLI for agent-based evaluation:

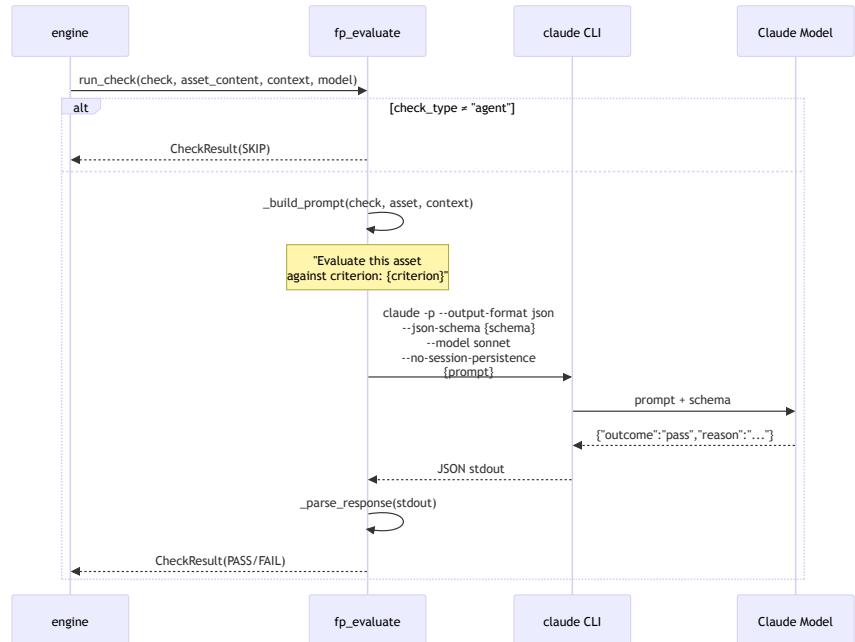


Diagram 16

Response Schema (JSON Schema)

```
{  
    "type": "object",  
    "properties": {  
        "outcome": {"type": "string", "enum": ["pass", "fail"]},  
        "reason": {"type": "string"}  
    "required": ["outcome", "reason"]  
}
```

Key limitation: Each agent check is a **separate claude -p subprocess** with no shared context. Check N does not know check N-1's result. This means 33 cold-start sessions per iteration for the standard profile.

Source: `fp_evaluate.py` — `run_check()`, `_build_prompt()`, `_parse_response()`.

16. Data Flow – Complete Pipeline [PARTIAL]

This diagram shows how data flows through the entire system for one iteration. The construct phase (dashed) is **not implemented** in the engine.

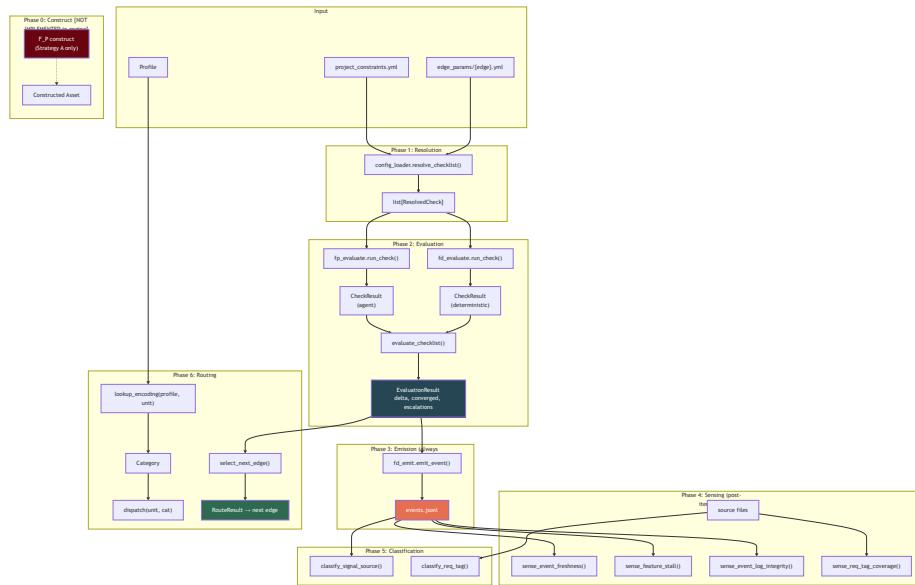


Diagram 17

17. Test Architecture [IMPLEMENTED]

Test Files

File	Tests	What it covers
test_config_loader.py	16	resolve_variable, resolve_variables, resolve_checklist, load_yaml, real config integration
test_functor_fd.py	39	Unit tests for all F_D modules: evaluate, emit, classify, sense, route, dispatch
test_functor_e2e.py	50	6 end-to-end scenarios wiring the full pipeline

E2E Test Scenarios

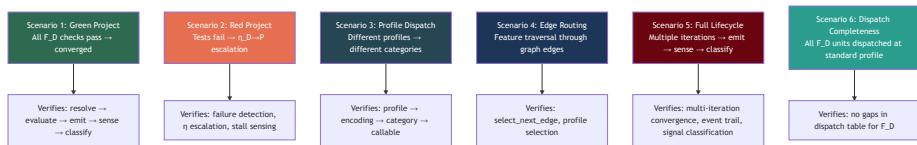


Diagram 18

18. Dependencies [IMPLEMENTED]

stdlib only (except PyYAML):

Dependency	Used by	Purpose
dataclasses	models.py	Data model definitions
enum	models.py	Category, FunctionalUnit, CheckOutcome
subprocess	fd_evaluate.py, fd_sense.py, fp_evaluate.py	Shell command execution
fcntl	fd_emit.py	Advisory file locking for JSONL append
json	fd_emit.py, fd_sense.py, fp_evaluate.py	Event serialization/parsing
re	config_loader.py, fd_classify.py, fd_sense.py	Pattern matching
yaml (PyYAML)	config_loader.py	YAML parsing
shutil	fp_evaluate.py	which() to find claude CLI

19. Process Model Decision: Synchronous Now, Actors Later [IMPLEMENTED]

The design docs (ADR-013, ADR-015, Design §1.11) describe an **actor model** with inbox staging, a single-writer serialiser, MCP sensory service, and observer agents. The current engine uses **synchronous direct function calls** instead.

This is a deliberate decision, not a gap.

Concern	Actor Model	Current Model	Verdict
Correctness	Same	Same	Both produce identical results
Testability	Harder (async, inboxes)	Easy (direct calls, deterministic)	Current wins
Debuggability	Distributed tracing needed	Stack traces just work	Current wins
Multi-agent	Required	Not needed yet	Defer
Robustness	Crash recovery, stale claim detection	Process dies, restart	Defer
MCP sensory service	Long-running monitors	On-demand sense functions	Defer

The actor model adds **robustness** (crash recovery, concurrent agent coordination, long-running monitors) but not **functionality**. Every functor, every η boundary, every event emission, every sense monitor is testable and exercised with the synchronous model.

When to revisit: when multi-agent coordination (ADR-013 claim protocol) or the sensory MCP server (ADR-015) become active implementation targets.

20. What's Not Implemented Yet

Gap	Category	Status	Notes
F_P Construct	F_P	[PLANNED]	Engine cannot generate artifacts. See Appendix A.
Batched F_P Evaluate	F_P	[PLANNED]	1 claudie -p per check (33 cold starts). Target: 1 per edge. See Appendix A.
Context accumulation	Engine	[PLANNED]	Each LLM call is stateless. Target: edge N context feeds edge N+1.
Algorithmic zoom	F_D Route	[FUTURE]	Uses named profiles, not continuous zoom. See Appendix C.
Recursive spawn	Engine	[FUTURE]	Spawn is manual via /gen-spawn. See Appendix D.

Gap	Category	Status	Notes
F_P modules (classify, route, sense)	F_P	[STUB]	fp_evaluate.py exists for evaluate only
F_H modules (all)	F_H	[STUB]	Interactive prompts — future work
Automatic η dispatch	η	[STUB]	Escalation signals recorded but not auto-dispatched
CLI entry point	Infra	[STUB]	No python -m genisis yet
Feature constraint merging	Config	[STUB]	feature.threshold_overrides + acceptance_criteria not yet composed
Actor model / inbox / serialiser	Infra	[FUTURE]	ADR-013 — defer until multi-agent needed
MCP sensory service	Infra	[FUTURE]	ADR-015 — defer until long-running monitors needed
Observer agents (dev, CI/CD, ops)	Infra	[FUTURE]	Design §1.11 — defer until hooks pipeline built

Part II: Roadmap

Everything in Part II is aspirational. These are target architectures, not current state. Appendices are ordered by dependency: A → B → C → D.

Appendix A: F_P Construct + Batched Evaluate [PLANNED – next implementation target]

The Critical Unlock

The engine today is an **evaluator** but not a **builder**. It can check whether assets meet criteria, but it cannot generate the assets. This appendix describes Strategy C (Hybrid Engine) — the minimum change that makes the engine a builder.

A.1: fp_construct.py — New Module

```
# Target interface (not yet implemented)
def run_construct(
    edge: str,
    asset_content: str,
    context: str,
    edge_config: dict,
    model: str = "sonnet",
    timeout: int = 120,
) -> str:
    """Call claude -p to construct/modify an asset for the given edge.

    Returns the constructed artifact as a string.
    """

```

One `claude -p` call per edge that: 1. Receives: current asset + accumulated context + edge checklist criteria 2. Returns: structured JSON with constructed artifact + all agent evaluations 3. Engine validates output deterministically (parsing, format, traceability)

A.2: Batched Construct + Evaluate

Instead of 33 separate `claude -p` calls for agent checks, **one call per edge** does both:

```
claude -p (per edge)

Input:
- Current asset
- Context from previous edges
- All agent criteria for this edge

Output (structured JSON):
- Constructed artifact
- Agent evaluation results (per check)
```

Projected cost: 4 LLM calls per iteration (vs 37 current, vs 1 E2E session).

A.3: Context Accumulation

Constructed artifacts from edge N become context for edge N+1:

Edge 1: intent→requirements	→ requirements.md	(context for Edge 2)
Edge 2: requirements→design	→ design.md + ADRs	(context for Edge 3)
Edge 3: design→code	→ source code	(context for Edge 4)
Edge 4: code↔unit_tests	→ test files	(final convergence)

Prerequisites

- fp_construct.py module with run_construct() function
 - JSON schema for combined construct+evaluate response
 - Engine loop modification: call construct before evaluate
 - Context threading between run_edge() calls
-

Appendix B: RLM Execution Model [FUTURE – target architecture]

The RLM Pattern

The Recursive Language Model (Zhang & Khattab, MIT 2025) solves “context rot” in LLMs by: 1. **Context as variable** — not in the prompt, explored programmatically 2. **REPL loop** — LLM writes code, code executes safely, output fed back 3. **Recursive decomposition** — sub-problems spawn child RLM instances 4. **Convergence via FINAL()** — explicit termination when answer is found

Where RLM Maps Today

The RLM analogy is **accurate for Strategy A (E2E agent)**. One LLM session constructs and evaluates in a continuous context window — this IS a REPL loop.

The RLM analogy is **not accurate for Strategy B (engine)** today. The engine has no construct step, makes 37 cold-start LLM calls with no shared context, and cannot converge independently. It is an evaluation orchestrator, not a REPL loop.

After Appendix A is implemented, the RLM analogy becomes accurate for Strategy C (Hybrid) as well — one LLM call per edge that constructs + evaluates, with context accumulation between edges.

RLM ↔ iterate() Mapping

RLM Concept	iterate() Mapping	Strategy A	Strategy C
context variable	Project assets (code, specs, tests)	In LLM context	Passed per-edge
LLM writes code	F_P construct — agent generates/modifies assets	In-session	Per-edge call
REPL executes code	F_D evaluate — deterministic checks run against assets	In-session	Subprocess

RLM Concept	iterate() Mapping	Strategy A	Strategy C
REPL output fed back	Evaluation results (delta, escalations) inform next iteration	In-session	Engine loop
FINAL("answer")	Convergence — delta = 0, all required checks pass	In-session	Engine detects
recursive_llm()	Spawn — child vector at deeper scope	Manual today	See Appendix D
max_iterations	Iteration budget — stuck detection, time-box	Profile config	max_iterations_per_edge

Three-Layer Cost Model

RLM uses a two-model strategy (expensive + cheap). The functor framework generalises this to three layers:

Layer	RLM Equivalent	Cost	When Used
F_D (deterministic)	REPL execution	Free	Every iteration
F_P (agent)	LLM call	\$\$\$	Construct + judgment tasks
F_H (human)	N/A in RLM	Blocking	When F_P fails (η escalation)

The three-layer model means **most evaluation is free** (F_D runs tests, lints, schemas). LLM calls happen for construction and semantic judgment (F_P). Human reviews happen when the agent can't resolve (F_H).

Prerequisites

- Appendix A complete (engine can construct)
 - At least one edge functional in hybrid mode (Strategy C)
-

Appendix C: Algorithmic Zoom [FUTURE – requires Appendix A]

Vision

Instead of selecting named profiles, use a **continuous zoom parameter** [0.0, 1.0] that determines graph subset, evaluator density, iteration budget, and recursion depth. Named profiles become discrete anchor points on this continuous space.

Zoom as Continuous Parameterisation

```
zoom:  
  graph_subset: [0.0, 1.0]      # 0.0 = minimal edges, 1.0 = all edges  
  evaluator_density: [0.0, 1.0] # 0.0 = agent-only, 1.0 = all types  
  context_density: [0.0, 1.0]   # 0.0 = minimal context, 1.0 = full  
    context  
  iteration_budget: [1, ∞]       # max iterations per edge  
  recursion_depth: [0, 5]        # max spawn depth
```

Named Profiles as Anchor Points

Profiles stay as named presets — zoom adds the ability to select intermediate values:

Profile	Equivalent Zoom	Characteristic
minimal	≈ 0.1	2 edges, 10-min time-box
hotfix	≈ 0.15	3 edges, 4-hour time-box, F_D route
spike	≈ 0.2	3 edges, 1-week time-box, F_P evaluate
poc	≈ 0.3	3 edges, 3-week time-box, F_P evaluate
standard	≈ 0.5	4-7 edges, per-sprint
full	≈ 1.0	All edges, unlimited

Zoom Selection

Zoom is driven by **existing data** — not arbitrary thresholds:

- `events.jsonl` — iteration history, delta trends
- `fd_sense` monitors — stall detection, freshness, coverage
- Feature trajectory — which edges converged, which are stuck
- Project constraints — time-box, risk level

Why This Requires Appendix A

Zoom parameterises the encoding (which units are F_D vs F_P). This only matters when the engine **does something different** based on encoding — specifically, when it can construct (F_P) at some zoom levels and skip construction (F_D template) at others. Without construct, zoom only affects which edges are traversed, which profiles already handle.

Prerequisites

- Appendix A complete (engine can construct)
 - At least one edge functional in hybrid mode
 - RouteResult extended with `zoom: float` field
 - `compute_encoding(zoom)` function replacing profile lookup
-

Appendix D: Recursive Spawn [FUTURE – requires Appendix A + C]

Vision

When the engine detects a stall (same delta for N iterations), it spawns a child vector at reduced scope to investigate the sub-problem. The child's findings fold back as context for the parent.

What Exists Today

- `/gen-spawn` creates child vector (manual command)
- Child vector has its own feature trajectory
- Fold-back via events: child emits `spawn_folded_back`, parent reads `events.jsonl`
- No engine code path for automatic spawning

Target: Automatic Recursive Spawn

```
# Target interface (not yet implemented)
def spawn(
    sub_feature_id: str,
    sub_context: str,
    parent_config: EngineConfig,
    current_depth: int,
) -> list[IterationRecord]:
    """Spawn child vector at depth+1 with reduced scope."""
    child_zoom = base_zoom * (0.7 ** (current_depth + 1))
    child_config = derive_config(parent_config, child_zoom)
    return run(sub_feature_id, child_config, sub_context,
              depth=current_depth + 1)
```

Zoom Decay at Depth

Each recursion level reduces effective scope — deeper investigations are inherently lighter:

Depth	Base Zoom 0.7	Effective Zoom	Approximate Profile
0	0.70	0.70	standard+
1	0.70	0.49	~standard
2	0.70	0.34	~poc
3	0.70	0.24	~spike

Fold-Back

Child's constructed artifacts become parent context: 1. Child runs at reduced scope, investigates sub-problem 2. Child converges, emits `spawn_folded_back` event with findings 3. Parent reads child's artifacts, adds to context 4. Parent resumes iteration with enriched context

Prerequisites

- Appendix A complete (engine can construct — required for child to produce anything)
 - Appendix C complete (zoom parameterisation — required for zoom decay at depth)
 - Engine `run()` accepting depth parameter
 - Max depth guard to prevent infinite recursion
-

Appendix E: Boundary Use Cases

These scenarios stress-test the design at the boundaries where technical choices become visible to the user. Each shows: what the user does, what they see, and how each phase changes the experience.

UC-1: New Developer, First Feature (Happy Path)

User does: `genesis start` on a fresh project.

Today (Strategy A — E2E agent):

State: UNINITIALISED

> Project name? `my-auth-service` (detected)
> Language? `Python 3.12` (detected)
> Intent? "Build OAuth2 authentication service"

State: IN_PROGRESS

Feature: REQ-F-AUTH-001 on intent→requirements (iteration 1)
[Agent generates requirements, evaluates, emits event]
[Routes to requirements→design, generates ADRs, evaluates, emits]
[Routes to design→code, generates code with REQ tags, evaluates, emits]
[Routes to code↔unit_tests, generates tests, runs pytest, iterates until green]

State: ALL_CONVERGED ✓

Time: ~10 min. Cost: ~\$3. Works today.

Today (Strategy B — Engine): Cannot do this. Engine has no construct step. All agent checks SKIP or ERROR. The engine declares convergence on empty assets.

After Appendix A (Strategy C — Hybrid):

Edge 1/4: intent→requirements [claude -p: 18s] ✓ delta=0
Edge 2/4: requirements→design [claude -p: 22s] ✓ delta=0
Edge 3/4: design→code [claude -p: 25s] ✓ delta=0
Edge 4/4: code↔unit_tests [claude -p: 15s] → δ=2 → [fix: 12s] → δ=0 ✓
Total: 92s, 5 LLM calls, \$0.60

Impact: Appendix A is required for the engine to be useful for new projects.

UC-2: Hotfix at 2am (Time Pressure)

User does: genesis start --feature REQ-F-HOTFIX-042 --profile hotfix

Today: Profile selects 3 edges, F_D routing, 4-hour time-box. Works.

After Appendix C (Zoom): User could say --zoom 0.15 instead of --profile hotfix. Mid-flight, if the fix is more complex, increase zoom to 0.5 without creating a new feature.

Impact: Zoom adds mid-flight flexibility. Profiles already work. Zoom is a convenience.

UC-3: Stuck Feature — Delta Unchanged 4 Iterations (Escalation)

User does: Feature won't converge after repeated iterations.

Today (E2E agent): Agent tries fixes, fails. After 4 iterations with same delta, reports STUCK. User manually spawns discovery vector via /gen-spawn.

Today (Engine): Engine detects stall via sense_feature_stall(), emits η escalation signals. But engine can't construct a fix and can't spawn automatically.

After Appendix A + D: Engine constructs fix attempts. After 3 identical deltas, spawns child at reduced scope automatically. Child investigates, folds back context. Parent resumes.

Impact: Appendix A (construct) enables auto-fix. Appendix D (spawn) enables recursive decomposition.

UC-4: Resume After Context Loss (Monday Morning)

User does: genesis status then genesis start.

Today: Works identically for both strategies. Status reads events.jsonl, reconstructs trajectory. Start detects IN_PROGRESS, selects closest-to-complete feature.

Impact: Minimal — resume doesn't change with any appendix.

UC-5: Autopilot Full Build (Unattended)

User does: genesis start --auto --feature REQ-F-DATAMAP-001 and walks away.

Today (E2E agent): One claude -p session, ~10 min, ~\$3-5. Works.

Today (Engine): Cannot do this (no construct).

After Appendix A (Hybrid): 4 LLM calls, ~92s, ~\$0.60. Fully testable, observable.

Impact: Appendix A makes autopilot viable in the engine. Strongest argument for prioritising it.

UC-6: Spike That Needs to Become a Feature (Scope Change)

User does: Creates a spike, investigates, realises it should be a full feature.

Today: Must create a new feature vector, manually copy findings, re-traverse with standard profile.

After Appendix C (Zoom): Change zoom from 0.2 to 0.5 — new edges activate, existing converged edges preserved.

After Appendix D (Spawn): Spike folds back to parent automatically.

Impact: Appendix C enables smooth scope changes. Appendix D enables automatic knowledge transfer.

UC-7: Multi-Feature Workspace (Two Features, Different Profiles)

User does: Two features active simultaneously — one standard, one spike.

Today: Works. Each feature has its own trajectory and profile. genesis start selects by priority (closest-to-complete).

Impact: Minimal change — multi-feature already works.

UC-8: Human Review Gate (Design Approval)

User does: Feature reaches a human gate (requirements→design requires approval).

Today (E2E agent): Agent constructs design candidate, presents for review with checklist.

Today (Engine): Engine can't construct the candidate — nothing to review.

After Appendix A: Engine constructs the design, then presents it for human review with full checklist status.

Impact: Appendix A is required for the engine to support human review gates meaningfully.

Summary: Which Appendices Matter for Which User Journeys

Use Case	Today	Appendix A	Appendix C	Appendix D
UC-1 New feature	E2E only	Engine viable	—	—
UC-2 Hotfix	Works	—	Mid-flight flex	—
UC-3 Stuck	Manual	Auto-fix	—	Auto-spawn
UC-4 Resume	Works	—	—	—
UC-5 Autopilot	E2E only	Engine viable, cheaper	—	—
UC-6 Scope change	Manual	—	Smooth transition	Auto fold-back
UC-7 Multi-feature	Works	—	—	—
UC-8 Human review	E2E only	Engine can construct	—	—

Appendix A (F_P construct) is the **critical unlock** — it makes the engine a viable builder. **Appendix C** (zoom) is a **convenience upgrade** — mid-flight scope changes. **Appendix D** (spawn) is **automation** — recursive decomposition.