

# AI SDLC – UAT Test Cases

**Version:** 1.0.0 **Date:** 2026-02-22 **Derived From:**

[AISDLC IMPLEMENTATION REQUIREMENTS.md](#) (v3.11.0),

[FEATURE VECTORS.md](#) (v1.8.0) **Method:** Asset Graph Model §9.2 — BDD  
(Given/When/Then), business language only

---

## Purpose & Approach

This document specifies **exhaustive functional use cases** for the AI SDLC Asset Graph Model. Every scenario is:

1. **Written in business language** (Given/When/Then) — no platform-specific syntax
2. **Traced to REQ keys** — every scenario carries `Validates: REQ-*` tags
3. **Event-verifiable** — expected outcomes verifiable via `events.jsonl`
4. **Fixture-based** — scenarios reference one of 5 base workspace fixtures
5. **Gap-annotated** — marked NEW (no existing coverage) or EXISTING (covered by which test)

This is a **spec-level artifact** — WHAT to test, not HOW. Platform-specific executable tests belong in `imp_<name>/tests/uat/`.

## Audience

Implementers building AI SDLC tooling for any platform (Claude, Gemini, Codex, or future implementations).

---

## Coverage Matrix

All 63 implementation requirements mapped to use case clusters and scenario counts.

REQ Key	Description	UC Cluster	Scenarios
REQ-INTENT-001	Intent Capture	UC-04	3
REQ-INTENT-002	Intent as Spec	UC-04	2

<b>REQ Key</b>	<b>Description</b>	<b>UC Cluster</b>	<b>Scenarios</b>
REQ-INTENT-003	Eco-Intent Generation	UC-06	2
REQ-INTENT-004	Spec Reproducibility	UC-03	3
REQ-GRAFH-001	Asset Type Registry	UC-01	3
REQ-GRAFH-002	Admissible Transitions	UC-01	4
REQ-GRAFH-003	Asset as Markov Object	UC-01	3
REQ-ITER-001	Universal Iteration Function	UC-01	4
REQ-ITER-002	Convergence and Promotion	UC-01	5
REQ-ITER-003	Functor Encoding Tracking	UC-01	5
REQ-EVAL-001	Three Evaluator Types	UC-02	4
REQ-EVAL-002	Evaluator Composition Per Edge	UC-02	4
REQ-EVAL-003	Human Accountability	UC-02	7
REQ-CTX-001	Context as Constraint Surface	UC-03	4
REQ-CTX-002	Context Hierarchy	UC-03	5
REQ-FEAT-001	Feature Vector Trajectories	UC-04	4
REQ-FEAT-002	Feature Dependencies	UC-04	4
REQ-FEAT-003	Task Planning as Trajectory Optimisation	UC-04	5
REQ-EDGE-001	TDD at Code/Tests Edges	UC-05	5
REQ-EDGE-002	BDD at Design/Test Edges	UC-05	4
REQ-EDGE-003	ADRs at Requirements/Design Edge	UC-05	4
REQ-EDGE-004	Code Tagging	UC-05	3
REQ-LIFE-001	CI/CD as Graph Edge	UC-06	3
REQ-LIFE-002	Telemetry and Homeostasis	UC-06	3
REQ-LIFE-003	Feedback Loop Closure	UC-06	3

REQ Key	Description	UC Cluster	Scenarios
REQ-LIFE-004	Feature Lineage in Telemetry	UC-06	2
REQ-LIFE-005	Intent Events as First-Class Objects	UC-06	3
REQ-LIFE-006	Signal Source Classification	UC-06	3
REQ-LIFE-007	Spec Change Events	UC-06	3
REQ-LIFE-008	Protocol Enforcement Hooks	UC-06	3
REQ-LIFE-009	Spec Review as Gradient Check	UC-06	2
REQ-LIFE-010	Development Observer Agent	UC-06	2
REQ-LIFE-011	CI/CD Observer Agent	UC-06	2
REQ-LIFE-012	Ops Observer Agent	UC-06	2
REQ-SENSE-001	Interoceptive Monitoring	UC-07	4
REQ-SENSE-002	Exteroceptive Monitoring	UC-07	3
REQ-SENSE-003	Affect Triage Pipeline	UC-07	4
REQ-SENSE-004	Sensory System Configuration	UC-07	2
REQ-SENSE-005	Review Boundary	UC-07	3
REQ-TOOL-001	Plugin Architecture	UC-08	2
REQ-TOOL-002	Developer Workspace	UC-08	3
REQ-TOOL-003	Workflow Commands	UC-08	3
REQ-TOOL-004	Release Management	UC-08	3
REQ-TOOL-005	Test Gap Analysis	UC-08	3
REQ-TOOL-006	Methodology Hooks	UC-08	2
REQ-TOOL-007	Project Scaffolding	UC-08	2
REQ-TOOL-008	Context Snapshot	UC-08	2
REQ-TOOL-009	Feature Views	UC-08	2
REQ-TOOL-010	Spec/Design Boundary Enforcement	UC-08	2
REQ-UX-001	State-Driven Routing	UC-09	8
REQ-UX-002	Progressive Disclosure	UC-09	4

REQ Key	Description	UC Cluster	Scenarios
REQ-UX-003	Project-Wide Observability	UC-09	4
REQ-UX-004	Automatic Feature/Edge Selection	UC-09	4
REQ-UX-005	Recovery and Self-Healing	UC-09	4
REQ-UX-006	Human Gate Awareness	UC-09	4
REQ-UX-007	Edge Zoom Management	UC-09	4
REQ-COORD-001	Agent Identity	UC-10	3
REQ-COORD-002	Feature Assignment via Events	UC-10	4
REQ-COORD-003	Work Isolation	UC-10	3
REQ-COORD-004	Markov-Aligned Parallelism	UC-10	3
REQ-COORD-005	Role-Based Evaluator Authority	UC-10	3
REQ-SUPV-001	IntentEngine Interface	UC-11	7
REQ-SUPV-002	Constraint Tolerances	UC-11	7

**Total: 64 REQ keys, 220 scenarios.**

---

## Test Fixture Strategy

All scenarios reference one of 5 base workspace fixtures. Each fixture represents a reproducible workspace state.

### Fixture 1: CLEAN

An empty directory with no `.ai-workspace/`. Represents a greenfield project.

**Contents:** Project source files only (e.g., `pyproject.toml`, `src/`). No methodology artifacts.

### Fixture 2: INITIALIZED

A freshly initialized workspace. `project_initialized` event emitted. No features created yet.

**Contents:** - .ai-workspace/ with graph topology, profiles, templates, project constraints - specification/INTENT.md with a concrete intent - events.jsonl with one project\_initialized event - No feature vectors in features/active/

## Fixture 3: IN\_PROGRESS

A workspace with 3 active features at various stages of progress.

**Contents:** - Everything in INITIALIZED - Feature A (REQ-F-ALPHA-001): converged through intent→requirements, iterating on requirements→design - Feature B (REQ-F-BETA-001): converged through design→code, iterating on code↔unit\_tests (iteration 3, delta=2) - Feature C (REQ-F-GAMMA-001): just started, at intent→requirements edge (iteration 1) - events.jsonl with ~15 events (edge\_started, iteration\_completed, edge\_converged) - Standard profile active

## Fixture 4: CONVERGED

A workspace where all features have completed all profile edges.

**Contents:** - Everything in INITIALIZED - 2 features, both fully converged (all edges in standard profile) - events.jsonl with complete trajectory (edge\_started → iteration\_completed → edge\_converged for each edge) - All assets tagged with REQ keys

## Fixture 5: STUCK

A workspace with features exhibiting stuck and blocked conditions.

**Contents:** - Everything in INITIALIZED - Feature X: same delta (3 failing checks) for 4 consecutive iterations on code↔unit\_tests - Feature Y: blocked — depends on spawn REQ-F-SPIKE-001 which has not converged - Feature Z: blocked — human review pending on requirements→design - events.jsonl with iteration events showing unchanging delta for Feature X

---

# UC-01: Asset Graph Engine

**Feature Vector:** REQ-F-ENGINE-001 **Satisfies:** REQ-GRAFH-001, REQ-GRAFH-002, REQ-GRAFH-003, REQ-ITER-001, REQ-ITER-002, REQ-ITER-003

## UC-01-01: Asset type registry contains all default types

**Validates:** REQ-GRAFH-001 | **Fixture:** INITIALIZED | **EXISTING** (TestGraphTopology)

Given an initialized workspace  
When I inspect the graph topology configuration  
Then I find exactly 10 asset types registered:  
Intent, Requirements, Design, Code, Unit Tests, Test Cases, UAT Tests, CI/CD, Running System, Telemetry

And each asset type has a schema definition  
And each asset type has Markov criteria for promotion

## **UC-01-02: Asset type registry is extensible**

Validates: REQ-GRAFH-001 | Fixture: INITIALIZED | NEW

Given an initialized workspace with the default 10 asset types  
When I add a custom asset type "security\_audit" to the graph topology  
with schema fields [audit\_id, findings, severity\_counts]  
and Markov criteria [has\_findings, human\_approved]  
Then the graph engine accepts the extended topology  
And I can define transitions involving "security\_audit"  
And existing transitions are unaffected

## **UC-01-03: Asset type interfaces are typed**

Validates: REQ-GRAFH-001 | Fixture: INITIALIZED | NEW

Given an initialized workspace  
When I create an asset of type "requirements"  
missing the mandatory field "acceptance\_criteria"  
Then the asset fails Markov criteria validation  
And the failure report identifies the missing field

## **UC-01-04: Only admissible transitions can be traversed**

Validates: REQ-GRAFH-002 | Fixture: IN\_PROGRESS | NEW

Given a workspace with Feature A at the "requirements" stage  
When I attempt to iterate on the edge "requirements→unit\_tests"  
which is not in the admissible transitions  
Then the system rejects the traversal  
And reports "No admissible transition from requirements to unit\_tests"

## **UC-01-05: Feedback loop edge makes the graph cyclic**

Validates: REQ-GRAFH-002 | Fixture: INITIALIZED | EXISTING  
(TestGraphTopology)

Given an initialized workspace  
When I inspect the graph topology  
Then I find a transition "Telemetry → Intent (Feedback Loop)"  
And the graph is marked as cyclic: true  
And the feedback loop edge has evaluators [agent, human]

## **UC-01-06: Transition registry is extensible**

Validates: REQ-GRAFH-002 | Fixture: INITIALIZED | NEW

Given an initialized workspace with 10 default transitions  
When I add a new transition "design→security\_audit"  
    with evaluators [agent, human]  
    and constructor: agent  
Then the graph topology contains 11 transitions  
And the new transition is traversable by the iterate function  
And all transitions are logged for audit

## **UC-01-07: Non-converged assets are candidates, not Markov objects**

Validates: REQ-GRAFH-003 | Fixture: IN\_PROGRESS | NEW

Given Feature B iterating on "code→unit\_tests" with delta=2  
When I inspect Feature B's code asset status  
Then the asset is marked as "candidate" (not "stable")  
And the asset cannot be promoted to the next graph node  
And the 2 failing evaluator checks are listed with remediation guidance

## **UC-01-08: Converged asset achieves Markov object status**

Validates: REQ-GRAFH-003 | Fixture: IN\_PROGRESS | NEW

Given Feature A whose "requirements" asset has delta=0  
    with all evaluators reporting convergence  
When I inspect Feature A's requirements asset status  
Then the asset is marked as "stable" (Markov object)  
And the asset is usable without its construction history  
And the asset's typed interface is valid

## **UC-01-09: Markov object is usable without construction history**

Validates: REQ-GRAFH-003 | Fixture: CONVERGED | NEW

Given a converged feature with a stable "design" asset  
When a new feature references Feature A's design as context  
Then the new feature can consume the design asset  
    without access to Feature A's iteration history  
And the design asset's typed interface is sufficient for downstream use

## **UC-01-10: Iterate function has universal signature**

Validates: REQ-ITER-001 | Fixture: IN\_PROGRESS | NEW

Given Feature A at the "requirements→design" edge  
When I invoke iterate(Asset<requirements>, Context[],  
Evaluators(requirements\_design))  
Then the function returns Asset<design.k+1> (next candidate)

And the same function signature works for all 10 edge types  
And the asset carries intent lineage and REQ keys

## **UC-01-11: Iterate behaviour is parameterised by edge config**

Validates: REQ-ITER-001 | Fixture: IN\_PROGRESS | NEW

Given two features at different edges:

  Feature A at "requirements→design" (evaluators: agent, human)

  Feature B at "code↔unit\_tests" (evaluators: agent, deterministic)

When I invoke iterate on each feature

Then Feature A's iteration uses ADR generation pattern

And Feature B's iteration uses TDD co-evolution pattern

And both use the same iterate function with different edge configs

## **UC-01-12: Iteration repeats until convergence**

Validates: REQ-ITER-001 | Fixture: IN\_PROGRESS | NEW

Given Feature B at "code↔unit\_tests" with delta=2

When I invoke iterate and the agent produces code that passes 1 more check

Then delta decreases to 1

And the feature remains at the same edge

And iterate can be invoked again to reduce delta further

And convergence is only declared when delta=0

## **UC-01-13: Constructor is edge-specific**

Validates: REQ-ITER-001 | Fixture: INITIALIZED | EXISTING (TestEdgeConfigs)

Given the graph topology with 10 transitions

When I inspect the constructor field for each transition

Then "intent→requirements" uses constructor: agent

And "code→cicd" uses constructor: deterministic

And "code↔unit\_tests" uses constructor: agent

And each constructor type is appropriate for its edge

## **UC-01-14: Convergence requires all evaluators to pass**

Validates: REQ-ITER-002 | Fixture: IN\_PROGRESS | NEW

Given Feature B at "code↔unit\_tests" with 3 evaluator checks

When 2 of 3 required checks pass but 1 fails

Then delta = 1

And the feature is NOT promoted to the next asset type

And the iteration report shows exactly which check failed

## **UC-01-15: Convergence threshold is configurable per evaluator**

Validates: REQ-ITER-002 | Fixture: IN\_PROGRESS | NEW

Given a feature with a coverage threshold of 80% (project default)  
And the feature vector overrides coverage threshold to 95%  
When the code achieves 85% coverage  
Then the project-level check would pass ( $85\% > 80\%$ )  
But the feature-level check fails ( $85\% < 95\%$ )  
And the effective threshold is 95% (feature override wins)

## **UC-01-16: Promotion creates next asset type version**

Validates: REQ-ITER-002 | Fixture: IN\_PROGRESS | NEW

Given Feature A at "requirements-design" with delta=0  
When all evaluators report convergence (including human approval)  
Then the requirements asset is promoted: ATn.m becomes ATn+1.0  
And a new design asset (version 0) is created  
And an "edge\_converged" event is emitted  
And the feature trajectory advances to the next edge

## **UC-01-17: Non-convergence after max iterations escalates to human**

Validates: REQ-ITER-002 | Fixture: STUCK | NEW

Given Feature X with delta=3 unchanged for 4 iterations  
And the iteration budget is 5 iterations before escalation  
When the 5th iteration also produces delta=3  
Then the system escalates to the human evaluator  
And the escalation includes: which checks are stuck, what has been tried  
And the human can: force-converge, spawn discovery, or continue iterating

## **UC-01-18: Extended convergence for discovery vectors (question answered)**

Validates: REQ-ITER-002 | Fixture: IN\_PROGRESS | NEW

Given a discovery vector investigating "Can technology X handle load Y?"  
And the vector has delta=1 (one non-required check remains)  
When the human evaluator judges "Yes, the question has been answered"  
Then the vector converges with convergence\_type: "question\_answered"  
And the findings are packaged as fold-back payload  
And a "spawn\_folded\_back" event is emitted to the parent feature

## **UC-01-19: Profile encoding section defines functor categories**

Validates: REQ-ITER-003 | Fixture: INITIALIZED | EXISTING  
(TestFunctorEncoding)

Given an initialized workspace with the standard profile  
When I inspect the profile encoding section  
Then I find strategy: "balanced", mode: "interactive", valence:  
"medium"  
And 8 functional units with category mappings:  
evaluate: F\_D, construct: F\_P, classify: F\_D, route: F\_H,  
propose: F\_P, sense: F\_D, emit: F\_D, decide: F\_H

## **UC-01-20: Feature vectors carry functor encoding**

Validates: REQ-ITER-003 | Fixture: IN\_PROGRESS | NEW

Given Feature B iterating on "code↔unit\_tests"  
When I inspect the iteration event for Feature B  
Then the event contains an "encoding" section with:  
encoding\_source (profile name), mode, valence, and active\_units  
And the active\_units map all 8 functional units to categories

## **UC-01-21: Category-fixed units are enforced**

Validates: REQ-ITER-003 | Fixture: INITIALIZED | NEW

Given a custom profile that attempts to set emit: F\_P (overriding the fixed category)  
When the profile is validated  
Then the system rejects the override for "emit" (must be F\_D)  
And the system rejects any override for "decide" (must be F\_H)  
And other functional units accept category overrides

## **UC-01-22: Encoding escalation emits event**

Validates: REQ-ITER-003 | Fixture: IN\_PROGRESS | NEW

Given Feature B on "code↔unit\_tests" with evaluate: F\_D (deterministic)  
When the deterministic evaluator cannot classify an ambiguous test failure  
And the system escalates evaluate from F\_D to F\_P (probabilistic agent)  
Then an "encoding\_escalated" event is emitted with:  
functional\_unit: "evaluate", from\_category: "F\_D", to\_category:  
"F\_P"  
And the trigger reason is recorded

## **UC-01-23: Escalation trajectory is recorded per feature**

Validates: REQ-ITER-003 | Fixture: IN\_PROGRESS | NEW

Given Feature B with 2 encoding escalations during its trajectory  
When I inspect the feature vector's trajectory  
Then each trajectory entry supports an "escalations" array  
And the escalations record: which unit, from/to category, iteration number  
And the escalation history is available for trajectory analysis

---

## **UC-02: Evaluator Framework**

**Feature Vector:** REQ-F-EVAL-001 **Satisfies:** REQ-EVAL-001, REQ-EVAL-002, REQ-EVAL-003

### **UC-02-01: Human evaluator captures judgment**

Validates: REQ-EVAL-001 | Fixture: IN\_PROGRESS | NEW

Given Feature A at "requirements→design" edge which requires human evaluation  
When the iterate function presents the design candidate for review  
Then the human can: approve, reject, or provide refinement guidance  
And the decision is recorded in the iteration history  
And a "review\_completed" event is emitted with the human's feedback

### **UC-02-02: Agent evaluator computes delta via LLM**

Validates: REQ-EVAL-001 | Fixture: IN\_PROGRESS | NEW

Given Feature B at "code↔unit\_tests" with agent evaluator configured  
When the agent evaluator assesses the code candidate  
Then it computes delta between current state and target state  
And provides specific, actionable feedback for the next iteration  
And does NOT declare convergence if known gaps exist

### **UC-02-03: Deterministic evaluator produces binary pass/fail**

Validates: REQ-EVAL-001 | Fixture: IN\_PROGRESS | NEW

Given Feature B at "code↔unit\_tests" with deterministic evaluator "tests\_pass"  
When the test runner executes and 3 of 5 tests fail  
Then the deterministic evaluator reports: FAIL  
And identifies the 3 failing tests with their error messages  
And provides remediation guidance for each failure

## **UC-02-04: Evaluators declare processing phase**

Validates: REQ-EVAL-001 | Fixture: INITIALIZED | **EXISTING** (TestProcessingPhases)

Given the evaluator configuration for all edge types  
When I inspect the processing\_phase field for each evaluator type  
Then Deterministic Tests are classified as "reflex" (autonomic)  
And Human and Agent deliberative are classified as "conscious"  
(deliberative)  
And Affect is a valence vector emitted by ANY evaluator on its gap  
finding (not an evaluator type assignment)

## **UC-02-05: Evaluator composition is configurable per edge**

Validates: REQ-EVAL-002 | Fixture: INITIALIZED | **EXISTING** (TestEdgeConfigs)

Given the graph topology with 10 transitions  
When I inspect the evaluator composition for each edge  
Then "intent→requirements" uses [agent, human]  
And "code↔unit\_tests" uses [agent, deterministic]  
And "code→cicd" uses [deterministic] only  
And each composition is read from config, not hard-coded

## **UC-02-06: Evaluator composition overridable at project level**

Validates: REQ-EVAL-002 | Fixture: INITIALIZED | **NEW**

Given the default "design→code" edge with evaluators [agent, deterministic]  
When the project overrides this edge to [agent, deterministic, human]  
Then the iterate function uses 3 evaluator types for this edge  
And convergence requires all 3 to pass  
And the override is stored in project configuration, not the graph topology

## **UC-02-07: Profile overrides evaluator composition**

Validates: REQ-EVAL-002 | Fixture: INITIALIZED | **EXISTING** (TestProfiles)

Given the standard profile with evaluator overrides:  
  "intent→requirements": [agent, human]  
  "requirements→design": [agent, human]  
When a feature using the standard profile reaches  
"intent→requirements"  
Then the iterate function uses [agent, human] evaluators (not the  
graph default)  
And the override source is the profile, not the edge default

## **UC-02-08: Default compositions match spec table**

Validates: REQ-EVAL-002 | Fixture: INITIALIZED | **EXISTING**  
(TestEvaluatorDefaults)

Given the graph topology configuration  
When I compare evaluator compositions against the Asset Graph Model §4.2 table  
Then every edge matches the specified default composition  
And the composition is complete (no edges missing evaluator assignments)

## **UC-02-09: AI suggestions require human review at human-configured edges**

Validates: REQ-EVAL-003 | Fixture: IN\_PROGRESS | **NEW**

Given Feature A at "requirements-design" with evaluators [agent, human]  
When the agent evaluator reports delta=0 (all agent checks pass)  
Then the feature does NOT auto-converge  
And the system presents the candidate for human review  
And convergence is blocked until the human explicitly approves

## **UC-02-10: Human can override any agent decision**

Validates: REQ-EVAL-003 | Fixture: IN\_PROGRESS | **NEW**

Given Feature A at "requirements-design" where the agent recommends convergence  
When the human reviews and disagrees with the agent's assessment  
And the human provides refinement feedback  
Then the human's rejection overrides the agent's approval  
And iteration continues with the human's feedback as guidance  
And the rejection is attributed to the human, not the AI

## **UC-02-11: Decisions attributed to humans, not AI**

Validates: REQ-EVAL-003 | Fixture: IN\_PROGRESS | **NEW**

Given a "review\_completed" event emitted after human evaluation  
When I inspect the event  
Then the decision is attributed to the human reviewer  
And the event records: who decided, what they decided, and their feedback  
And the AI's prior recommendation is logged but not authoritative

## **UC-02-12: Human evaluator override always available**

Validates: REQ-EVAL-003 | Fixture: IN\_PROGRESS | **NEW**

Given Feature B at "code->unit\_tests" with evaluators [agent, deterministic]  
where human evaluator is NOT in the default composition  
When the user requests a human review on this edge  
Then the system allows the human to review and provide feedback  
And the human's feedback is recorded in the iteration history  
And the human override capability is always available regardless of edge config

### **UC-02-13: Spec mutations always require human approval**

Validates: REQ-EVAL-003 | Fixture: IN\_PROGRESS | NEW

Given an agent evaluator that detects a spec gap during iteration  
When the agent proposes a specification modification  
Then the modification is presented as a draft to the human  
And the specification is NOT modified until the human approves  
And the approval is recorded as a "spec\_modified" event with human attribution

### **UC-02-14: Human approval required at human-required edges before convergence**

Validates: REQ-EVAL-003 | Fixture: IN\_PROGRESS | NEW

Given Feature A at "requirements->design" where human\_required: true  
When all agent and deterministic evaluators report delta=0  
Then convergence status shows "PENDING\_HUMAN\_REVIEW"  
And the edge does NOT converge until human explicitly approves  
And auto-mode pauses at this gate

### **UC-02-15: Human review records feedback in iteration history**

Validates: REQ-EVAL-003 | Fixture: IN\_PROGRESS | NEW

Given Feature A at "requirements->design" pending human review  
When the human provides refinement guidance: "Add error handling for edge case X"  
Then the feedback is stored in the feature's iteration history  
And the next iteration uses this feedback as construction guidance  
And the feedback is traceable in the event log

---

## **UC-03: Context Management**

**Feature Vector:** REQ-F-CTX-001 **Satisfies:** REQ-CTX-001, REQ-CTX-002, REQ-INTENT-004

## **UC-03-01: Context types include required categories**

Validates: REQ-CTX-001 | Fixture: INITIALIZED | **EXISTING** (TestFixtureSources)

Given an initialized workspace

When I inspect the context directory structure

Then I find directories for: ADRs, data models, templates, policy, standards

And the context system accepts open-ended context types

And context elements are version-controlled in the workspace

## **UC-03-02: Context narrows admissible constructions**

Validates: REQ-CTX-001 | Fixture: IN\_PROGRESS | **NEW**

Given Feature A at "requirements->design" edge

And Context[] contains an ADR mandating "PostgreSQL for persistence"

When the iterate function generates a design candidate

Then the design references PostgreSQL (not an arbitrary database)

And the ADR constraint narrowed the space of admissible designs

And the context element is recorded in the iteration's context hash

## **UC-03-03: Context subset configurable per edge**

Validates: REQ-CTX-001 | Fixture: INITIALIZED | **NEW**

Given a workspace with 5 context elements:

ADR-001, DataModel-User, Template-API, Policy-Security, Standard-Naming

When iterating on "design->code" edge

Then the iterate function loads only the context elements relevant to this edge

And the relevance is configurable (not all context applies to all edges)

And the loaded context is recorded in the context hash

## **UC-03-04: Context version control**

Validates: REQ-CTX-001 | Fixture: INITIALIZED | **NEW**

Given a workspace with ADR-001 v1 in Context[]

When ADR-001 is updated to v2 (changed decision)

Then the context system preserves v1 (immutable)

And subsequent iterations use v2

And the context hash changes to reflect the new version

And the version history is traceable

## **UC-03-05: Context hierarchy — global to project override**

Validates: REQ-CTX-002 | Fixture: INITIALIZED | **NEW**

Given a context hierarchy:

```
Global: {coverage_threshold: 80%, style: "PEP8"}  
Organisation: {coverage_threshold: 90%}  
Project: {style: "Google Python Style"}  
When the iterate function resolves the effective context  
Then coverage_threshold = 90% (organisation overrides global)  
And style = "Google Python Style" (project overrides global)  
And unoverridden values inherit from parent levels
```

## UC-03-06: Context deep merge for objects

Validates: REQ-CTX-002 | Fixture: INITIALIZED | NEW

```
Given a global context with deployment: {cloud: "AWS", region: "us-east-1"}  
And a project context with deployment: {region: "eu-west-1"}  
When the iterate function resolves the effective context  
Then deployment = {cloud: "AWS", region: "eu-west-1"} (deep merge)  
And the cloud field is inherited, the region field is overridden
```

## UC-03-07: Later context overrides earlier

Validates: REQ-CTX-002 | Fixture: INITIALIZED | NEW

```
Given context loaded in order: global → organisation → team → project  
When the same key appears at multiple levels  
Then the project value wins (last loaded)  
And the override chain is traceable for audit
```

## UC-03-08: Customisation without forking

Validates: REQ-CTX-002 | Fixture: INITIALIZED | NEW

```
Given a team-level context with standard evaluator thresholds  
When a project needs stricter coverage (95% instead of 80%)  
Then the project overrides only the coverage threshold  
And all other team-level context is inherited unchanged  
And the team context is not modified (no fork)
```

## UC-03-09: Context hierarchy supports four levels

Validates: REQ-CTX-002 | Fixture: INITIALIZED | NEW

```
Given context defined at all four levels: global, organisation, team, project  
When I request the effective context for a feature  
Then all four levels are composed in order  
And conflicts resolve by the "later overrides earlier" rule  
And the composition is deterministic (same inputs = same output)
```

## **UC-03-10: Spec canonical serialisation is deterministic**

Validates: REQ-INTENT-004 | Fixture: IN\_PROGRESS | NEW

Given the same Intent + Context[] provided twice  
When the spec is canonically serialised each time  
Then both serialisations produce identical byte sequences  
And the content-addressable hash is the same  
And independent tools would compute the same hash

## **UC-03-11: Spec hash recorded at each iteration**

Validates: REQ-INTENT-004 | Fixture: IN\_PROGRESS | NEW

Given Feature B on "code↔unit\_tests" at iteration 3  
When I inspect the iteration\_completed event  
Then the event contains a "context\_hash" field with "sha256:..." format  
And the hash changes if and only if the effective context changes

## **UC-03-12: Spec versions are immutable**

Validates: REQ-INTENT-004 | Fixture: IN\_PROGRESS | NEW

Given a spec version V1 with hash H1  
When a requirement is modified (creating spec version V2 with hash H2)  
Then V1 remains accessible with its original hash H1  
And V2 is a new version, not a mutation of V1  
And the evolution from V1 to V2 is traceable via "spec\_modified" events

---

## **UC-04: Feature Vector Traceability**

**Feature Vector:** REQ-F-TRACE-001 **Satisfies:** REQ-INTENT-001, REQ-INTENT-002, REQ-FEAT-001, REQ-FEAT-002, REQ-FEAT-003

### **UC-04-01: Intent captured in structured format**

Validates: REQ-INTENT-001 | Fixture: CLEAN | NEW

Given a clean project directory  
When the user provides an intent: "Build a user authentication system"  
Then an intent is created with:  
    unique identifier (INT-001)  
    description: "Build a user authentication system"  
    source: "human"  
    timestamp: current time  
    priority: assigned by user

And the intent is persisted in specification/INTENT.md  
And the intent is version-controlled

## UC-04-02: Intent from runtime feedback

Validates: REQ-INTENT-001 | Fixture: CONVERGED | NEW

Given a converged project with running telemetry  
When the ops observer detects a latency deviation  
Then an intent is generated with:  
  identifier: INT-002  
  source: "runtime\_feedback"  
  description including the deviation details  
And the intent enters the asset graph as a new feature vector candidate

## UC-04-03: Intent from ecosystem changes

Validates: REQ-INTENT-001 | Fixture: CONVERGED | NEW

Given a converged project using library X v2.0  
When the exteroceptive monitor detects library X v2.0 has a critical CVE  
Then an intent is generated with:  
  identifier: INT-EC0-001  
  source: "ecosystem"  
  description: "Critical CVE in library X v2.0"  
And the intent enters the graph for human triage

## UC-04-04: Intent composes with Context to form Spec

Validates: REQ-INTENT-002 | Fixture: IN\_PROGRESS | NEW

Given INT-001 ("Build a user authentication system")  
And Context[] containing: ADR-001 (OAuth2), DataModel-User, Policy-Security  
When the intent composes with context to form the spec  
Then the spec is the fitness landscape against which evaluators measure convergence  
And the spec includes both the intent and all context constraints  
And the spec evolves as intent lineage accumulates through traversal

## UC-04-05: Spec is the fitness landscape for evaluators

Validates: REQ-INTENT-002 | Fixture: IN\_PROGRESS | NEW

Given a spec formed from INT-001 + Context[]  
When evaluators assess a candidate design  
Then the evaluators measure convergence against the spec  
And delta is the distance between the candidate and the spec's

constraints  
And the spec bounds what counts as "correct"

## **UC-04-06: REQ key format is enforced**

Validates: REQ-FEAT-001 | Fixture: INITIALIZED | EXISTING  
(TestReqKeyCoverage)

Given the requirement key format: REQ-{TYPE}-{DOMAIN}-{SEQ}  
When I create requirements with valid keys (REQ-F-AUTH-001, REQ-NFR-PERF-001)  
Then the keys are accepted  
And when I create a requirement with invalid key format ("AUTH-001")  
Then the key is rejected with a format error  
And REQ key types include: F, NFR, DATA, BR

## **UC-04-07: REQ keys propagate through all assets on trajectory**

Validates: REQ-FEAT-001 | Fixture: CONVERGED | EXISTING  
(TestEndToEndTraceability)

Given a converged feature REQ-F-AUTH-001  
When I trace the REQ key through all artifacts  
Then the spec references REQ-F-AUTH-001  
And the design references REQ-F-AUTH-001  
And the code contains "Implements: REQ-F-AUTH-001"  
And the tests contain "Validates: REQ-F-AUTH-001"  
And bidirectional navigation works: forward (intent→runtime) and backward (runtime→intent)

## **UC-04-08: REQ key is immutable identifier**

Validates: REQ-FEAT-001 | Fixture: IN\_PROGRESS | NEW

Given REQ-F-AUTH-001 assigned to a requirement  
When the requirement text is modified (creating version REQ-F-AUTH-001.1.0.0)  
Then the key REQ-F-AUTH-001 remains the same (immutable trajectory identifier)  
And the version suffix tracks how the requirement statement evolves  
And all artifacts continue referencing REQ-F-AUTH-001 (not the versioned form)

## **UC-04-09: Bidirectional navigation works**

Validates: REQ-FEAT-001 | Fixture: CONVERGED | NEW

Given a converged feature REQ-F-AUTH-001 with artifacts at all stages  
When I navigate forward from intent  
Then I can trace: Intent → Requirements → Design → Code → Tests

And when I navigate backward from a test  
Then I can trace: Test → Code → Design → Requirements → Intent  
And the REQ key is the thread connecting all stages

## **UC-04-10: Cross-feature dependencies tracked**

Validates: REQ-FEAT-002 | Fixture: IN\_PROGRESS | EXISTING  
(TestFeatureVectorConsistency)

Given Feature A (REQ-F-AUTH-001) depends on Feature B (REQ-F-DB-001)  
at the design stage  
When I inspect Feature A's dependency graph  
Then the dependency is recorded: A.design depends on B.design  
And the dependency is between trajectories, not individual assets  
And the dependency graph shows the relationship

## **UC-04-11: Circular dependencies detected**

Validates: REQ-FEAT-002 | Fixture: IN\_PROGRESS | NEW

Given Feature A depends on Feature B at the code stage  
And Feature B depends on Feature A at the code stage  
When the dependency graph is validated  
Then the system detects the circular dependency  
And flags it with both feature IDs and the stage where the cycle occurs  
And the cycle blocks concurrent iteration on both features

## **UC-04-12: Dependency graph is visualisable**

Validates: REQ-FEAT-002 | Fixture: IN\_PROGRESS | NEW

Given 3 features with dependencies: A→B, A→C (no dependency between B and C)  
When I request the dependency graph  
Then the system produces a visual representation showing:  
A depends on B and C  
B and C are independent (parallelisable)  
And the visualisation identifies the critical path

## **UC-04-13: Feature views show per-REQ cross-artifact status**

Validates: REQ-FEAT-002 | Fixture: IN\_PROGRESS | NEW

Given a feature REQ-F-AUTH-001 with artifacts at spec, design, and code stages  
When I request the feature view for REQ-F-AUTH-001  
Then I see which stages have the REQ key tagged:  
Spec: tagged, Design: tagged, Code: tagged, Tests: not yet,  
Telemetry: not yet

And missing stages are flagged  
And the coverage summary shows "3 of 5 stages tagged"

### **UC-04-14: Tasks emerge from feature vector decomposition**

Validates: REQ-FEAT-003 | Fixture: INITIALIZED | NEW

Given an intent "Build a user authentication system"  
When the system decomposes the intent into feature vectors  
Then each feature vector maps to a trajectory through the asset graph  
And the features are: REQ-F-AUTH-001, REQ-F-SESSION-001, REQ-F-RBAC-001  
And each feature has a defined path through graph edges

### **UC-04-15: Task graph identifies parallelisation opportunities**

Validates: REQ-FEAT-003 | Fixture: IN\_PROGRESS | NEW

Given 3 features with dependency graph: A→B, A→C, B and C independent  
When the task graph is generated  
Then B and C are marked as parallelisable (can execute concurrently)  
And A must complete its dependency edge before B and C can proceed  
And the compressed task graph shows the optimal execution order

### **UC-04-16: Task graph batches parallel edges**

Validates: REQ-FEAT-003 | Fixture: IN\_PROGRESS | NEW

Given 3 independent features all at the "design→code" edge  
When the task graph is generated  
Then the 3 design→code iterations are batched as parallel work  
And dependencies at later edges are sequenced correctly  
And the batch execution reduces total development time

### **UC-04-17: Inter-vector dependencies identified at each node**

Validates: REQ-FEAT-003 | Fixture: IN\_PROGRESS | NEW

Given Feature A's code depends on Feature B's code (shared module)  
When the task graph maps trajectories through the asset graph  
Then the dependency is identified at the "code" node specifically  
And Feature A's code edge is sequenced after Feature B's code edge  
And their design edges may still execute in parallel (no dependency at design)

### **UC-04-18: Compression produces minimal task graph**

Validates: REQ-FEAT-003 | Fixture: IN\_PROGRESS | NEW

Given 5 features with various inter-dependencies  
When the task graph is compressed  
Then parallel-eligible edges are batched  
And sequential dependencies are ordered  
And the result is the minimum spanning task graph with dependency order  
And no unnecessary sequencing is introduced

---

## UC-05: Edge Parameterisations

**Feature Vector:** REQ-F-EDGE-001 **Satisfies:** REQ-EDGE-001, REQ-EDGE-002, REQ-EDGE-003, REQ-EDGE-004

### UC-05-01: TDD RED phase — write failing test first

Validates: REQ-EDGE-001 | Fixture: IN\_PROGRESS | NEW

Given Feature B at "code↔unit\_tests" beginning a new TDD cycle  
When the iterate function starts the TDD co-evolution  
Then the first action is to write a failing test  
And the test references "Validates: REQ-F-BETA-001"  
And the deterministic evaluator confirms: test exists, test fails (RED state)

### UC-05-02: TDD GREEN phase — minimal code to pass

Validates: REQ-EDGE-001 | Fixture: IN\_PROGRESS | NEW

Given Feature B in RED state (failing test exists)  
When the iterate function writes code to pass the test  
Then the code is minimal (just enough to make the test pass)  
And the code references "Implements: REQ-F-BETA-001"  
And the deterministic evaluator confirms: test passes (GREEN state)

### UC-05-03: TDD REFACTOR phase — quality improvement

Validates: REQ-EDGE-001 | Fixture: IN\_PROGRESS | NEW

Given Feature B in GREEN state (all tests pass)  
When the iterate function refactors for quality  
Then the refactored code still passes all tests  
And all evaluators re-confirm convergence after refactoring  
And the refactoring does not change behaviour

### UC-05-04: TDD co-evolution oscillates between test and code

Validates: REQ-EDGE-001 | Fixture: IN\_PROGRESS | NEW

Given Feature B at "code↔unit\_tests" (bidirectional edge)  
When I inspect the iteration history  
Then iterations alternate between writing tests and writing code  
And this is a single edge with bidirectional construction  
And convergence requires both test and code assets to be stable

## **UC-05-05: TDD coverage threshold configurable**

Validates: REQ-EDGE-001 | Fixture: IN\_PROGRESS | NEW

Given the project default coverage threshold of 80%  
When Feature B's vector overrides the threshold to 95%  
Then the TDD edge uses 95% as the convergence criterion  
And coverage below 95% reports delta > 0 even if code compiles and tests pass  
And the threshold source (feature override) is recorded in the checklist

## **UC-05-06: BDD scenarios in Gherkin format**

Validates: REQ-EDGE-002 | Fixture: IN\_PROGRESS | NEW

Given Feature A at the "design→uat\_tests" edge  
When the iterate function generates UAT test scenarios  
Then each scenario is in Given/When/Then format  
And each scenario is tagged with "Validates: REQ-F-ALPHA-001"  
And scenarios use business language only (no technical jargon)  
And scenarios are executable, not just documentation

## **UC-05-07: Every REQ key has at least 1 BDD scenario**

Validates: REQ-EDGE-002 | Fixture: IN\_PROGRESS | NEW

Given Feature A with 3 REQ keys: REQ-F-ALPHA-001, REQ-F-ALPHA-002, REQ-F-ALPHA-003  
When the BDD edge evaluator checks scenario coverage  
Then each REQ key has at least 1 BDD scenario  
And missing scenarios are reported as delta > 0

## **UC-05-08: System test BDD vs UAT BDD differentiation**

Validates: REQ-EDGE-002 | Fixture: IN\_PROGRESS | NEW

Given Feature A at "design→test\_cases" (system test BDD)  
And Feature A at "design→uat\_tests" (UAT BDD)  
When I compare the generated scenarios  
Then system test BDD includes technical integration details  
And UAT BDD uses pure business language (no technical jargon)  
And both formats trace to the same REQ keys

## **UC-05-09: BDD human approval for business scenarios**

Validates: REQ-EDGE-002 | Fixture: IN\_PROGRESS | NEW

Given Feature A at "design->uat\_tests" with evaluators [agent, human]

When the agent generates BDD scenarios

Then the scenarios are presented to the human (business stakeholder)

And the human validates the business language accuracy

And convergence requires human approval (business language review)

## **UC-05-10: ADR generation at requirements to design edge**

Validates: REQ-EDGE-003 | Fixture: IN\_PROGRESS | NEW

Given Feature A at "requirements->design" edge

And mandatory constraint dimensions: ecosystem, deployment, security, build

When the iterate function produces a design candidate

Then ADRs are generated for each mandatory constraint dimension

And each ADR documents: decision, context, consequences, alternatives considered

And each ADR references the requirement keys it addresses

And ADRs become part of Context[] for downstream edges

## **UC-05-11: ADRs acknowledge ecosystem constraints**

Validates: REQ-EDGE-003 | Fixture: IN\_PROGRESS | NEW

Given the ecosystem constraint: "Python 3.12 + Django 5.0"

When ADRs are generated at the requirements->design edge

Then at least one ADR acknowledges and resolves the ecosystem constraint

And the ADR documents why Django 5.0 was chosen over alternatives

And the ADR references the requirement keys it satisfies

## **UC-05-12: ADRs are versioned context artifacts**

Validates: REQ-EDGE-003 | Fixture: IN\_PROGRESS | NEW

Given ADR-001 generated during Feature A's requirements->design iteration

When Feature B reaches the same edge

Then Feature B can reference ADR-001 as part of its Context[]

And ADR-001 is version-controlled and immutable once accepted

## **UC-05-13: ADR referencing requirement keys**

Validates: REQ-EDGE-003 | Fixture: IN\_PROGRESS | NEW

Given ADR-001 documenting the database choice  
When I inspect ADR-001  
Then it references the requirement keys that drove the decision  
And each referenced REQ key is traceable back to the original intent

### **UC-05-14: Code tagged with Implements**

**Validates:** REQ-EDGE-004 | **Fixture:** IN\_PROGRESS | **EXISTING**  
(TestEndToEndTraceability)

Given Feature B at "design->code" producing code files  
When the iterate function generates code  
Then each code file contains "Implements: REQ-F-BETA-001" in language-appropriate comments  
And the tag format is the contract (comment syntax varies by language)

### **UC-05-15: Tests tagged with Validates**

**Validates:** REQ-EDGE-004 | **Fixture:** IN\_PROGRESS | **EXISTING**  
(TestEndToEndTraceability)

Given Feature B at "code↔unit\_tests" producing test files  
When the iterate function generates tests  
Then each test file contains "Validates: REQ-F-BETA-001"  
And the tag format matches the code tagging convention

### **UC-05-16: Commit messages include REQ keys**

**Validates:** REQ-EDGE-004 | **Fixture:** IN\_PROGRESS | **NEW**

Given Feature B completing a TDD cycle (RED→GREEN→REFACTOR→COMMIT)  
When the iterate function produces a commit  
Then the commit message includes the REQ key: "Implements: REQ-F-BETA-001"  
And the REQ key in the commit message is validatable by tooling

---

## **UC-06: Full Lifecycle Closure**

**Feature Vector:** REQ-F-LIFE-001 **Satisfies:** REQ-LIFE-001 through REQ-LIFE-012, REQ-INTENT-003

### **UC-06-01: CI/CD is a first-class graph edge**

**Validates:** REQ-LIFE-001 | **Fixture:** INITIALIZED | **EXISTING**  
(TestGraphTopology)

Given the graph topology  
When I inspect transitions involving CI/CD

Then "Code → CI/CD" exists with evaluators: [deterministic],  
constructor: deterministic  
And "CI/CD → Running System" exists with evaluators: [deterministic],  
constructor: deterministic  
And both are iterative edges (failed deployment retries with evaluator  
feedback)

## **UC-06-02: CI/CD iteration retries on failure**

Validates: REQ-LIFE-001 | Fixture: IN\_PROGRESS | NEW

Given a feature at the "code→cicd" edge  
When the build fails (deterministic evaluator reports FAIL)  
Then delta > 0 and the iteration can be retried  
And the failure feedback guides the next iteration  
And the CI/CD edge behaves like any other iterate() call

## **UC-06-03: Release manifest lists feature vector IDs**

Validates: REQ-LIFE-001 | Fixture: CONVERGED | NEW

Given all features converged and a release is created  
When I inspect the release manifest  
Then the manifest lists all feature vector IDs (REQ keys) included in  
the release  
And a "release\_created" event is emitted with features\_included count  
and coverage\_pct

## **UC-06-04: Telemetry tagged with REQ keys**

Validates: REQ-LIFE-002 | Fixture: CONVERGED | NEW

Given a running system with telemetry configured  
When the system emits telemetry (logs, metrics, traces)  
Then telemetry records contain req="REQ-F-AUTH-001" as a structured  
field  
And the REQ keys in telemetry match those in code ("Implements: REQ-  
\*")

## **UC-06-05: Homeostasis check at runtime**

Validates: REQ-LIFE-002 | Fixture: CONVERGED | NEW

Given a running system with SLA bounds configured  
When the system's latency exceeds the p99 SLA threshold  
Then the deterministic evaluator (health probes) reports deviation  
And the deviation is tagged with the affected REQ keys  
And a drift detection signal is generated

## **UC-06-06: Runtime evaluators span all three types**

Validates: REQ-LIFE-002 | Fixture: CONVERGED | NEW

Given a running system with telemetry monitoring  
When evaluators are invoked at the telemetry stage  
Then Deterministic Tests check: alerting rules, SLA thresholds, health probes  
And Agent evaluators check: anomaly detection, trend analysis  
And Human evaluators respond to: incident triage, capacity planning

## **UC-06-07: Deviation generates new intent**

Validates: REQ-LIFE-003 | Fixture: CONVERGED | NEW

Given a running system that detects a latency deviation on REQ-F-AUTH-001  
When the deviation is confirmed by the ops observer  
Then a new intent INT-003 is generated with:  
    source: "runtime\_feedback"  
    deviation details: what metric, expected vs actual  
    impacted feature vectors: REQ-F-AUTH-001  
And the intent enters the graph as a new feature vector  
And the specification becomes a living encoding updated from runtime observation

## **UC-06-08: New intent enters graph as new feature vector**

Validates: REQ-LIFE-003 | Fixture: CONVERGED | NEW

Given INT-003 generated from runtime deviation  
When the human approves the intent for action  
Then a new feature vector is spawned from INT-003  
And the feature vector enters the graph at the intent node  
And the trajectory starts: intent → requirements → design → ...

## **UC-06-09: Feedback loop closure from production to spec**

Validates: REQ-LIFE-003 | Fixture: CONVERGED | NEW

Given a deviation detected at telemetry stage  
When the deviation generates an intent that leads to a spec update  
Then the loop is closed: telemetry → intent → spec update → new feature → code → deployment → telemetry  
And the full loop is traceable via events.jsonl  
And the specification is a living encoding updated from runtime observation

## **UC-06-10: Feature lineage queryable in production**

Validates: REQ-LIFE-004 | Fixture: CONVERGED | **NEW**

Given a running system with REQ-tagged telemetry  
When I query production metrics filtered by REQ-F-AUTH-001  
Then I get: latency per feature, error rate per feature, incident count per feature  
And incidents are traceable to originating requirement via telemetry REQ key

## **UC-06-11: REQ keys in telemetry match code tags**

Validates: REQ-LIFE-004 | Fixture: CONVERGED | **NEW**

Given code with "Implements: REQ-F-AUTH-001"  
And telemetry with req="REQ-F-AUTH-001"  
When I cross-reference code tags with telemetry tags  
Then every req tag in telemetry matches an Implements tag in code  
And there are no orphaned telemetry REQ keys

## **UC-06-12: Intent raised event has full causal chain**

Validates: REQ-LIFE-005 | Fixture: IN\_PROGRESS | **EXISTING**  
(TestConsciousnessLoopIntentRaised)

Given an observer that detects a non-trivial delta during iteration  
When an "intent\_raised" event is emitted  
Then the event includes:  
    intent\_id, trigger, delta, signal\_source, vector\_type,  
    affected\_req\_keys, prior\_intents (causal chain), edge\_context,  
    severity  
And the prior\_intents field traces: if intent A led to intent B, chain A→B is visible

## **UC-06-13: Intent raised from development-time signals**

Validates: REQ-LIFE-005 | Fixture: IN\_PROGRESS | **NEW**

Given Feature B at "code↔unit\_tests" with the same test failing for 4 iterations  
When the iterate agent detects the stuck condition  
Then an "intent\_raised" event is emitted with signal\_source: "test\_failure"  
And the delta describes: expected (test passes) vs observed (test keeps failing)  
And the human is presented with the intent for decision (create vector, acknowledge, dismiss)

## **UC-06-14: Human decision on raised intent**

Validates: REQ-LIFE-005 | Fixture: IN\_PROGRESS | NEW

Given an "intent\_raised" event presented to the human  
When the human decides to "create feature vector"  
Then a new feature vector is spawned from the intent  
When the human decides to "dismiss"  
Then the intent is logged as dismissed with reason  
And no new feature vector is created

## **UC-06-15: Seven signal source types recognised**

Validates: REQ-LIFE-006 | Fixture: IN\_PROGRESS | EXISTING  
(TestSignalSourceClassification)

Given all observer points in the system  
When I enumerate the signal source types  
Then exactly 7 types exist: gap, test\_failure, refactoring,  
source\_finding, process\_gap, runtime\_feedback, ecosystem  
And each has an intent template in the feedback loop edge  
configuration  
And the signal source is recorded in every intent\_raised event

## **UC-06-16: Gap analysis emits intent with signal\_source: gap**

Validates: REQ-LIFE-006 | Fixture: IN\_PROGRESS | NEW

Given a gap analysis run that finds REQ-F-AUTH-001 has no test coverage  
When the gap is significant enough to warrant action  
Then an "intent\_raised" event is emitted with signal\_source: "gap"  
And affected\_req\_keys includes REQ-F-AUTH-001  
And the delta describes: expected (test coverage) vs observed (none)

## **UC-06-17: TDD stuck emits test\_failure signal**

Validates: REQ-LIFE-006 | Fixture: STUCK | NEW

Given Feature X with the same test failing for 4+ iterations  
When the iterate agent detects the stuck condition  
Then an "intent\_raised" event is emitted with signal\_source:  
"test\_failure"  
And the delta includes: which test, how many iterations stuck, what was tried

## **UC-06-18: Spec modified event emitted on spec change**

Validates: REQ-LIFE-007 | Fixture: IN\_PROGRESS | EXISTING  
(TestSpecChangeEvents)

Given a signal that leads to a specification update  
When a requirement is added, modified, or deprecated  
Then a "spec\_modified" event is emitted with:  
  trigger\_intent, signal\_source, what\_changed (list of REQ key changes),  
  affected\_req\_keys, spawned\_vectors, prior\_intents

## **UC-06-19: Spec change history reconstructable from events**

Validates: REQ-LIFE-007 | Fixture: IN\_PROGRESS | NEW

Given multiple spec modifications over the project lifecycle  
When I replay the "spec\_modified" events from events.jsonl  
Then the full history of every REQ key modification is reconstructable  
And I can determine which spec sections are volatile vs stable

## **UC-06-20: Feedback loop detection in spec changes**

Validates: REQ-LIFE-007 | Fixture: IN\_PROGRESS | NEW

Given intent A that caused spec change S1  
And spec change S1 that generated intent B  
And intent B that caused spec change S2  
When I trace prior\_intents for intent B  
Then the chain A→B is visible  
And the feedback loop (A → S1 → B → S2) is detectable from the event log

## **UC-06-21: Every iteration emits mandatory event**

Validates: REQ-LIFE-008 | Fixture: IN\_PROGRESS | EXISTING  
(TestProtocolEnforcementHooks)

Given any iteration on any edge  
When the iteration completes (whether converged or not)  
Then an "iteration\_completed" event is emitted to events.jsonl  
And the event includes: feature, edge, iteration number, delta, evaluator results  
And event emission is a reflex-phase operation (fires unconditionally)

## **UC-06-22: Protocol enforcement detects missing side effects**

Validates: REQ-LIFE-008 | Fixture: IN\_PROGRESS | NEW

Given a completed iteration that should have emitted an event  
When the protocol enforcement hook verifies side effects  
Then it checks: event emitted, feature vector updated, project state view regenerated,  
  source\_findings array present, process\_gaps array present

And any missing side effect is flagged as a PROTOCOL\_VIOLATION process gap

### **UC-06-23: Circuit breaker prevents infinite regression**

Validates: REQ-LIFE-008 | Fixture: IN\_PROGRESS | NEW

Given an iteration where event emission itself fails (filesystem error)

When the protocol enforcement hook detects the failure

Then the failure is logged as a TELEMETRY signal in the iteration report

And the iteration does NOT block on the observability failure

And no infinite regression occurs (enforcement failure doesn't trigger enforcement)

### **UC-06-24: Spec review computes delta against workspace**

Validates: REQ-LIFE-009 | Fixture: IN\_PROGRESS | NEW

Given a workspace with 3 features and 15 events

When spec review is invoked

Then it reads: feature vectors, event log, convergence status, telemetry signals

And computes delta against spec: what spec asserts vs what workspace contains

And classifies non-zero deltas by signal source

And generates draft intent\_raised events with full causal chain

And presents draft intents to human for approval

### **UC-06-25: Spec review is stateless and idempotent**

Validates: REQ-LIFE-009 | Fixture: IN\_PROGRESS | NEW

Given the same workspace state and the same spec

When spec review is invoked twice with no intervening changes

Then both invocations produce the same intents

And the review does not modify workspace state

### **UC-06-26: Dev observer watches event stream**

Validates: REQ-LIFE-010 | Fixture: IN\_PROGRESS | NEW

Given the dev observer agent specification

When an "iteration\_completed" event is emitted

Then the dev observer is triggered (via hooks)

And it reads: events.jsonl, feature vector state, convergence status

And computes delta(workspace\_state, spec) → intents

And emits "observer\_signal" events with observer\_id, signal\_source, severity

## **UC-06-27: Dev observer is a Markov object**

Validates: REQ-LIFE-010 | Fixture: IN\_PROGRESS | NEW

Given the dev observer agent  
When it executes  
Then it reads inputs (event log) and emits events (observer\_signal)  
And it has no shared mutable state (event log is the mailbox)  
And draft intents require human approval (actor model – no autonomous spec modification)  
And it is idempotent: same workspace state + same spec = same observations

## **UC-06-28: CI/CD observer maps failures to REQ keys**

Validates: REQ-LIFE-011 | Fixture: IN\_PROGRESS | NEW

Given the CI/CD observer agent specification  
When a CI/CD pipeline fails after a push  
Then the observer reads build logs and test results  
And maps failures back to REQ keys via Implements/Validates tags in failing code  
And emits "observer\_signal" with observer\_id: "cicd\_observer"  
And generates draft intent\_raised events for test gaps and build regressions

## **UC-06-29: CI/CD observer drafts rollback intent**

Validates: REQ-LIFE-011 | Fixture: IN\_PROGRESS | NEW

Given a deployment where health checks fail  
When the CI/CD observer detects the failure  
Then it generates a draft rollback intent (not automatic – human approves)  
And the draft includes: which deployment, what failed, proposed rollback steps

## **UC-06-30: Ops observer correlates telemetry with REQ keys**

Validates: REQ-LIFE-012 | Fixture: CONVERGED | NEW

Given the ops observer agent specification  
And production telemetry with req="REQ-F-AUTH-001" tags  
When the ops observer detects a latency anomaly  
Then it correlates the anomaly with REQ-F-AUTH-001 via the req= tag  
And emits "observer\_signal" with observer\_id: "ops\_observer", affected\_req\_keys  
And generates draft intent\_raised events for SLA breaches

## **UC-06-31: Ops observer runs on schedule**

Validates: REQ-LIFE-012 | Fixture: CONVERGED | NEW

Given the ops observer configured with a daily schedule  
When the scheduled interval elapses  
Then the observer reads production telemetry  
And computes delta against spec constraints (SLA targets, resource bounds)  
And emits `observer_signal` events  
And is stateless: same telemetry + same spec = same observations

## **UC-06-32: Eco-intent generated from ecosystem changes**

Validates: REQ-INTENT-003 | Fixture: CONVERGED | NEW

Given the exteroceptive monitor detects a deprecated API in a dependency  
When the ecosystem signal is classified as requiring action  
Then an INT-ECO-\* intent is generated with:  
  source: "ecosystem"  
  ecosystem context (which dependency, what changed)  
And the intent feeds into the graph as a new feature vector

## **UC-06-33: Eco-intent covers security, deprecation, API, compliance**

Validates: REQ-INTENT-003 | Fixture: CONVERGED | NEW

Given the monitoring system watches for ecosystem changes  
When any of these changes are detected:  
  security vulnerability, API deprecation, upstream API change,  
  compliance update  
Then the corresponding INT-ECO-\* intent is generated  
And each intent type carries the appropriate ecosystem context

---

## **UC-07: Sensory Systems**

**Feature Vector:** REQ-F-SENSE-001 **Satisfies:** REQ-SENSE-001, REQ-SENSE-002, REQ-SENSE-003, REQ-SENSE-004, REQ-SENSE-005

### **UC-07-01: Interoceptive monitors run on schedule**

Validates: REQ-SENSE-001 | Fixture: IN\_PROGRESS | NEW

Given a sensory service with interoceptive monitors configured  
When the configured schedule triggers (default: daily or on workspace open)  
Then all 7 minimum interoceptive monitors execute:

INTRO-001 event freshness, INTRO-002 feature vector stall,  
INTRO-003 test health, INTRO-004 state view freshness,  
INTRO-005 build health, INTRO-006 spec/code drift, INTRO-007 event  
log integrity  
And each monitor produces a typed signal with:  
monitor\_id, observation\_timestamp, metric\_value, threshold, severity  
And monitors are observation-only (do not modify workspace)

## **UC-07-02: Interoceptive signal logged to events**

Validates: REQ-SENSE-001 | Fixture: IN\_PROGRESS | NEW

Given the INTRO-002 monitor detects a feature stalled for 5 days  
When the monitor produces its signal  
Then an "interoceptive\_signal" event is logged to events.jsonl  
And the signal feeds into the affect triage pipeline  
And the signal does NOT go directly to conscious review

## **UC-07-03: Monitor thresholds configurable per project and profile**

Validates: REQ-SENSE-001 | Fixture: INITIALIZED | NEW

Given a project using the hotfix profile  
And the hotfix profile suppresses non-critical interoceptive signals  
When the INTRO-002 (feature stall) monitor fires with severity: info  
Then the signal is suppressed (not escalated)  
When the INTRO-003 (test health) monitor fires with severity: critical  
Then the signal is escalated (critical signals are never suppressed)

## **UC-07-04: Full profile runs all interoceptive monitors**

Validates: REQ-SENSE-001 | Fixture: INITIALIZED | NEW

Given a project using the full profile  
When the sensory service runs its schedule  
Then all 7 interoceptive monitors execute without suppression  
And all signals (info, warning, critical) are processed by the triage  
pipeline

## **UC-07-05: Exteroceptive monitors run on schedule**

Validates: REQ-SENSE-002 | Fixture: IN\_PROGRESS | NEW

Given a sensory service with exteroceptive monitors configured  
When the configured schedule triggers  
Then all 4 minimum exteroceptive monitors execute:  
EXTRO-001 dependency freshness, EXTRO-002 CVE scanning,  
EXTRO-003 runtime telemetry deviation, EXTRO-004 API contract  
changes  
And each monitor produces a typed signal with:

```
    monitor_id, observation_timestamp, external_source, finding,  
severity  
And monitors are observation-only (do not modify workspace)
```

## **UC-07-06: Exteroceptive signal logged to events**

Validates: REQ-SENSE-002 | Fixture: IN\_PROGRESS | NEW

```
Given the EXTR0-002 CVE scanner detects a critical vulnerability  
When the monitor produces its signal  
Then an "exteroceptive_signal" event is logged to events.jsonl  
And the signal feeds into the affect triage pipeline
```

## **UC-07-07: Spike profile disables exteroception**

Validates: REQ-SENSE-002 | Fixture: INITIALIZED | NEW

```
Given a project using the spike profile  
When the sensory service runs its schedule  
Then exteroceptive monitors are disabled entirely  
And no exteroceptive_signal events are emitted  
And interoceptive monitors may still run (if configured for spike)
```

## **UC-07-08: Affect triage — rule-based classification first**

Validates: REQ-SENSE-003 | Fixture: IN\_PROGRESS | NEW

```
Given a sensory signal from INTRO-003 (test health): "3 tests failing"  
When the affect triage pipeline processes the signal  
Then rule-based classification fires first (fast, pattern-matching)  
And the signal is classified: signal_source=test_failure,  
severity=warning  
And if the rule matches, no agent classification is invoked  
And an "affect_triage" event is logged with the classification  
decision
```

## **UC-07-09: Affect triage — agent classification for ambiguous signals**

Validates: REQ-SENSE-003 | Fixture: IN\_PROGRESS | NEW

```
Given a sensory signal that doesn't match any classification rule  
When the affect triage pipeline processes the signal  
Then the probabilistic agent classifier is invoked  
And the agent assigns: signal_source, severity, recommended_action  
And an "affect_triage" event is logged with classification method:  
"agent"
```

## **UC-07-10: Escalation thresholds vary by profile**

Validates: REQ-SENSE-003 | Fixture: INITIALIZED | NEW

Given the hotfix profile with low escalation threshold (escalate aggressively)  
And the spike profile with high escalation threshold (suppress most)  
When the same warning-severity signal is processed  
Then the hotfix profile escalates the signal (below threshold = escalate)  
And the spike profile suppresses the signal (above threshold = defer)  
And below-threshold signals are logged but not escalated

## **UC-07-11: Above-threshold signals generate draft proposals**

Validates: REQ-SENSE-003 | Fixture: IN\_PROGRESS | NEW

Given a critical-severity signal above the escalation threshold  
When the affect triage pipeline processes the signal  
Then a draft proposal is generated via probabilistic agent evaluation  
And the proposal is surfaced to human via the review boundary  
And the proposal does NOT modify workspace files (draft only)

## **UC-07-12: Sensory monitor registry in configuration**

Validates: REQ-SENSE-004 | Fixture: INITIALIZED | EXISTING  
(TestSensorySystemsAcceptanceCriteria)

Given an initialized workspace  
When I inspect the sensory configuration  
Then sensory\_monitors.yml defines which monitors are active, their schedules, and thresholds  
And affect\_triage.yml defines classification patterns, severity mapping, and escalation thresholds  
And profile-level overrides can enable/disable specific monitors

## **UC-07-13: Monitor health — meta-monitoring**

Validates: REQ-SENSE-004 | Fixture: IN\_PROGRESS | NEW

Given the INTRO-003 test health monitor  
When the monitor itself fails to run (e.g., test runner not installed)  
Then an interoceptive meta-signal is generated:  
    "The system senses that its own sensing has failed"  
And the meta-signal has severity: critical (sensing failure is always critical)  
And the meta-signal feeds into the triage pipeline

## **UC-07-14: Review boundary separates sensing from changes**

Validates: REQ-SENSE-005 | Fixture: IN\_PROGRESS | NEW

Given the sensory service has generated a draft proposal  
When the proposal reaches the review boundary  
Then the human can: view the proposal, approve it, or dismiss it  
And approved proposals are applied by the interactive session (NOT the sensory service)  
And the sensory service cannot modify workspace files  
And dismissed proposals are logged with reason

## **UC-07-15: Two event categories enforced at review boundary**

Validates: REQ-SENSE-005 | Fixture: IN\_PROGRESS | NEW

Given the review boundary between sensory service and interactive session  
When events are classified  
Then sensor/evaluate events are autonomous and observation-only  
And change-approval events are conscious and human-approved  
And no sensor event can trigger a file modification  
And human approval is required for all workspace changes

## **UC-07-16: Review boundary preserves human accountability**

Validates: REQ-SENSE-005 | Fixture: IN\_PROGRESS | NEW

Given a draft proposal from the sensory service  
When the human reviews and approves the proposal  
Then the file modification is attributed to the human decision  
And REQ-EVAL-003 (Human Accountability) is preserved  
And the approval event records: who approved, what was proposed, what was applied

---

# **UC-08: Developer Tooling**

**Feature Vector:** REQ-F-TOOL-001 **Satisfies:** REQ-TOOL-001 through REQ-TOOL-010

## **UC-08-01: Plugin is installable and discoverable**

Validates: REQ-TOOL-001 | Fixture: CLEAN | EXISTING (TestPluginJson)

Given a clean project directory  
When I install the AI SDLC methodology plugin  
Then the plugin is discovered by the platform  
And the plugin includes: agent configurations, commands, templates,

edge configs  
And the plugin is versioned (semver)

## UC-08-02: Plugin is versioned with semver

Validates: REQ-TOOL-001 | Fixture: INITIALIZED | EXISTING  
(TestVersionConsistency)

Given an installed plugin  
When I inspect the plugin version  
Then the version follows semantic versioning (MAJOR.MINOR.PATCH)  
And the plugin version is consistent across all plugin artifacts

## UC-08-03: Workspace supports task tracking

Validates: REQ-TOOL-002 | Fixture: INITIALIZED | EXISTING  
(TestTaskUpdateOnConvergence)

Given an initialized workspace  
When I inspect the workspace structure  
Then I find task tracking directories: active, completed, archived  
And context preservation across sessions via the checkpoint mechanism  
And the workspace is version-controlled (git)

## UC-08-04: Context preserved across sessions

Validates: REQ-TOOL-002 | Fixture: IN\_PROGRESS | NEW

Given a feature in progress with 5 iterations completed  
When I close the session and start a new one  
Then the feature's iteration history is preserved  
And the event log contains all previous iterations  
And the workspace state is reconstructable from the event log

## UC-08-05: Workspace is git-integrated

Validates: REQ-TOOL-002 | Fixture: INITIALIZED | NEW

Given an initialized workspace  
When I check the version control status  
Then the .ai-workspace/ directory is within the git repository  
And workspace artifacts are tracked by git  
And context elements are version-controlled

## UC-08-06: Workflow commands for task management

Validates: REQ-TOOL-003 | Fixture: IN\_PROGRESS | NEW

Given an initialized workspace with active features  
When I invoke task management commands

Then I can: create tasks, update task status, complete tasks  
And task status changes are reflected in ACTIVE\_TASKS.md

## **UC-08-07: Workflow commands for context operations**

Validates: REQ-TOOL-003 | Fixture: IN\_PROGRESS | NEW

Given an in-progress workspace  
When I invoke checkpoint  
Then a context snapshot is saved with current timestamp and git ref  
When I invoke restore from a previous checkpoint  
Then the workspace context is restored to that point  
And the restore is non-destructive (no data loss)

## **UC-08-08: Status command shows progress and coverage**

Validates: REQ-TOOL-003 | Fixture: IN\_PROGRESS | NEW

Given a workspace with 3 features at various stages  
When I invoke the status command  
Then I see: progress per feature, coverage gaps, pending actions  
And the status is derived from the event log (not stored state)

## **UC-08-09: Release management with semver**

Validates: REQ-TOOL-004 | Fixture: CONVERGED | NEW

Given all features converged  
When I create a release with version "1.0.0"  
Then the release is tagged with semantic version  
And a changelog is generated from the feature vectors and events  
And the release manifest includes feature vector coverage summary  
And a "release\_created" event is emitted

## **UC-08-10: Release includes feature coverage summary**

Validates: REQ-TOOL-004 | Fixture: CONVERGED | NEW

Given a release manifest for version 1.0.0  
When I inspect the manifest  
Then it lists all included feature vectors with their REQ keys  
And shows coverage percentage (how many REQ keys have full trajectory)  
And flags known gaps (REQ keys without test or telemetry coverage)

## **UC-08-11: Release tagging creates git tag**

Validates: REQ-TOOL-004 | Fixture: CONVERGED | NEW

Given a release created for version 1.0.0  
When I inspect the git repository

Then a tag "v1.0.0" exists  
And the tag points to the commit with the release manifest

## **UC-08-12: Test gap analysis identifies uncovered REQ keys**

Validates: REQ-TOOL-005 | Fixture: IN\_PROGRESS | EXISTING  
(TestTraceabilityPipeline)

Given 5 features with 15 REQ keys total  
When I run the gap analysis command  
Then the analysis identifies which REQ keys have "Validates:" tags in tests  
And which REQ keys are missing test coverage  
And uncovered trajectories are reported with suggested test cases

## **UC-08-13: Gap analysis suggests test cases for gaps**

Validates: REQ-TOOL-005 | Fixture: IN\_PROGRESS | NEW

Given a gap analysis that finds REQ-F-AUTH-002 has no test coverage  
When the analysis report is generated  
Then it suggests test cases that would cover REQ-F-AUTH-002  
And the suggestions reference the acceptance criteria from the requirement

## **UC-08-14: Gap analysis reports as events**

Validates: REQ-TOOL-005 | Fixture: IN\_PROGRESS | NEW

Given a gap analysis run  
When the analysis completes  
Then a "gaps\_validated" event is emitted with:  
  layers\_run, total\_req\_keys, full\_coverage count, test\_gaps count,  
  telemetry\_gaps count  
And the event is queryable from the event log

## **UC-08-15: Methodology hooks trigger on commit**

Validates: REQ-TOOL-006 | Fixture: IN\_PROGRESS | NEW

Given methodology hooks configured for commit events  
When a commit is made  
Then the hook validates: REQ-\* tags present in changed files  
And the hook validates: evaluator convergence recorded for changed features  
And the hook is configurable per project

## **UC-08-16: Methodology hooks trigger on edge transition**

Validates: REQ-TOOL-006 | Fixture: IN\_PROGRESS | NEW

Given methodology hooks configured for edge transitions  
When an edge converges (edge\_converged event)  
Then the hook validates: all required side effects completed  
And the hook fires unconditionally (reflex phase)

## **UC-08-17: Project scaffolding creates structure**

Validates: REQ-TOOL-007 | Fixture: CLEAN | EXISTING (TestInitWorkflow)

Given a clean project directory  
When I initialize the AI SDLC workspace  
Then the scaffolding creates:  
    asset graph configuration, context directories, workspace structure  
    And templates with placeholder guidance are populated  
    And a design-implementation binding manifest is created

## **UC-08-18: Scaffolding generates templates**

Validates: REQ-TOOL-007 | Fixture: CLEAN | NEW

Given a clean project directory  
When I run the project scaffolding  
Then template files are generated:  
    feature\_vector\_template.yml, project\_constraints.yml,  
    graph\_topology.yml  
And each template has placeholder guidance explaining what to fill in  
And the templates are ready for customisation

## **UC-08-19: Context snapshot is immutable**

Validates: REQ-TOOL-008 | Fixture: IN\_PROGRESS | NEW

Given a workspace checkpoint created at timestamp T1  
When I continue working (creating new events, modifying features)  
Then the checkpoint at T1 remains unchanged  
And the checkpoint captures: active tasks, current work context,  
timestamp, git ref  
And a "checkpoint\_created" event is emitted

## **UC-08-20: Context snapshot integrates with task checkpoint**

Validates: REQ-TOOL-008 | Fixture: IN\_PROGRESS | NEW

Given a workspace checkpoint  
When I inspect the checkpoint contents  
Then it includes: active tasks from ACTIVE\_TASKS.md  
And current feature vector states  
And the context hash at checkpoint time  
And the git reference for the commit state

## **UC-08-21: Feature view for a specific REQ key**

Validates: REQ-TOOL-009 | Fixture: IN\_PROGRESS | NEW

Given Feature A with REQ-F-ALPHA-001 referenced in spec, design, and code

When I invoke the feature view for REQ-F-ALPHA-001

Then I see a cross-artifact report:

Spec: tagged (line N), Design: tagged (component X), Code: tagged (file.py:42)

Tests: NOT tagged, Telemetry: NOT tagged

And missing stages are flagged: "No telemetry tagging"

And coverage summary: "3 of 5 stages tagged"

## **UC-08-22: Feature view aggregated across all features**

Validates: REQ-TOOL-009 | Fixture: IN\_PROGRESS | NEW

Given 3 features with a total of 10 REQ keys

When I invoke the feature view for all REQ keys

Then I see a summary table: REQ key × stage (spec, design, code, tests, telemetry)

And each cell shows: tagged / not yet / N/A

And the total coverage percentage is computed

## **UC-08-23: Spec/Design boundary — technology leakage detected**

Validates: REQ-TOOL-010 | Fixture: IN\_PROGRESS | NEW

Given a requirements document that mentions "Django REST Framework"

When the spec/design boundary enforcement evaluator runs

Then it flags "Django REST Framework" as technology leakage in requirements

And the requirement is rejected until the technology reference is removed

And the technology choice belongs in a design ADR, not in the spec

## **UC-08-24: Multiple design variants supported**

Validates: REQ-TOOL-010 | Fixture: INITIALIZED | NEW

Given REQ-F-AUTH-001 in the shared specification

When two implementations (Claude and Gemini) create designs for the same REQ key

Then each implementation has its own design trajectory

And the same REQ key appears in both design documents

And the spec remains technology-agnostic (shared across implementations)

---

# **UC-09: User Experience**

**Feature Vector:** REQ-F-UX-001 **Satisfies:** REQ-UX-001 through REQ-UX-007

## **UC-09-01: State detection — UNINITIALISED**

**Validates:** REQ-UX-001 | **Fixture:** CLEAN | **EXISTING** (TestTwoCommandUX)

Given a project directory with no .ai-workspace/  
When I invoke the start command  
Then the system detects state: UNINITIALISED  
And routes to progressive initialization (5 questions)  
And does NOT require the user to know command names

## **UC-09-02: State detection — NEEDS\_CONSTRAINTS**

**Validates:** REQ-UX-001 | **Fixture:** INITIALIZED | **NEW**

Given an initialized workspace where Feature A has reached "requirements->design"  
And mandatory constraint dimensions (ecosystem, deployment, security, build) are empty  
When I invoke the start command  
Then the system detects state: NEEDS\_CONSTRAINTS  
And prompts for each unresolved mandatory dimension  
And mentions advisory dimensions as optional

## **UC-09-03: State detection — NEEDS\_INTENT**

**Validates:** REQ-UX-001 | **Fixture:** INITIALIZED | **NEW**

Given an initialized workspace with no specification/INTENT.md (or empty placeholder)  
When I invoke the start command  
Then the system detects state: NEEDS\_INTENT  
And prompts: "Describe what you want to build"  
And writes the response to specification/INTENT.md

## **UC-09-04: State detection — NO\_FEATURES**

**Validates:** REQ-UX-001 | **Fixture:** INITIALIZED | **NEW**

Given an initialized workspace with intent but no features in features/active/  
When I invoke the start command  
Then the system detects state: NO\_FEATURES  
And routes to feature creation (spawn command)  
And informs: "Intent captured. Let's create your first feature vector."

## **UC-09-05: State detection — IN\_PROGRESS**

Validates: REQ-UX-001 | Fixture: IN\_PROGRESS | **EXISTING**  
(TestTwoCommandUX)

Given a workspace with 3 active features, 1 converged  
When I invoke the start command  
Then the system detects state: IN\_PROGRESS  
And selects the highest-priority actionable feature  
And determines the next unconverged edge  
And delegates to the iterate function

## **UC-09-06: State detection — ALL\_CONVERGED**

Validates: REQ-UX-001 | Fixture: CONVERGED | **EXISTING**  
(TestTwoCommandUX)

Given a workspace where all active features are fully converged  
When I invoke the start command  
Then the system detects state: ALL\_CONVERGED  
And suggests: gap analysis, release, or new feature creation

## **UC-09-07: State detection — STUCK**

Validates: REQ-UX-001 | Fixture: STUCK | **NEW**

Given Feature X with delta=3 unchanged for 4 consecutive iterations  
When I invoke the start command  
Then the system detects state: STUCK (before IN\_PROGRESS)  
And surfaces the stuck feature with: which checks are failing, how long stuck  
And recommends: spawn discovery, human review, or force-iterate

## **UC-09-08: State detection — ALL\_BLOCKED**

Validates: REQ-UX-001 | Fixture: STUCK | **NEW**

Given all active features are blocked:  
    Feature Y: blocked by spawn dependency  
    Feature Z: blocked by pending human review  
When I invoke the start command  
Then the system detects state: ALL\_BLOCKED  
And surfaces each blocked feature with its reason  
And recommends: work on blocking dependency, or complete pending review

## **UC-09-09: State derived, never stored**

Validates: REQ-UX-001 | Fixture: IN\_PROGRESS | **NEW**

Given a workspace with events.jsonl and feature vector files  
When the start command detects state  
Then the state is computed from workspace filesystem + event log  
And there is no "current\_state" variable stored anywhere  
And the same workspace state always produces the same detected state

## **UC-09-10: Progressive init requires 5 or fewer inputs**

Validates: REQ-UX-002 | Fixture: CLEAN | EXISTING  
(TestTwoCommandUXAcceptanceCriteria)

Given a clean project directory  
When I run progressive initialization  
Then the system asks at most 5 questions:  
    project name (auto-detected, confirm), project kind, language (auto-detected, confirm),  
    test runner (auto-detected, confirm), intent description  
And constraint dimensions are NOT prompted (deferred to design edge)

## **UC-09-11: Constraints deferred until design edge**

Validates: REQ-UX-002 | Fixture: INITIALIZED | EXISTING  
(TestTwoCommandUXAcceptanceCriteria)

Given a freshly initialized workspace  
When I start working on "intent→requirements" edge  
Then the system does NOT ask for ecosystem, deployment, security, or build constraints  
And constraints are only prompted when a feature reaches "requirements→design"

## **UC-09-12: Sensible defaults inferred from project detection**

Validates: REQ-UX-002 | Fixture: CLEAN | NEW

Given a project directory containing pyproject.toml with pytest configuration  
When progressive init runs  
Then language is auto-detected as Python  
And test runner is auto-detected as pytest  
And the user confirms (or overrides) the detected values  
And the system does not ask for information it can infer

## **UC-09-13: Advisory dimensions optional**

Validates: REQ-UX-002 | Fixture: IN\_PROGRESS | NEW

Given a feature reaching the "requirements→design" edge  
When mandatory constraints are prompted  
Then advisory dimensions (data\_governance, performance\_envelope,

```
observability, error_handling)
    are mentioned but skippable ("Press Enter to skip")
    And skipping advisory dimensions does not block convergence
```

## UC-09-14: Status shows “you are here” per feature

Validates: REQ-UX-003 | Fixture: IN\_PROGRESS | **EXISTING**  
(TestTwoCommandUXAcceptanceCriteria)

```
Given a workspace with 3 features at different stages
When I invoke the status command
Then each feature shows its current position in the graph:
    Feature A: intent→req ✓ req→design [iterating] design→code ○
    code↔tests ○
    Feature B: intent→req ✓ req→design ✓ design→code ✓ code↔tests
    [iterating δ=2]
    Feature C: intent→req [iterating] ...
```

## UC-09-15: Status shows cross-feature rollup

Validates: REQ-UX-003 | Fixture: IN\_PROGRESS | **NEW**

```
Given 3 features with various convergence states
When I invoke the status command
Then I see aggregate edge convergence counts:
    "12 of 20 edges converged across 3 features"
And blocked features shown with reason
And unactioned intent_raised events surfaced as signals
```

## UC-09-16: Status previews next action

Validates: REQ-UX-003 | Fixture: IN\_PROGRESS | **NEW**

```
Given an IN_PROGRESS workspace
When I invoke the status command
Then I see a preview of what the start command would do next:
    "Next: REQ-F-ALPHA-001 on requirements→design (iteration 2)"
And the reasoning for the selection is shown
```

## UC-09-17: Status shows workspace health

Validates: REQ-UX-003 | Fixture: IN\_PROGRESS | **NEW**

```
Given a workspace with some health issues
When I invoke the status command
Then the health check reports:
    event log integrity: OK
    feature vector consistency: OK
    orphaned spawns: 1 found
```

stuck features: Feature X (4 iterations, 6 unchanged)  
And remediation suggestions are provided for each issue

## **UC-09-18: Feature selection — time-boxed spawns first**

Validates: REQ-UX-004 | Fixture: IN\_PROGRESS | NEW

Given Feature A (standard, 2 edges remaining)  
And Feature B (time-boxed spike, 25% time remaining, 1 edge remaining)  
When the start command selects a feature  
Then Feature B is selected (time-boxed spawn approaching expiry takes priority)  
And the selection reasoning is displayed

## **UC-09-19: Feature selection — closest-to-complete reduces WIP**

Validates: REQ-UX-004 | Fixture: IN\_PROGRESS | NEW

Given Feature A (3 edges remaining)  
And Feature C (1 edge remaining, same priority)  
And no time-boxed spawns  
When the start command selects a feature  
Then Feature C is selected (closest-to-complete to reduce WIP)

## **UC-09-20: Edge determination — topological walk**

Validates: REQ-UX-004 | Fixture: IN\_PROGRESS | NEW

Given Feature A with converged edges: intent→requirements,  
requirements→design  
And the next edge in topological order is "design→code"  
When the start command determines the next edge  
Then "design→code" is selected  
And converged edges are skipped  
And edges not in the active profile are skipped

## **UC-09-21: User can override feature and edge selection**

Validates: REQ-UX-004 | Fixture: IN\_PROGRESS | NEW

Given the automatic selection would choose Feature A on edge X  
When the user provides --feature "REQ-F-BETA-001" --edge  
"code↔unit\_tests"  
Then the user's override is used instead of automatic selection  
And the system proceeds with the specified feature and edge

## **UC-09-22: Recovery detects corrupted event log**

Validates: REQ-UX-005 | Fixture: IN\_PROGRESS | NEW

Given an events.jsonl file with 3 corrupted lines (malformed JSON)  
When the start command runs its health check  
Then the corruption is detected  
And the system offers: "Event log has 3 corrupted lines. Truncate at last valid line?"  
And the recovery is non-destructive (asks before modifying)

### **UC-09-23: Recovery detects orphaned spawns**

Validates: REQ-UX-005 | Fixture: IN\_PROGRESS | NEW

Given a spawn REQ-F-SPIKE-001 whose parent feature no longer exists  
When the start command runs its health check  
Then the orphaned spawn is detected  
And the system offers: "Spawn REQ-F-SPIKE-001 has no parent. Link to feature or archive?"  
And the user decides the disposition

### **UC-09-24: Workspace rebuildable from event log**

Validates: REQ-UX-005 | Fixture: IN\_PROGRESS | NEW

Given a workspace with corrupted feature vector files  
When the user chooses to rebuild from the event log  
Then feature vector states are reconstructed from events.jsonl  
And task tracking is regenerated  
And the rebuilt workspace matches the event log history

### **UC-09-25: Recovery detects missing feature vectors**

Validates: REQ-UX-005 | Fixture: IN\_PROGRESS | NEW

Given the event log references feature REQ-F-ALPHA-001  
But the feature vector file features/active/REQ-F-ALPHA-001.yml is missing  
When the start command runs its health check  
Then the missing vector is detected  
And the system offers to regenerate the vector from events

### **UC-09-26: Escalation displayed inline in interactive mode**

Validates: REQ-UX-006 | Fixture: IN\_PROGRESS | NEW

Given an IntentEngine escalation during an iteration  
When the escalation occurs in interactive CLI mode  
Then the escalation is displayed inline with clear visual distinction  
And the escalation is NOT buried in normal output  
And the user can act on it immediately

## **UC-09-27: Pending escalations shown in status**

Validates: REQ-UX-006 | Fixture: IN\_PROGRESS | NEW

Given 2 pending escalations in the review queue  
When I invoke the status command  
Then I see: "2 pending escalations" with their summaries  
And the start command checks escalations before feature selection  
And escalations take priority over normal iteration

## **UC-09-28: Unactioned escalations re-surface with elevated urgency**

Validates: REQ-UX-006 | Fixture: IN\_PROGRESS | NEW

Given an escalation that has been pending for 3 iterations of the affected feature  
When the next iteration begins  
Then the escalation is re-surfaced with elevated urgency  
And the urgency level is visually distinct from the original  
And no escalation is silently dropped – the review queue guarantees delivery

## **UC-09-29: Escalation queue in async mode**

Validates: REQ-UX-006 | Fixture: IN\_PROGRESS | NEW

Given an IntentEngine escalation during sensory service processing (async mode)  
When the escalation occurs outside an interactive session  
Then the escalation is written to .ai-workspace/reviews/pending/ as structured YAML  
And the YAML includes: source level, triggering signal, affected features, timestamp  
And the next /gen-status displays the pending count

## **UC-09-30: Zoom in expands edge to sub-graph**

Validates: REQ-UX-007 | Fixture: IN\_PROGRESS | NEW

Given the edge "design→code" for Feature A  
When the user requests zoom: --edge "design→code" --level in  
Then the system proposes intermediate node types:  
    design → module\_decomp → basis\_projections → code\_per\_module → code  
And each intermediate node gets its own asset type in the workspace  
And sub-edges inherit parent evaluators (overridable)  
And a "graph\_zoom\_in" event is emitted with full before/after topology

## **UC-09-31: Zoom out collapses sub-graph**

Validates: REQ-UX-007 | Fixture: IN\_PROGRESS | NEW

Given a previously zoomed "design→code" edge with 3 intermediate nodes  
And all intermediate nodes are converged  
When the user requests zoom: --edge "design→code" --level out  
Then the sub-graph collapses back to a single "design→code" edge  
And intermediate assets are archived (not deleted)  
And convergence state reflects sub-graph completion  
And a "graph\_zoom\_out" event is emitted

### **UC-09-32: Selective zoom retains waypoints**

Validates: REQ-UX-007 | Fixture: IN\_PROGRESS | **NEW**

Given a zoomed edge with 3 intermediate nodes  
When the user selects node "module\_decomp" as a mandatory waypoint  
And collapses the other intermediates  
Then "module\_decomp" becomes an explicit convergence checkpoint  
And the other intermediates are treated as internal to encapsulating edges  
And a "graph\_zoom\_selective" event is emitted

### **UC-09-33: Zoom preserves graph invariants**

Validates: REQ-UX-007 | Fixture: IN\_PROGRESS | **NEW**

Given a zoom operation on any edge  
When the zoom completes  
Then the graph remains directed, typed, and all edges have evaluators  
And feature vectors traversing the zoomed edge see the sub-graph topology  
And the zoom is reversible and auditable via the event log

---

## **UC-10: Multi-Agent Coordination**

**Feature Vector:** REQ-F-COORD-001 **Satisfies:** REQ-COORD-001 through REQ-COORD-005

### **UC-10-01: Agent identity on all events**

Validates: REQ-COORD-001 | Fixture: IN\_PROGRESS | **EXISTING**  
(TestMultiAgentCoordinationAcceptanceCriteria)

Given multi-agent mode is active  
When any agent emits an event to events.jsonl  
Then the event includes agent\_id (unique per instance) and agent\_role  
And agent\_role is drawn from the project-defined role registry  
And agent identity is self-declared (not centrally assigned)

## **UC-10-02: Single-agent backward compatibility**

Validates: REQ-COORD-001 | Fixture: IN\_PROGRESS | NEW

Given single-agent mode (default)  
When events are emitted  
Then agent\_id and agent\_role fields are optional  
And if present, default to agent\_id: "primary"  
And all existing single-agent workflows continue to function

## **UC-10-03: Agent role registry is simple YAML**

Validates: REQ-COORD-001 | Fixture: INITIALIZED | NEW

Given an initialized workspace in multi-agent mode  
When I inspect the agent role registry  
Then it is a simple YAML list: [architect, tdd\_engineer, documentation, ...]  
And there is no role hierarchy or inheritance  
And the event log is the source of truth for which agents have participated

## **UC-10-04: Feature assignment via edge\_claim event**

Validates: REQ-COORD-002 | Fixture: IN\_PROGRESS | NEW

Given Agent A wants to work on Feature X's "code->unit\_tests" edge  
When Agent A emits an "edge\_claim" event with agent\_id, feature, and edge  
Then the serialiser resolves the claim:  
  If no other agent holds this edge: emit "edge\_started" (claim granted)  
  If another agent holds this edge: emit "claim\_rejected" with reason and holding\_agent  
And the claim resolution is atomic (first claim wins)

## **UC-10-05: No lock files for assignment**

Validates: REQ-COORD-002 | Fixture: IN\_PROGRESS | NEW

Given multi-agent feature assignment  
When I inspect the workspace for assignment state  
Then there are no lock files, lease files, or mutex mechanisms  
And assignment state is entirely derivable from the event log  
And current assignments are a derived projection (replay edge\_claim, edge\_started, edge\_converged)

## **UC-10-06: Stale claims detectable**

Validates: REQ-COORD-002 | Fixture: IN\_PROGRESS | NEW

Given Agent A claimed Feature X's edge 2 hours ago  
And no follow-up events from Agent A since then  
And the configured timeout is 1 hour  
When the health check runs  
Then a "claim\_expired" telemetry signal is generated  
And the edge becomes available for other agents to claim

## **UC-10-07: Single-agent mode skips claim step**

Validates: REQ-COORD-002 | Fixture: IN\_PROGRESS | NEW

Given single-agent mode (one agent writing events)  
When the agent starts work on an edge  
Then it emits "edge\_started" directly (no claim step needed)  
And the single-agent workflow is simpler than multi-agent

## **UC-10-08: Agent working state isolated from shared state**

Validates: REQ-COORD-003 | Fixture: IN\_PROGRESS | NEW

Given Agent A working on Feature X's "design->code" edge  
When Agent A generates code drafts  
Then drafts are stored in Agent A's private directory  
(agents/<id>/drafts/)  
And the shared project state (features/, events/, spec/) is not  
written directly  
And only converged assets enter the shared workspace (via promotion)

## **UC-10-09: Promotion requires evaluator gate**

Validates: REQ-COORD-003 | Fixture: IN\_PROGRESS | NEW

Given Agent A has produced a code candidate in its private workspace  
When Agent A requests promotion to shared state  
Then the configured evaluators for this edge must pass  
And if human review is configured, the human must approve  
And only after all evaluators pass does the asset enter shared paths

## **UC-10-10: Agent crash only loses private state**

Validates: REQ-COORD-003 | Fixture: IN\_PROGRESS | NEW

Given Agent A has emitted 5 iteration events and has private drafts  
When Agent A crashes mid-iteration  
Then the 5 emitted events persist in events.jsonl  
And Agent A's private working state may be discarded  
And shared project state is unaffected  
And another agent (or a restarted Agent A) can continue from the event log

## **UC-10-11: Orthogonal features freely parallelisable**

Validates: REQ-COORD-004 | Fixture: IN\_PROGRESS | NEW

Given Feature A (authentication) and Feature B (reporting)  
And the inner product (shared module count) is zero  
When two agents are assigned to Feature A and Feature B respectively  
Then they can work in parallel with no coordination  
And their iterations do not conflict (orthogonal trajectories)

## **UC-10-12: Non-orthogonal features trigger warning**

Validates: REQ-COORD-004 | Fixture: IN\_PROGRESS | NEW

Given Feature A and Feature C that share a "user\_service" module  
And the inner product (shared module count) is 1  
When both features are assigned to different agents concurrently  
Then the system triggers a warning: "Non-orthogonal features assigned concurrently"  
And suggests sequencing: build shared module first, then diverge

## **UC-10-13: Inner product computed from dependency DAG**

Validates: REQ-COORD-004 | Fixture: IN\_PROGRESS | NEW

Given the module dependency DAG for the project  
When the inner product is computed for two features  
Then it uses the actual module dependency graph (not a heuristic)  
And the result is the count of shared modules  
And zero shared modules = safe parallelism

## **UC-10-14: Role-based authority scopes convergence**

Validates: REQ-COORD-005 | Fixture: IN\_PROGRESS | NEW

Given the role registry:  
tdd\_engineer: [code\_unit\_tests]  
architect: [requirements\_design, design\_code]  
When Agent A (role: tdd\_engineer) attempts to converge  
"requirements->design"  
Then the convergence is rejected (outside role authority)  
And a "convergence\_escalated" event is emitted  
And the edge is escalated to an agent with architect role or a human

## **UC-10-15: Human authority is universal**

Validates: REQ-COORD-005 | Fixture: IN\_PROGRESS | NEW

Given any edge type in the graph  
When a human evaluator reviews and approves

Then the convergence is accepted regardless of the edge type  
And human authority is universal (can converge any edge)

## **UC-10-16: Role definitions are project-configurable**

Validates: REQ-COORD-005 | Fixture: INITIALIZED | NEW

Given a project that defines custom roles:

```
security_reviewer: [design_code, code_unit_tests]  
data_engineer: [design_code]
```

When the role registry is loaded

Then the custom roles are used for authority scoping  
And the roles are not hard-coded in the engine

---

## **UC-11: IntentEngine / Supervision**

Feature Vector: REQ-F-SUPV-001 Satisfies: REQ-SUPV-001, REQ-SUPV-002

### **UC-11-01: Every edge traversal produces IntentEngine output**

Validates: REQ-SUPV-001 | Fixture: IN\_PROGRESS | NEW

Given any edge traversal (e.g., Feature A on "requirements→design")  
When the iterate function completes  
Then a classified observation is produced (observer output)  
And a typed routing decision is produced (evaluator output)  
And the output is exactly one of: reflex.log, specEventLog, or  
escalate

### **UC-11-02: Ambiguity classified into three regimes**

Validates: REQ-SUPV-001 | Fixture: IN\_PROGRESS | NEW

Given an observation from an edge evaluator  
When the evaluator classifies the ambiguity  
Then the classification is exactly one of:  
zero ambiguity → reflex (fire-and-forget event)  
bounded nonzero → probabilistic disambiguation (deferred intent)  
persistent/unbounded → escalate to higher consciousness level  
And the classification maps to the three evaluator types

### **UC-11-03: Three output types are exhaustive**

Validates: REQ-SUPV-001 | Fixture: IN\_PROGRESS | NEW

Given any observation at any scale (iteration, edge, feature,  
production)

When the IntentEngine processes the observation

Then the output is exactly one of:

- reflex.log (fire-and-forget event → events.jsonl)
- specEventLog (deferred intent for further processing)
- escalate (push to higher consciousness level)

And these three types map to existing event types (no new types required)

And every possible observation routes to exactly one output type

### **UC-11-04: Affect propagates through chained invocations**

Validates: REQ-SUPV-001 | Fixture: IN\_PROGRESS | NEW

Given an IntentEngine invocation at Level 1 (single iteration)

And the invocation carries affect: urgency=high, valence=negative

When the output is "escalate" (pushing to Level 2: edge convergence)

Then Level 2 receives the intent + affect from Level 1

And the urgency/valence propagates and transforms at each level

And the affect influences the evaluation at Level 2

### **UC-11-05: Level N escalate becomes Level N+1 reflex input**

Validates: REQ-SUPV-001 | Fixture: IN\_PROGRESS | NEW

Given Level 1 (iteration) produces output: escalate

When Level 2 (edge convergence) receives the escalation

Then Level 2 treats it as reflex input (automatic processing)

And Level 1's reflex.log events are invisible to Level 2 (handled at origin)

And this implements consciousness-as-relative (each level observes a different scale)

### **UC-11-06: IntentEngine applies at every scale**

Validates: REQ-SUPV-001 | Fixture: IN\_PROGRESS | NEW

Given the fractal application table:

- Level 1: single iteration
- Level 2: edge convergence
- Level 3: feature traversal
- Level 4: sensory monitoring
- Level 5: production homeostasis
- Level 6: spec review

When each level processes an observation

Then each uses the same IntentEngine pattern: observer → evaluator → typed\_output

And the pattern composes the four primitives (Graph, Iterate, Evaluators, Spec+Context)

## **UC-11-07: IntentEngine interface on all actors**

Validates: REQ-SUPV-001 | Fixture: IN\_PROGRESS | NEW

Given actors: iterate agent, sensory monitors, protocol hooks, observers

When each actor processes an observation

Then each exposes the IntentEngine interface:

observer → evaluator → typed\_output(reflex.log | specEventLog | escalate)

And the interface is parameterised by intent + affect

And this is not a fifth primitive – it is the composition law

## **UC-11-08: Every constraint has a measurable tolerance**

Validates: REQ-SUPV-002 | Fixture: INITIALIZED | NEW

Given the spec constraints, design bindings, and edge evaluator thresholds

When I inspect each constraint

Then every constraint has an associated measurable threshold

And a constraint without a tolerance is flagged by gap analysis as incomplete

And the tolerance makes the gradient operational: delta(state, constraint) is computable

## **UC-11-09: Tolerance breach produces classified signal**

Validates: REQ-SUPV-002 | Fixture: IN\_PROGRESS | NEW

Given a constraint with tolerance: coverage >= 80%

When the measured coverage is 75% (within drift range, say 70–80%)

Then the signal is classified as: specEventLog (optimization intent deferred)

When the measured coverage is 50% (below breach threshold)

Then the signal is classified as: escalate (corrective intent raised)

When the measured coverage is 85% (within bounds)

Then the signal is classified as: reflex.log (everything fine)

## **UC-11-10: Sensory monitors use tolerances as evaluation criteria**

Validates: REQ-SUPV-002 | Fixture: IN\_PROGRESS | NEW

Given an interoceptive monitor INTRO-003 (test health)

And the tolerance: all tests must pass, warning at >0 failures

When the monitor observes 2 test failures

Then the monitor reports: metric\_value=2, threshold=0, severity=warning

And the IntentEngine classifies the delta using the tolerance

And the tolerance defines the threshold, not the monitor

## **UC-11-11: Design bindings have performance/cost tolerances**

Validates: REQ-SUPV-002 | Fixture: IN\_PROGRESS | NEW

Given a design ADR choosing PostgreSQL with tolerance: query latency p99 < 100ms

When production telemetry shows query latency p99 = 150ms

Then the tolerance is breached

And an optimization intent is raised: "PostgreSQL latency exceeds tolerance"

And the intent includes: current value, threshold, recommended action

## **UC-11-12: Tolerances at every scale of the gradient**

Validates: REQ-SUPV-002 | Fixture: IN\_PROGRESS | NEW

Given tolerances defined at multiple scales:

Iteration: max 10 iterations per edge

Edge: all required checks must pass

Feature: all profile edges must converge

Production: SLA p99 < 200ms

Spec review: no more than 5 unresolved intents

When each scale monitors its tolerance

Then breach at any scale triggers the appropriate IntentEngine level

And each level classifies the breach into reflex/specEventLog/escalate

## **UC-11-13: Tolerance pressure balances escalation pressure**

Validates: REQ-SUPV-002 | Fixture: IN\_PROGRESS | NEW

Given escalation pressure: "This edge has too much ambiguity → add more structure"

And tolerance pressure: "This edge has too many checks → simplify"

When both pressures are active on the same edge

Then the graph topology reaches equilibrium:

enough structure to handle ambiguity, not so much that cost is breached

And tolerance breach (complexity/cost) triggers simplification intents

And escalation (persistent ambiguity) triggers structure-adding intents

## **UC-11-14: Constraint without tolerance flagged as incomplete**

Validates: REQ-SUPV-002 | Fixture: IN\_PROGRESS | NEW

Given a requirement REQ-NFR-PERF-001: "The system shall be fast"  
with no measurable threshold defined

When gap analysis runs

Then the requirement is flagged: "Constraint without tolerance – inert"

And the gap report recommends: define a measurable threshold (e.g., p99 < 200ms)

And without a tolerance, the sensory system has nothing to measure

---

## Appendix A: Existing Test Coverage Map

Summary of existing tests and their relationship to UAT scenarios.

Test File	Tests	Coverage Focus	UC Ove
imp_claude/tests/test_config_validation.py	~100	YAML syntax, graph topology, edge configs, profiles, plugin structure	UC-01 (structure) UC-02 (structure) UC-05 (structure)
imp_claude/tests/test_methodology_bdd.py	~160	BDD acceptance for all feature vectors — checks design/agent docs reference REQ keys	UC-01–1 (structure)
imp_claude/tests/test_spec_validation.py	~40	Spec-level requirements existence, lineage, processing phases	UC-02, 1 UC-07, 1 (existence)
imp_claude/tests/test_integration_uat.py	~45	End-to-end traceability, event log integrity, feature vector consistency	UC-04 (integrity) UC-06 (integrity)
imp_claude/tests/e2e/	22	Headless CLI convergence (standard)	UC-01, 1 UC-09 (convergence scenario)

Test File	Tests	Coverage Focus	UC Ove
		profile, temperature converter)	

## Coverage Gap Summary

The existing 345+ non-E2E tests validate **specification content** and **configuration structure**. They confirm that:

- Config files parse correctly and have required fields
- Design documents reference the right REQ keys
- Agent specifications mention the right concepts
- Feature vector templates have the right shape

They do **NOT** validate:

- State machine transitions execute correctly (UC-09-01 through UC-09-08)
- Iterate function actually reduces delta (UC-01-10 through UC-01-18)
- Evaluators actually compute delta against spec (UC-02-01 through UC-02-03)
- TDD/BDD patterns produce correct artifacts (UC-05-01 through UC-05-09)
- Multi-agent claim/reject/expire flow works (UC-10-04 through UC-10-07)
- IntentEngine classification produces correct output types (UC-11-01 through UC-11-07)
- Sensory monitors detect actual workspace conditions (UC-07-01 through UC-07-07)
- Recovery procedures repair corrupted workspaces (UC-09-22 through UC-09-25)
- Observer agents compute delta and generate intents (UC-06-26 through UC-06-31)
- Constraint tolerances trigger correct signal classifications (UC-11-09 through UC-11-14)

**~162 scenarios are genuinely NEW (no existing functional coverage). ~49 scenarios overlap with existing tests** (marked EXISTING with test class reference).

---

## Appendix B: Fixture Specifications

### Fixture File Structure

Each fixture is a reproducible directory tree. Implementations should provide fixture generators that produce these states deterministically.

#### CLEAN Fixture

```
project/
└── src/
    └── main.py          # placeholder source
└── pyproject.toml      # Python project config
└── (no .ai-workspace/)
```

#### INITIALIZED Fixture

```
project/
└── src/main.py
└── pyproject.toml
```

```

└── specification/
    └── INTENT.md          # "Build a temperature converter CLI"
└── .ai-workspace/
    ├── graph/
    │   └── graph_topology.yml
    ├── edges/              # edge param files
    ├── profiles/           # 6 profile YAMLs
    ├── features/
    │   └── active/          # empty
    │       └── feature_vector_template.yml
    ├── context/
    │   └── project_constraints.yml
    ├── events/
    │   └── events.jsonl     # 1 event: project_initialized
    ├── tasks/
    │   └── active/ACTIVE_TASKS.md
    └── STATUS.md

```

## IN\_PROGRESS Fixture

Extends INITIALIZED with:

```

.ai-workspace/features/active/
├── REQ-F-ALPHA-001.yml      # at requirements→design, iteration 2
├── REQ-F-BETA-001.yml        # at code↔unit_tests, iteration 3, δ=2
└── REQ-F-GAMMA-001.yml       # at intent→requirements, iteration 1

.ai-workspace/events/events.jsonl  # ~15 events

```

Event sequence for IN\_PROGRESS: 1. project\_initialized 2. edge\_started (ALPHA, intent→req) 3. iteration\_completed (ALPHA, intent→req, δ=1) 4. iteration\_completed (ALPHA, intent→req, δ=0) 5. edge\_converged (ALPHA, intent→req) 6. edge\_started (ALPHA, req→design) 7. iteration\_completed (ALPHA, req→design, δ=3) 8. iteration\_completed (ALPHA, req→design, δ=2) — current 9. edge\_started (BETA, intent→req) ... through edge\_converged (BETA, design→code) 10. edge\_started (BETA, code↔unit\_tests) 11-13. iteration\_completed (BETA, code↔unit\_tests, δ=2) × 3 iterations 14. edge\_started (GAMMA, intent→req) 15. iteration\_completed (GAMMA, intent→req, δ=2) — current

## CONVERGED Fixture

Extends INITIALIZED with 2 fully converged features and complete event trajectories.

## STUCK Fixture

Extends IN\_PROGRESS with: - Feature X: 4 iteration events all showing δ=3 on same checks - Feature Y: spawn\_created event for REQ-F-SPIKE-001, no spawn\_folded\_back - Feature Z: review\_completed pending (no decision event)

---

# Summary

UC Cluster	Feature Vector	REQ Keys	Scenarios	New	Existing
UC-01: Asset Graph Engine	REQ-F-ENGINE-001	6	23	19	4
UC-02: Evaluator Framework	REQ-F-EVAL-001	3	15	11	4
UC-03: Context Management	REQ-F-CTX-001	3	12	11	1
UC-04: Feature Traceability	REQ-F-TRACE-001	5	18	14	4
UC-05: Edge Parameterisations	REQ-F-EDGE-001	4	16	14	2
UC-06: Full Lifecycle Closure	REQ-F-LIFE-001	13	33	27	6
UC-07: Sensory Systems	REQ-F-SENSE-001	5	16	15	1
UC-08: Developer Tooling	REQ-F-TOOL-001	10	24	17	7
UC-09: User Experience	REQ-F-UX-001	7	33	27	6
UC-10: Multi-Agent Coordination	REQ-F-COORD-001	5	16	15	1
UC-11: IntentEngine / Supervision	REQ-F-SUPV-001	2	14	14	0
<b>Total</b>	<b>11 vectors</b>	<b>64</b>	<b>220</b>	<b>184</b>	<b>36</b>

**220 scenarios. 64 REQ keys. 100% coverage. 184 genuinely new functional scenarios.**

---

**Traces To:** [AISDLC IMPLEMENTATION REQUIREMENTS.md](#),  
[FEATURE VECTORS.md](#), [AI SDLC ASSET GRAPH MODEL.md](#)