

AI SDLC – User Experience

Version: 1.0.0 **Date:** 2026-02-23 **Derived From:**
[AI SDLC ASSET GRAPH MODEL.md](#) (v2.8.0),
[AISDLC IMPLEMENTATION REQUIREMENTS.md](#) (v3.11.0)

1. Purpose

This document defines the **user experience** of the AI SDLC methodology. It answers three questions:

- 1. What are the fundamentally different ways a human uses this system?** (User Journeys)
- 2. Which features are required for minimum viability?** (MVP tagging per journey)
- 3. Do the requirements compose into coherent end-to-end paths?** (Validation Scenarios)

Relationship to Other Spec Documents

Document	What it specifies	This document's relationship
<u>AI SDLC ASSET GRAPH MODEL.md</u>	The formal system (constraints, invariants, symmetries)	Journeys must be satisfiable within the formal system
<u>AISDLC IMPLEMENTATION REQUIREMENTS.md</u>	Platform-agnostic WHAT (including §11 REQ-UX-* requirements)	Journeys exercise requirement end-to-end; gaps drive new REQ keys
<u>UAT TEST CASES.md</u>	Exhaustive functional use cases per REQ key	Scenarios here are cross-cutting — each touches many REQ keys simultaneously

Document	What it specifies	This document's relationship
<u>FEATURE VECTORS.md</u>	Feature decomposition	Journeys may reveal missing feature vectors
<u>PROJECTIONS AND INVARIANTS.md</u>	Projection profiles and vector types	Journeys exercise specific profiles end-to-end

This document is a gap amplifier. Where UAT_TEST_CASES.md validates individual requirements in isolation, the scenarios here validate that requirements compose into coherent user journeys. A scenario that cannot be completed exposes integration gaps.

2. User Journey Map

A user journey is a mode of engagement. It defines who the user is, what they want, how they interact (or don't), and what guarantees the system provides.

2.1 Journey Overview

ID	Journey	Interaction	Actor	Profile(s)	MVP
UJ-1	<u>Interactive Build</u>	Human-driven, edge-by-edge	Developer	standard, full	Yes
UJ-2	<u>Unattended Build</u>	Autonomous, full trajectory	Developer / CI	full-auto, poc, spike	Yes
UJ-3	<u>Resume & Orient</u>	Human queries state	Returning developer	any	Yes
UJ-4	<u>Review & Approve</u>	Human gate at edge boundary	Tech lead / reviewer	full, standard	Yes
UJ-5	<u>Background Sensing</u>	Autonomous observation, human reviews proposals	Ops / maintainer	full, standard	No
UJ-6	<u>CI/CD Integration</u>	Pipeline-triggered, results to queue	Build system	full-auto, hotfix	No

ID	Journey	Interaction	Actor	Profile(s)	MVP
UJ-7	<u>Observation & Monitoring</u>	Read-only dashboard	Any stakeholder	any	Yes

Journey Dependency Map

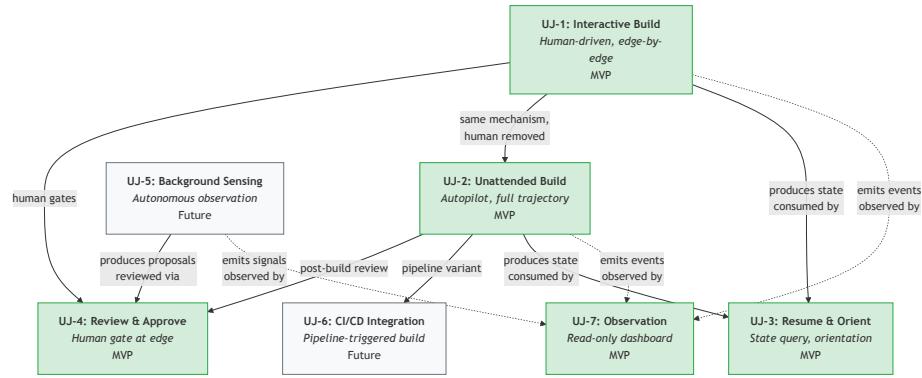


Diagram 0

UJ-1 and UJ-2 are the two primary build journeys. Everything else is a supporting journey (UJ-3 orient, UJ-4 review, UJ-7 observe) or a future extension (UJ-5 sensing, UJ-6 CI/CD).

The critical insight: **UJ-2 is not a degraded UJ-1.** It is a distinct journey with stronger pre-conditions, weaker in-flight guarantees, and an explicit post-build verification step. The user makes a deliberate choice to trade in-flight human judgment for velocity, accepting the post-build review responsibility.

2.2 UJ-1: Interactive Build

The canonical human-in-the-loop development journey.

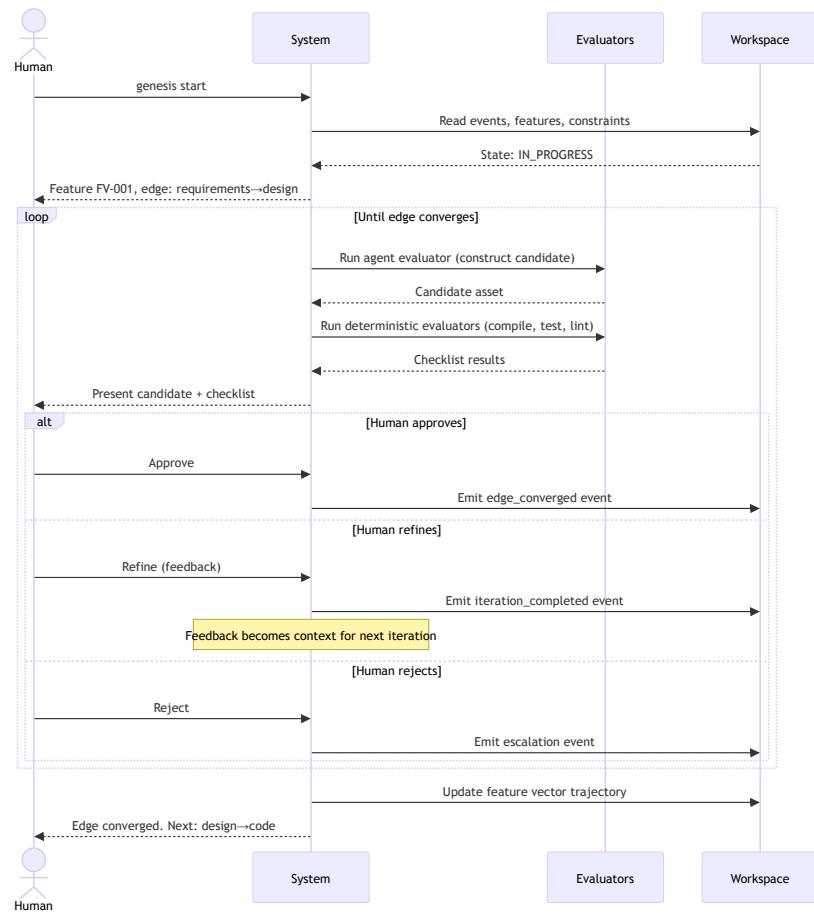


Diagram 1

Who

Developer working on a feature. May be first-time or experienced.

When

- New project, new feature, or resuming existing work
- Developer wants control over each edge transition
- Requirements are still forming — human judgment needed at every stage

Key Characteristics

Aspect	Detail
Edge traversal	One at a time, human confirms transitions
Evaluators	Human + Agent + Deterministic (per profile)
Conscious phase	Human decides — approves, refines, rejects
Escalations	Inline — displayed in session
Feedback loop	Immediate — human sees delta, adjusts

Aspect	Detail
State routing	System suggests, human confirms

Invariants (all hold)

- [MVP] Evaluator Existence — every edge has evaluators
- [MVP] Markov Property — state derived from workspace, not session memory
- [MVP] Convergence Monotonicity — delta decreases or stuck detection triggers
- [MVP] IntentEngine Composition — observer → evaluator → typed output at every scale
- [MVP] Review Boundary — human crosses it explicitly (approve/refine/reject)
- [MVP] Event Sourcing — every iteration emits events, state reconstructible

Pre-conditions

- [MVP] Workspace initialised (.ai-workspace/ exists)
- [MVP] Intent defined (specification/INTENT.md or INT-* event)
- [MVP] At least one feature vector created

Post-conditions

- [MVP] Feature vector trajectory shows all edges converged
- [MVP] Event log contains complete audit trail
- [MVP] REQ key traceability: intent → requirements → design → code → tests
- /gen-gaps Layer 1 + Layer 2 at target coverage

MVP Features

- [MVP] State-driven routing (/gen-start detects state, routes to action) — REQ-UX-001
 - [MVP] Progressive disclosure (ask only what's needed now) — REQ-UX-002
 - [MVP] Human gate pause and review — REQ-UX-006, REQ-EVAL-003
 - [MVP] Feature and edge auto-selection — REQ-UX-004
 - [MVP] Event emission at every iteration boundary — REQ-LIFE-004
 - [MVP] Evaluator checklist execution — REQ-EVAL-001, REQ-EVAL-002
 - Spec-boundary review (/gen-spec-review) — REQ-EDGE-003
 - Edge zoom management — REQ-UX-007
-

2.3 UJ-2: Unattended Build (Autopilot)

Agent traverses the full graph autonomously. No human gates. Human verifies afterward.

First demonstrated in `data_mapper.test06`: a pre-populated spec + constraints + feature vector, launched with full-auto profile, producing a complete Scala/Spark codebase (8 sbt modules, 95 passing tests, 13 events) with zero human interaction.

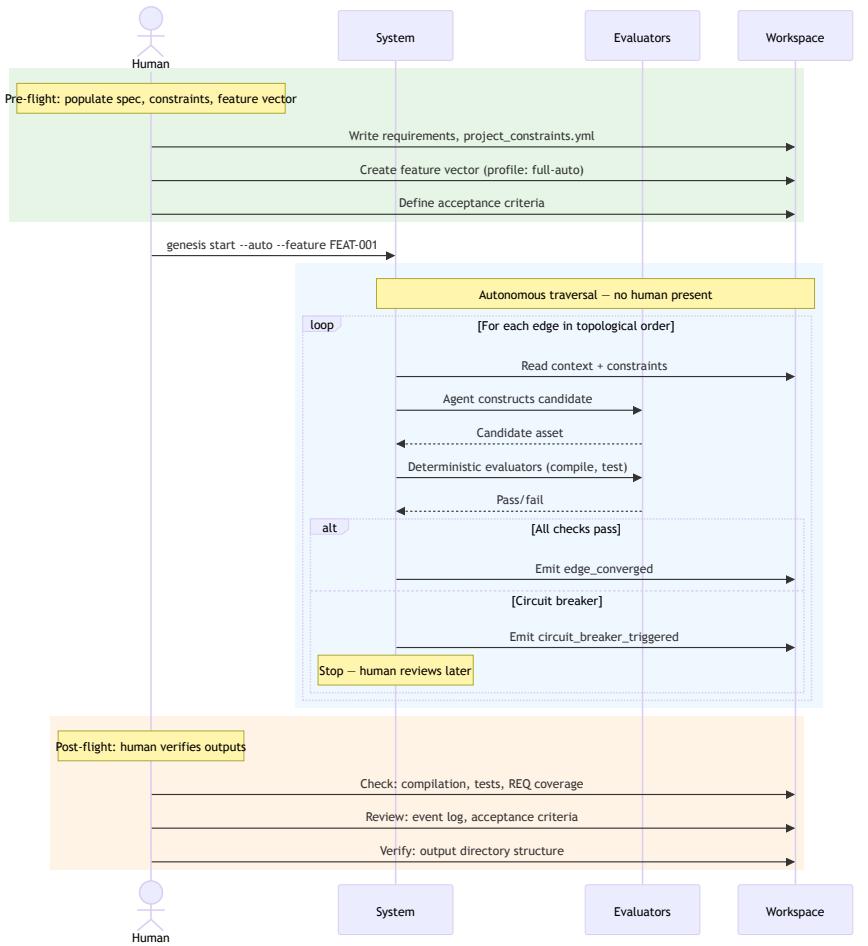


Diagram 2

Who

- Developer with a well-defined spec who wants to see what the agent produces end-to-end
- CI/CD pipeline running a build-from-spec job
- Methodology tester validating end-to-end convergence

When

- Spec is complete and stable — requirements, constraints, and acceptance criteria defined
- Human review would add latency without proportional value (PoC, spike, re-build, demo)
- User wants to evaluate agent capability on a known-good spec
- CI/CD pipeline builds from spec on commit

Key Characteristics

Aspect	Detail
Edge traversal	All edges in topological order, no pause

Aspect	Detail
Evaluators	Agent + Deterministic only — human evaluator removed
Conscious phase	Agent decides — no human judgment during build
Escalations	Agent self-resolves or circuit-breaker triggers
Feedback loop	Deferred — human reviews after completion
State routing	Automatic — system selects next edge, no confirmation
Duration	Minutes to hours depending on project complexity

Invariants

Hold: - [MVP] Evaluator Existence — every edge still has evaluators (agent + deterministic) - [MVP] Markov Property — state derived from workspace + events - [MVP] Convergence Monotonicity — delta decreases or circuit-breaker stops - [MVP] Event Sourcing — every iteration emits events - [MVP] IntentEngine Composition — observer → evaluator → typed output (agent replaces human at conscious level) - [MVP] REQ key traceability — tags in code, tests, events

Relaxed: - Review Boundary — agent crosses it autonomously (human is not present) - Human Accountability (REQ-EVAL-003) — deferred to post-build verification - Escalation delivery (REQ-UX-006) — agent self-resolves; unresolvable items recorded in event log for post-build review

Pre-conditions (more demanding than UJ-1)

- [MVP] Workspace initialised with graph topology scaffolded
- [MVP] Spec complete: requirements document with REQ keys, acceptance criteria, priority
- [MVP] Constraints bound: `project_constraints.yml` with mandatory dimensions filled (ecosystem, deployment, security, build)
- [MVP] Structure defined: `structure.design_tenants` with output directory binding — REQ-TOOL-012, REQ-TOOL-013
- [MVP] Feature vector created with profile set to `full-auto` (or any profile with `human_required_on_all_edges: false`)
- [MVP] Acceptance criteria defined in feature vector `constraints.acceptance_criteria`
- Circuit breaker configured (max iterations per edge, stall detection)

Post-conditions (human verifies these)

- [MVP] **Compilation:** all generated code compiles/parses without error

- [MVP] **Tests pass:** unit tests + UAT tests green
- [MVP] **REQ coverage:** grep -r "Implements: REQ-" covers all REQ keys from spec
- [MVP] **Event log complete:** 13+ events (project_initialized, edge_started/converged pairs for each edge, iteration_completed events)
- [MVP] **Feature vector converged:** all trajectory edges show status: converged
- [MVP] **Acceptance criteria met:** each AC-* check in the feature vector is satisfiable against the outputs
- **Output directory correct:** code placed in `imp_<name>/` per structure binding, not at project root — REQ-TOOL-012
- **No silent failures:** event log contains no circuit_breaker_triggered events (or if it does, user investigates)

MVP Features

- [MVP] Profile mechanism: full-auto profile removes human evaluator from all edges — REQ-EVAL-001
- [MVP] Auto-mode (--auto flag): iterate without pause — REQ-UX-001
- [MVP] State-driven routing drives edge selection automatically — REQ-UX-004
- [MVP] Circuit breaker: max iterations per edge, stall detection — REQ-ITER-003
- [MVP] Event emission: complete audit trail for post-build review — REQ-LIFE-004
- [MVP] Structure binding: output directory resolved from constraints — REQ-TOOL-012, REQ-TOOL-013
- [MVP] Graph topology scaffolded by installer — REQ-TOOL-007, REQ-TOOL-014
- Acceptance criteria verification (automated check of AC-* against outputs)
- Post-build report generation (summary of what was built, coverage, gaps)

The Autopilot Contract

When a user launches an unattended build, they are asserting:

“I have provided sufficient constraints (spec, requirements, acceptance criteria, ecosystem binding, structure) that I trust the agent to construct a candidate solution. I will verify the outputs afterward. I accept that the agent may make design choices I would have made differently — the post-build review is where I exercise judgment.”

This is analogous to a CI/CD pipeline: the build runs autonomously; the human reviews the results. The quality of the output is bounded by the quality of the input constraints.

What makes this safe: 1. All outputs are in a clean directory (`imp_<name>/`) — no existing code modified 2. Event log provides full audit trail — every decision traceable 3. Deterministic evaluators still run — code must compile, tests must pass 4. Circuit breaker prevents infinite loops 5. Post-build verification is explicit — the journey includes the human review step

What makes this risky if pre-conditions are weak: 1. Vague requirements → agent fills gaps with assumptions (hallucination vector) 2. Missing constraint dimensions → agent chooses technology stack arbitrarily 3. No acceptance criteria → no way to verify fitness of output 4. No structure binding → code at project root blocks future parallel designs

Related Profile Configuration

The `full-auto` profile is the mechanism:

```
evaluators:  
  default: [agent, deterministic]      # no human  
convergence:  
  human_required_on_all_edges: false   # agent crosses review boundary  
encoding:  
  mode: autopilot                      # agent routes and decides  
  decide: F_P                          # probabilistic decisions  
iteration:  
  escalation: agent                   # self-resolve, don't block
```

2.4 UJ-3: Resume & Orient

Returning developer queries project state to understand where they are.

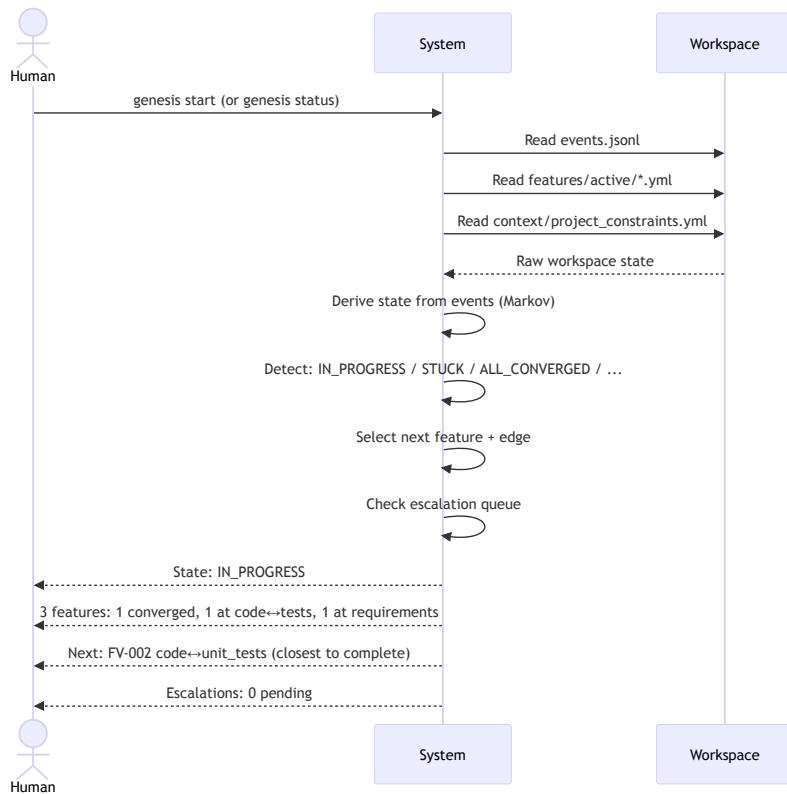


Diagram 3

Who

Developer who closed their IDE, lost context, and returns to the project.

When

- Start of a new session
- After a break (hours, days, weeks)

- After a context window reset (new conversation)

Key Characteristics

Aspect	Detail
Read-only	No iteration, no state mutation (until user acts)
State derivation	From filesystem + event log — never from stored variable
Orientation	“You are here” indicator, blocked features, pending escalations

Invariants (all hold)

- [MVP] Markov Property — state derived, not recalled
- [MVP] Event Sourcing — workspace rebuildable from events

Pre-conditions

- [MVP] Workspace exists with at least one event

Post-conditions

- User knows: current state, active features, next recommended action, pending escalations

MVP Features

- [MVP] State detection and routing — REQ-UX-001
 - [MVP] Project-wide observability (/gen-status) — REQ-UX-003
 - [MVP] Feature and edge auto-selection with reasoning — REQ-UX-004
 - [MVP] Escalation queue visibility — REQ-UX-006
 - Recovery and self-healing guidance — REQ-UX-005
-

2.5 UJ-4: Review & Approve

A human evaluator reviews work at an edge boundary.

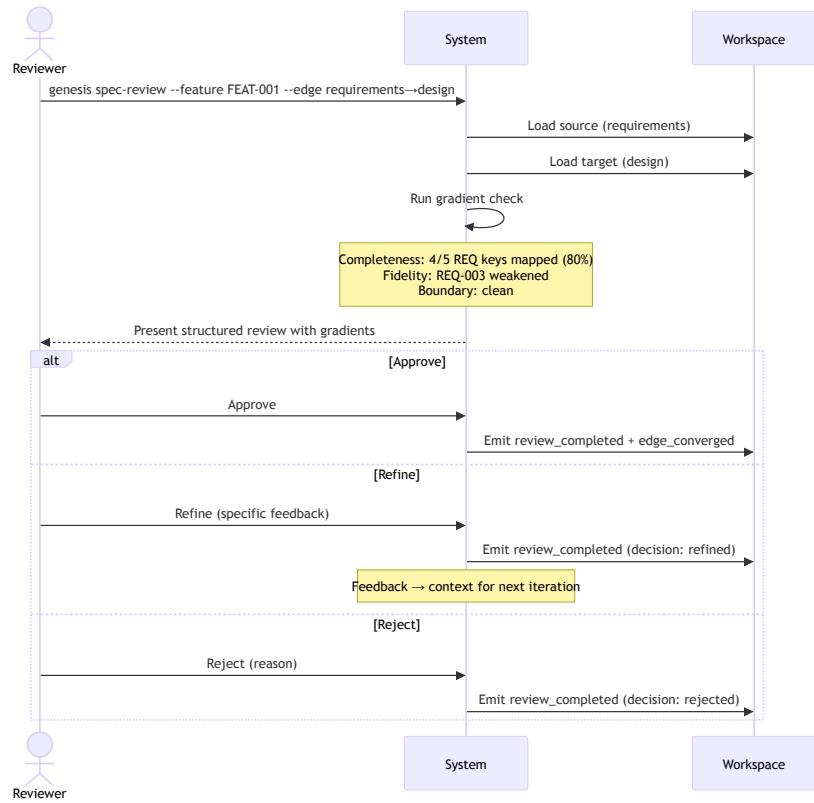


Diagram 4

Who

- Tech lead reviewing design against requirements
- Product owner reviewing UAT test scenarios
- Security reviewer checking threat model
- Any human evaluator at any edge where `human_required: true`

When

- Interactive build (UJ-1) reaches a human gate
- Unattended build (UJ-2) produced outputs that need post-build review
- Escalation queue has pending items

Key Characteristics

Aspect	Detail
Judgment	Human exercises judgment — approve, refine, reject
Gradient check	Completeness (all REQ keys mapped), fidelity (intent preserved), boundary (spec/design clean)
Feedback loop	Refinement feedback becomes context for next iteration

Invariants (all hold)

- [MVP] Review Boundary — human explicitly crosses it
- [MVP] Human Accountability — decision recorded with reasoning

Pre-conditions

- [MVP] Edge has produced a candidate (source + target assets exist)
- [MVP] Evaluator checklist has been run (automated checks done before human reviews)

Post-conditions

- [MVP] Decision recorded in feature vector + event log
- Edge converges (approve) or iterates with feedback (refine) or blocks (reject)

MVP Features

- [MVP] Human evaluator gate detection and pause — REQ-EVAL-003, REQ-UX-006
- [MVP] Gradient check (completeness, fidelity, boundary) — REQ-EDGE-003
- [MVP] Decision recording with event emission — REQ-LIFE-004
- Spec-boundary review command — REQ-EDGE-003
- Feedback-as-context for re-iteration

2.6 UJ-5: Background Sensing

Sensory service watches the project continuously. Human reviews proposals when convenient.

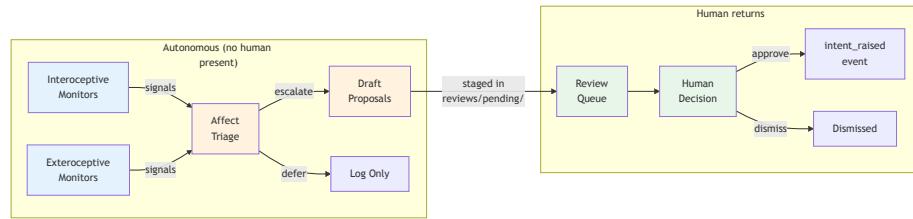


Diagram 5

Who

- Maintainer of a running system
- Ops engineer monitoring health
- Developer between active work sessions

When

- System is deployed and running
- Project is between active development phases
- Long-running maintenance mode

Key Characteristics

Aspect	Detail
Continuous	Runs on schedule, independent of iterate()
Observe-only	Service cannot modify spec, code, or emit intent events
Proposals	Draft proposals staged in review queue
Asynchronous	Human reviews when convenient, not immediately

Invariants

- Review Boundary — architectural invariant, service cannot cross it
- Sensor autonomy — observations are autonomous, changes are not

MVP Features

- This journey is **not MVP**. It requires the sensory service infrastructure.
- Future: REQ-SENSE-001, REQ-SENSE-002, REQ-SUPV-001, REQ-SUPV-002

2.7 UJ-6: CI/CD Integration

Pipeline triggers a build or validation from spec, results go to review queue.

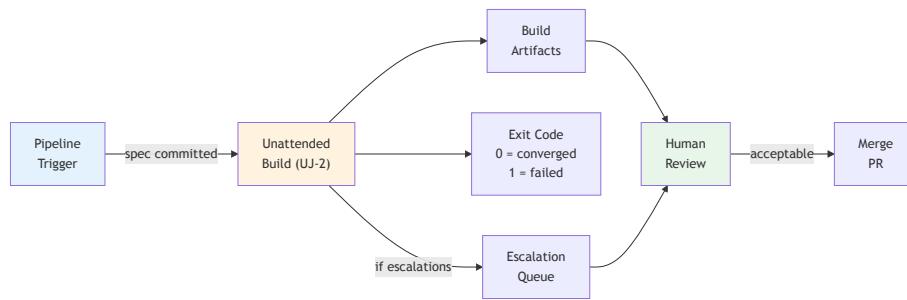


Diagram 6

Who

- CI/CD pipeline (GitHub Actions, Jenkins, etc.)
- Scheduled cron job
- Git hook on spec change

When

- Spec file committed → pipeline triggers regeneration

- Nightly build validates spec-to-code alignment
- Release pipeline runs full gap validation

Key Characteristics

Aspect	Detail
Triggered	Pipeline event, not human invocation
Headless	No interactive session, no TTY
Artifacts	Build outputs as pipeline artifacts (code, tests, events)
Exit code	Machine-readable success/failure

Invariants

- Same as UJ-2 (Unattended Build)
- Additionally: escalations guaranteed delivery to review queue (never lost in pipeline output)

MVP Features

- This journey is **not MVP** in its full form.
 - However, [MVP] the `--auto` flag and `full-auto` profile (from UJ-2) are the foundation that CI/CD builds on.
 - Future: pipeline integration, artifact packaging, exit code semantics
-

2.8 UJ-7: Observation & Monitoring

Stakeholder observes project state through a dashboard. Read-only.

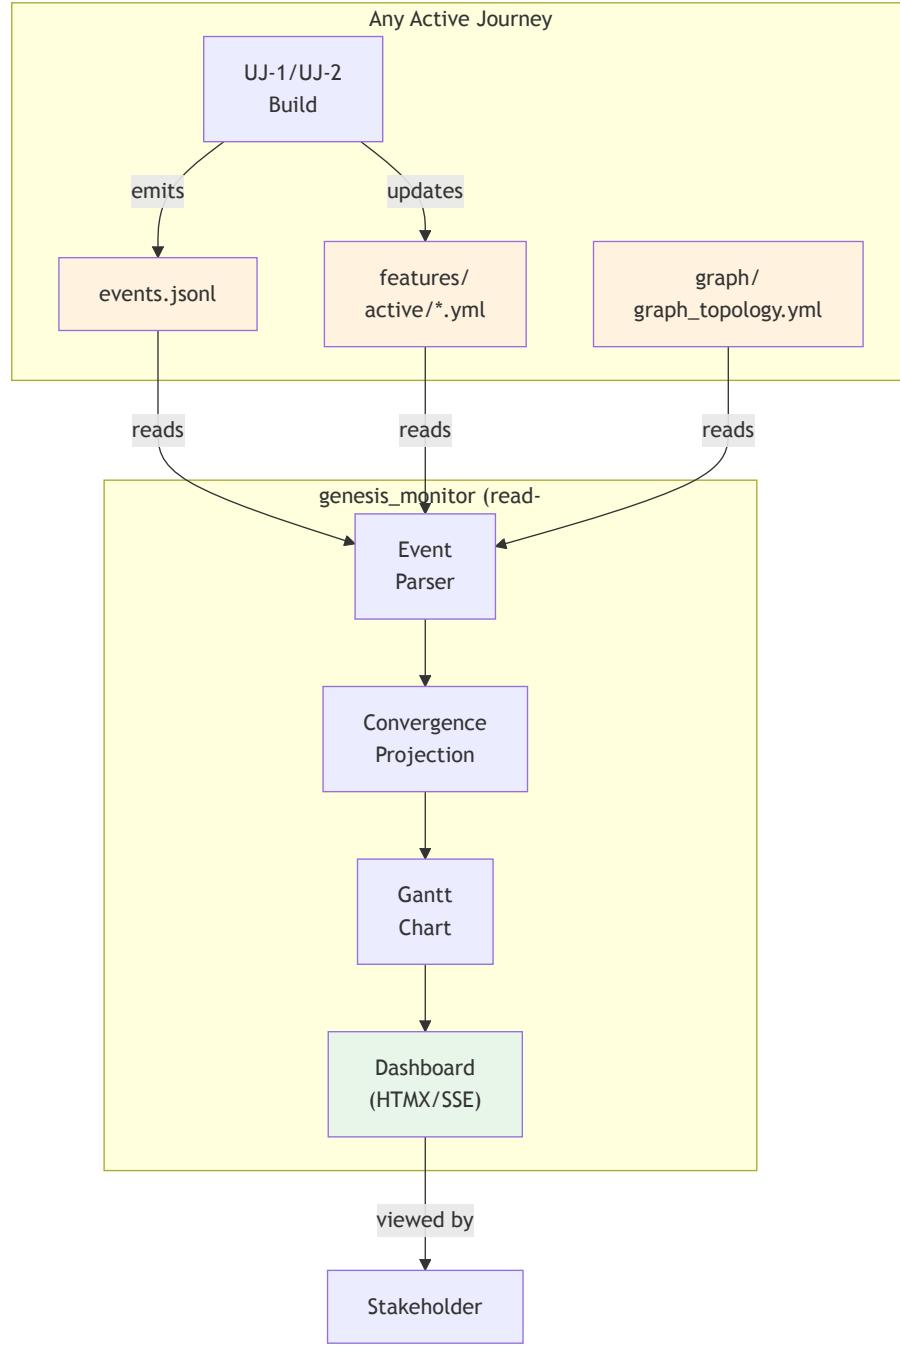


Diagram 7

Who

- Project manager checking progress
- Tech lead reviewing convergence across features
- Developer checking if their build completed
- Anyone with dashboard access

When

- During active development (monitor build progress)
- During unattended build (watch convergence in real-time)
- During maintenance (check project health)

Key Characteristics

Aspect	Detail
Read-only	Dashboard reads event stream, never writes
Real-time	SSE updates as events arrive
Multi-feature	Shows all features, their trajectories, convergence state
Integration contract	Requires: events.jsonl, features/active/*.yml, graph/graph_topology.yml

Invariants

- Observer does not disturb the observed system
- Dashboard state derivable from event log (same Markov property)

MVP Features

- [MVP] Event log as integration contract — REQ-LIFE-004
 - [MVP] Graph topology scaffolded by installer (monitor requires it) — REQ-TOOL-007, REQ-TOOL-014
 - [MVP] Feature vector status visible — REQ-UX-003
 - Convergence Gantt chart
 - Regime classification (reflex/affect/conscious event tagging)
 - Multi-project discovery
-

3. MVP Feature Summary

Features tagged [MVP] across all journeys, deduplicated and grouped by category.

3.1 Core Mechanism

Feature	REQ Keys	Journeys
State-driven routing	REQ-UX-001	UJ-1, UJ-2, UJ-3
Progressive disclosure	REQ-UX-002	UJ-1
Feature and edge auto-selection	REQ-UX-004	UJ-1, UJ-2, UJ-3
Evaluator checklist execution	REQ-EVAL-001, REQ-EVAL-002	UJ-1, UJ-2, UJ-4

Feature	REQ Keys	Journeys
Event emission at every boundary	REQ-LIFE-004	All
Markov state derivation	REQ-GRAFH-003	All

3.2 Human Interaction

Feature	REQ Keys	Journeys
Human gate detection and pause	REQ-EVAL-003, REQ-UX-006	UJ-1, UJ-4
Gradient check (completeness, fidelity, boundary)	REQ-EDGE-003	UJ-1, UJ-4
Escalation queue visibility	REQ-UX-006	UJ-1, UJ-3
Project-wide observability (/gen-status)	REQ-UX-003	UJ-1, UJ-3, UJ-7

3.3 Autonomous Operation

Feature	REQ Keys	Journeys
full-auto profile (no human evaluators)	REQ-EVAL-001	UJ-2, UJ-6
Auto-mode (--auto flag)	REQ-UX-001	UJ-2, UJ-6
Circuit breaker (stall + max iteration)	REQ-ITER-003	UJ-2, UJ-6

3.4 Structure & Scaffolding

Feature	REQ Keys	Journeys
Workspace initialisation	REQ-TOOL-011	All
Graph topology scaffolded by installer	REQ-TOOL-007, REQ-TOOL-014	UJ-2, UJ-7

Feature	REQ Keys	Journeys
Structure binding (output directory)	REQ-TOOL-012, REQ-TOOL-013	UJ-2

3.5 Traceability

Feature	REQ Keys	Journeys
REQ key traceability (intent → telemetry)	REQ-LIFE-002	UJ-1, UJ-2
Feature vector trajectory tracking	REQ-FEAT-001	UJ-1, UJ-2, UJ-3
Event-sourced audit trail	REQ-LIFE-004	All

MVP Feature Map

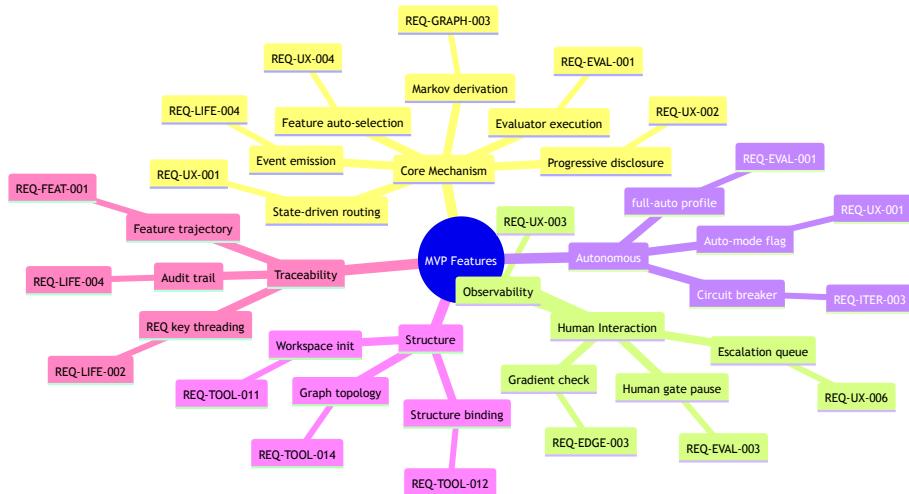


Diagram 8

4. Validation Scenarios

End-to-end scenarios that exercise the user journeys. Each scenario is a gap-analysis driver — mismatches between the scenario and the current implementation surface work items.

Scenario Format

SC-{NNN}: {Title}

Journey: UJ-{N}

****Profile**:** {projection profile}
****Actor**:** {who is performing the scenario}
****Precondition**:** {starting state}
****REQ Keys Exercised**:** {list}

Steps

| # | User Action | System Response | REQ Keys | Gap Status |

Gaps / Invariant Checks

SC-001: First-Time Developer — Green Field to Running System

Journey: UJ-1 (Interactive Build) **Profile:** standard **Actor:** Developer with no prior Genesis experience, has a new project idea **Precondition:** Empty directory, development environment (Python, git). Genesis installed. **REQ Keys Exercised:** REQ-INTENT-001, REQ-INTENT-002, REQ-UX-001, REQ-UX-002, REQ-UX-003, REQ-UX-004, REQ-UX-005, REQ-GRAFH-001, REQ-GRAFH-002, REQ-ITER-001, REQ-ITER-002, REQ-EVAL-001, REQ-EVAL-002, REQ-EVAL-003, REQ-CTX-001, REQ-CTX-002, REQ-FEAT-001, REQ-FEAT-002, REQ-EDGE-001, REQ-EDGE-002, REQ-EDGE-003, REQ-EDGE-004, REQ-LIFE-001, REQ-LIFE-002, REQ-LIFE-003, REQ-LIFE-004, REQ-LIFE-005, REQ-TOOL-001, REQ-TOOL-002, REQ-TOOL-003, REQ-TOOL-004, REQ-TOOL-005, REQ-TOOL-006, REQ-TOOL-008, REQ-TOOL-009, REQ-TOOL-010, REQ-SUPV-001, REQ-SUPV-002, REQ-SENSE-001, REQ-SENSE-002

The canonical end-to-end scenario. A developer takes a project from nothing to a running, monitored system using only Genesis tooling. Every edge of the standard profile graph is traversed.

Steps

#	User Action	System Response	REQ Keys
1	Install Genesis — single command	Tooling available, version verifiable, commands/agents/configs loaded	REQ-TOOL-001
2	Start in empty directory — genesis start	Progressive init: detect project name, language, test runner. Create .ai-workspace/, emit project_initialized event.	REQ-UX-001, REQ-UX-002, REQ-TOOL-003
3	Describe intent	Write specification/INTENT.md. Emit intent_raised. State → NO_FEATURES.	REQ-INTENT-001, REQ-

#	User Action	System Response	REQ Keys
			INTENT-002, REQ-LIFE-005
4	System creates feature vectors	Parse intent, propose feature decomposition. Create feature vectors. State → IN_PROGRESS.	REQ-FEAT-001, REQ-FEAT-002, REQ-TOOL-004
5	Iterate intent→requirements	Generate structured requirements. Human reviews via gradient check. Edge converges.	REQ-ITER-001, REQ-EVAL-001, REQ-EVAL-003
6	Iterate requirements→design	Prompt for constraint dimensions. Generate ADRs. Human approves. Edge converges.	REQ-CTX-001, REQ-EDGE-003, REQ-EVAL-003
7	Iterate design→code + code↔unit_tests	TDD cycle. Deterministic evaluators. Edge converges.	REQ-EDGE-001, REQ-EDGE-004, REQ-EVAL-001
8	Iterate code→uat_tests	BDD acceptance tests. Edge converges.	REQ-EDGE-002, REQ-

#	User Action	System Response	REQ Keys
			EVAL-002
9	Iterate uat_tests→cicd	CI/CD pipeline config. Deterministic evaluator. Edge converges.	REQ-LIFE-001
10	Iterate cicd→running_system→telemetry	Deploy. Telemetry tagging. Edge converges.	REQ-LIFE-002, REQ-LIFE-004
11	Observe — genesis_monitor dashboard	Feature progress, convergence Gantt, event stream, feedback loop.	REQ-LIFE-003, REQ-SENSE-001

Post-Condition

- All features converged. events.jsonl complete. REQ traceability end-to-end. /gen-status shows ALL_CONVERGED.

Gap Analysis Summary

Steps: 11 | COVERED: 2 | GAP-001: 1 | GAP-002: 1 | GAP-003: 6 |

PARTIAL: 1

Methodology completeness: 18% executable today

Invariant Checks

- Evaluator Existence — every edge traversal invokes evaluators
- Markov Property — each step depends only on current asset state
- Convergence Monotonicity — delta decreases within each edge
- IntentEngine Composition — step 11 closes the loop

SC-002: Returning Developer — Resume After Context Loss

Journey: UJ-3 (Resume & Orient) **Profile:** standard **Actor:** Developer who returns with no memory of where they left off **Precondition:** .ai-workspace/ with 3 features at various stages. ~50 events. No IDE state. **REQ Keys Exercised:** REQ-UX-001, REQ-UX-003, REQ-UX-004, REQ-UX-005, REQ-TOOL-002, REQ-TOOL-009, REQ-GRAFH-003, REQ-SUPV-001

Steps

#	User Action	System Response	REQ Keys	Gap Status
1	Open project — genesis start	SessionStart hook: validate events, check consistency, detect staleness.	REQ-UX-005, REQ-TOOL-002	COVERED
2	State detection	IN_PROGRESS. 3 features displayed with “you are here” indicators.	REQ-UX-001, REQ-UX-003	COVERED
3	Check status — genesis status	Feature trajectories, convergence markers, rollup, Gantt.	REQ-TOOL-002, REQ-UX-003	COVERED
4	Resume iteration	Load context from events. Resume on selected edge.	REQ-UX-004, REQ-GRAF-003	GAP-003
5	Check blockers	Blocked features with reasons and recommended unblocking actions.	REQ-UX-005, REQ-SUPV-001	COVERED

Invariant Checks

- Markov Property** — central: state derived from filesystem + events, not from session memory
-

SC-003: Stuck Feature — Escalation and Recovery

Journey: UJ-1 (Interactive Build) **Profile:** standard **Actor:** Developer whose feature has been iterating with no progress **Precondition:** Feature on code+unit_tests. Delta=3 unchanged for 4 iterations. 3 failing checks. **REQ Keys Exercised:** REQ-UX-005, REQ-SUPV-001, REQ-SUPV-002, REQ-LIFE-005, REQ-LIFE-006, REQ-FEAT-002, REQ-FEAT-003, REQ-ITER-003

Steps

#	User Action	System Response	REQ Keys	Gap Status
1	Start	State: STUCK. Display stuck feature with delta history.	REQ-UX-005, REQ-SUPV-001	COVERED
2	View escalation queue	1 stuck feature (critical). Context and recommended actions.	REQ-SUPV-001, REQ-SUPV-002	COVERED

#	User Action	System Response	REQ Keys	Gap Status
3	Spawn discovery vector	Create spike child. Emit spawn_created.	REQ-FEAT-002, REQ-LIFE-005	GAP-003
4	Investigate via spike	Spike iterates, investigates root cause, produces report.	REQ-ITER-003	GAP-003
5	Fold back findings	Update parent context. Emit spawn_folded_back. Delta decreases.	REQ-FEAT-002, REQ-LIFE-005	GAP-003
6	Parent unsticks	Failing checks resolve. Edge converges.	REQ-ITER-001, REQ-ITER-002	GAP-003

Invariant Checks

- Evaluator Existence — stuck detection is itself an evaluator (meta-level)
 - Markov Property — spike depends on current parent state
 - Convergence Monotonicity — delta stuck, then decreases after fold-back
 - IntentEngine Composition — stuck → escalate → spawn → fold-back → resume
-

SC-004: Multi-Feature Coordination — Parallel Work with Dependencies

Journey: UJ-1 (Interactive Build) **Profile:** standard **Actor:** Developer with 3 features, one blocking another **Precondition:** 3 features active. REQ-F-API-001 blocked on REQ-F-DB-001. REQ-F-UI-001 independent. **REQ Keys Exercised:** REQ-FEAT-002, REQ-FEAT-003, REQ-UX-004, REQ-COORD-001 through REQ-COORD-005

Steps

#	User Action	System Response	REQ Keys	Gap Status
1	Start	Select REQ-F-DB-001 (closest + blocks API). Show dependency graph.	REQ-UX-004, REQ-FEAT-002	COVERED
2	Complete DB schema	DB feature converges. Emit edge_converged.	REQ-ITER-001, REQ-ITER-002	GAP-003
3	Dependency unblocks	API feature unblocked. Routing shifts. Explain decision.	REQ-FEAT-002, REQ-FEAT-003	COVERED

#	User Action	System Response	REQ Keys	Gap Status
4	Status check	Three features at different positions. Gantt shows parallel tracks.	REQ-TOOL-002, REQ-UX-003	COVERED
5	Multi-agent coordination	Event-sourced edge assignment. Work isolation.	REQ-COORD-001 through 005	GAP-006

Invariant Checks

- Evaluator Existence — each edge has evaluators
 - Markov Property — API depends on DB's converged state, not how DB converged
-

SC-005: Spec-Boundary Review — Human Gate at Requirements→Design

Journey: UJ-4 (Review & Approve) **Profile:** full **Actor:** Tech lead reviewing design against requirements **Precondition:** Requirements converged. Design drafted.

human_required: true. **REQ Keys Exercised:** REQ-EVAL-003, REQ-EDGE-003, REQ-CTX-001, REQ-CTX-002, REQ-LIFE-005, REQ-UX-006

Steps

#	User Action	System Response	REQ Keys	Gap Status
1	Auto-mode pauses at human gate	Announce: "Human gate: requirements→design requires spec-boundary review."	REQ-UX-006, REQ-EVAL-003	COVERED
2	Invoke spec review	Gradient check: completeness (80%), fidelity (REQ-003 weakened), boundary (clean).	REQ-EVAL-003, REQ-EDGE-003	COVERED
3	Reviewer finds gap	"REQ-F-AUTH-002 has no design binding."	REQ-CTX-001	COVERED
4	Refine — specific feedback	Record decision. Emit review_completed. Edge does NOT converge.	REQ-EVAL-003, REQ-LIFE-005	GAP-003
5	Design iterates with feedback	Revised design. Re-run gradients. Delta decreases.	REQ-ITER-001, REQ-ITER-002	GAP-003

#	User Action	System Response	REQ Keys	Gap Status
6	Second review — approve	All gradients pass. Edge converges.	REQ-EVAL-003	GAP-003

Invariant Checks

- Evaluator Existence — human evaluator is the focus
 - Markov Property — review depends on current design + requirements
 - Convergence Monotonicity — delta decreases from iteration 1 to 2
-

SC-006: Projection Switch — Standard to Spike Mid-Flight

Journey: UJ-1 (Interactive Build) **Profile:** standard → spike **Actor:** Developer who discovers a requirement needs feasibility investigation **Precondition:** Feature at design→code. ML approach may not work. **REQ Keys Exercised:** REQ-FEAT-001, REQ-FEAT-002, REQ-FEAT-003, REQ-ITER-003, REQ-TOOL-004, REQ-LIFE-005

Steps

#	User Action	System Response	REQ Keys	Gap Status
1	Spawn spike	genesis spawn --type spike --parent FEAT-001 --reason feasibility	REQ-FEAT-002, REQ-TOOL-004	COVERED
2	Spike created	Spike profile (code+tests only, time-boxed). Parent paused.	REQ-FEAT-001, REQ-ITER-003	COVERED
3	Spike iterates	Reduced graph. Evaluate feasibility. Time-box: 4 hours.	REQ-FEAT-003, REQ-ITER-001	GAP-003
4	Spike concludes	Feasibility report, prototype, go/no-go recommendation.	REQ-ITER-002	GAP-003
5	Fold back	If go: resume with prototype. If no-go: pivot. Emit spawn_folded_back.	REQ-FEAT-002, REQ-LIFE-005	GAP-003

Invariant Checks

- Evaluator Existence — spike has evaluators (relaxed)
- Markov Property — spike depends on current parent state
- IntentEngine Composition — no-go → intent_raised → design pivot

SC-007: Telemetry Feedback — Production Anomaly Closes the Loop

Journey: UJ-5 (Background Sensing) + UJ-1 (Interactive Build) **Profile:** full Actor: Operations engineer monitoring a running system **Precondition:** System deployed, emitting feature-keyed telemetry. genesis_monitor watching. **REQ Keys Exercised:** REQ-LIFE-002, REQ-LIFE-003, REQ-LIFE-004, REQ-SENSE-001, REQ-SENSE-002, REQ-SUPV-001, REQ-SUPV-002, REQ-INTENT-003

Steps

#	User Action	System Response	REQ Keys	Gap Status
1	Anomaly detected — latency spike	Exteroceptive monitor detects tolerance breach. Emit exteroceptive_signal.	REQ-SENSE-001, REQ-SUPV-002	GAP-008
2	IntentEngine classifies	bounded_nonzero ambiguity. Output: spec_event_log. Emit affect_triage.	REQ-SUPV-001	GAP-003
3	Intent raised	intent_raised: trigger=runtime_feedback, severity=high.	REQ-LIFE-005, REQ-INTENT-003	GAP-003
4	Developer sees intent	Escalation queue: 1 high-severity item. Recommended: spawn hotfix.	REQ-SUPV-001, REQ-LIFE-003	COVERED
5	Spawn hotfix	Hotfix vector, profile hotfix (code + tests + cicd only).	REQ-FEAT-002, REQ-LIFE-005	GAP-003
6	Hotfix iterates and deploys	Minimal edges, fast convergence. Deploy fix.	REQ-ITER-001, REQ-LIFE-001	GAP-003
7	Loop closes — telemetry confirms fix	p99 restored. Hotfix converges.	REQ-LIFE-002, REQ-LIFE-003	GAP-003, GAP-008

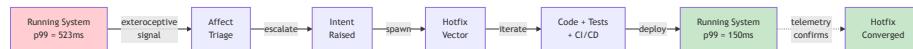


Diagram 9

Invariant Checks

- Evaluator Existence — hotfix edges have evaluators

- Markov Property — hotfix depends on current telemetry state
 - Convergence Monotonicity — deviation → fix → resolved
 - IntentEngine Composition — anomaly → signal → triage → intent → spawn → iterate → deploy → confirm
-

5. Gap Registry

Consolidated across all scenarios. Each gap appears once; scenarios reference it.

Gap ID	Title	Scenarios	Journeys	Severity	Drives
GAP-001	No installability requirement	SC-001	UJ-1	High	REQ-TOOL-011 (now implemented)
GAP-002	No automated feature decomposition from intent	SC-001	UJ-1	Low	REQ-FEAT-001, REQ-UX-001
GAP-003	No executable orchestration engine	SC-001..007	All	Critical	Phase 2 — see §5.1
GAP-004	genesis_monitor schema drift + single-project discovery	SC-001	UJ-7	Medium	genesis_monitor upgrade
GAP-005	No automated root-cause investigation for stuck features	SC-003	UJ-1	Low	Spike diagnostic heuristics
GAP-006	Multi-agent coordination is spec-only	SC-004	UJ-1	Low	Phase 3 (post-engine)
GAP-007	No fold-back command	SC-006	UJ-1	Medium	New /gen-fold-back command
GAP-008	No runtime telemetry integration	SC-007	UJ-5	Medium	REQ-SENSE-001 scope clarification

Gap Impact Matrix

	SC-001	SC-002	SC-003	SC-004	SC-005	SC-006	SC-007
GAP-001	X						
GAP-002	X						
GAP-003	X	X	X	X	X	X	X
GAP-004	X						
GAP-005			X				
GAP-006				X			
GAP-007					X		
GAP-008						X	

GAP-003 dominates. It appears in every scenario. This is expected — the methodology is fully specified but not yet executable. The remaining gaps are spec-level and independently addressable.

Note: GAP-001 (installability) has been resolved — REQ-TOOL-011 was added and `gen-setup.py` implements it.

5.1 The F_P Question (GAP-003 Commentary)

GAP-003 — “no executable engine” — deserves deeper analysis because it may not be a gap at all. It depends on which encoding is considered primary:

If F_D (deterministic) is the target encoding: GAP-003 is a real gap. The specs describe a state machine, and no code implements that state machine. The human must manually drive the loop.

If F_P (probabilistic) is the target encoding: GAP-003 is a *feature*. The specs are the program — loaded into an LLM’s context, the LLM executes them. The iterate agent markdown IS the code. This is the current de facto architecture.

If both encodings are valid (ADR-017 functor model): the question is which edges benefit from F_D and which from F_P. State detection, event emission, and evaluator checklists are good F_D candidates. Asset generation (requirements, design, code) is inherently F_P. The engine is a hybrid — and partially exists today.

This is a design decision that belongs in implementation design docs (per spec/design separation). The scenarios must be satisfiable regardless of encoding choice.

6. Scenario Backlog

Future scenarios to add as the methodology matures:

ID	Title	Journey	Profile	Key Focus
SC-008	Enterprise Compliance — Audit Trail Generation	UJ-1	full	REQ-LIFE-004, regulatory traceability

ID	Title	Journey	Profile	Key Focus
SC-009	Library Release — Full Profile with Versioning	UJ-1	full	REQ-TOOL-008, semver, changelog
SC-010	Cross-Platform — Same Spec, Claude + Gemini Implementations	UJ-2	standard	Multi-tenancy, spec/design separation
SC-011	Methodology Self-Application — Genesis Builds Genesis	UJ-2	full	\$0 recursive compliance, dogfooding
SC-012	Time-Boxed Exploration — PoC with Expiry	UJ-1	poc	Time-box mechanics, fold-back-or-discard
SC-013	Context Window Pressure — Large Project Scaling	UJ-1	standard	REQ-CTX-002, context selection
SC-014	Ecosystem Shock — Major Dependency Breaks	UJ-5	standard	REQ-INTENT-003, exteroceptive response
SC-015	CI/CD Pipeline — Build from Spec on Commit	UJ-6	full-auto	Pipeline integration, exit codes

Appendix A: test06 as UJ-2 Reference Implementation

The `data_mapper.test06` build (2026-02-23) is the first complete execution of UJ-2. Key facts:

Aspect	Value
Spec	CDME (Categorical Data Mapping & Computation Engine)
Language	Scala 2.13 + Spark 3.5

Aspect	Value
Profile	full-auto (custom, in .ai-workspace/profiles/)
Edges traversed	5 (intent→requirements→design→code↔unit_tests→UAT)
Events emitted	13
Tests passing	95/95
Compiled artifacts	352 .class files
REQ coverage	29/30 files tagged
Human interaction	Zero during build
Post-build verification	Manual: compile check, test run, REQ grep, event log review

Gaps discovered during test06 that drove spec changes: - GAP-042: Installer missing `graph_topology.yml` → REQ-TOOL-007 strengthened, REQ-TOOL-014 added - GAP-043: No constraint enforcing folder structure → REQ-TOOL-012 added - GAP-044: No output directory binding → REQ-TOOL-013 added - Profile mechanism: `evaluator_overrides` in `project_constraints.yml` is unimplemented; profiles are the actual mechanism for removing human evaluators

These gaps were fed back into the spec and installer — demonstrating the methodology's own feedback loop (telemetry → intent → spec modification).